



COSE474 Deep Learning

Lecture 13: Generative Adversarial Networks (GANs)

Seungryong Kim

Computer Vision Lab. (CVLAB)

Department of Computer Science and Engineering

Korea University

Progress for Image Generation

2014



GAN

2015



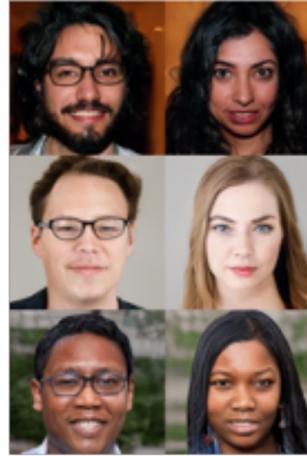
DCGAN

2017



PG-GAN

2018



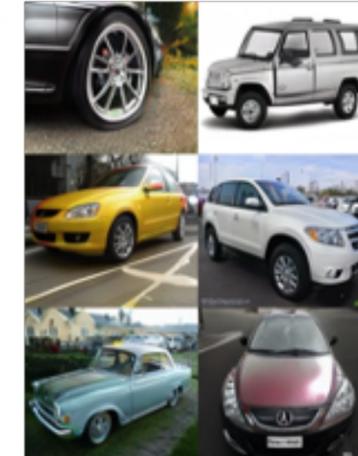
StyleGAN

2018



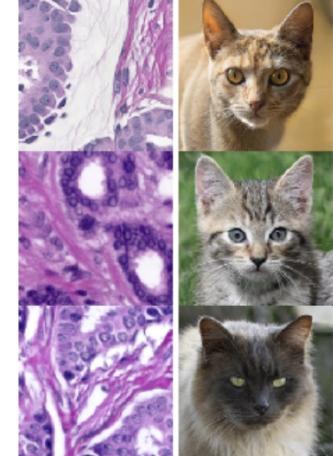
BigGAN

2019



StyleGANv2

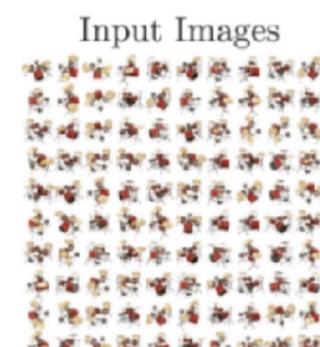
2020



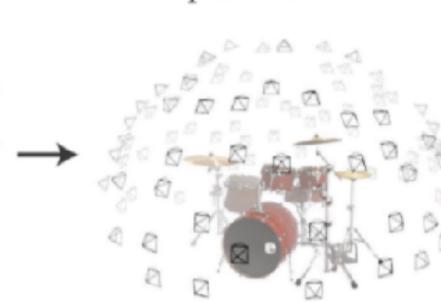
StyleGAN-ADA

2020: NeRF (Neural Radiance Fields)

2021: OpenAI DALLE (VQ-VAE)
Arm chair in shape of avocado



Optimize NeRF



Render new views



What is Generative Model?

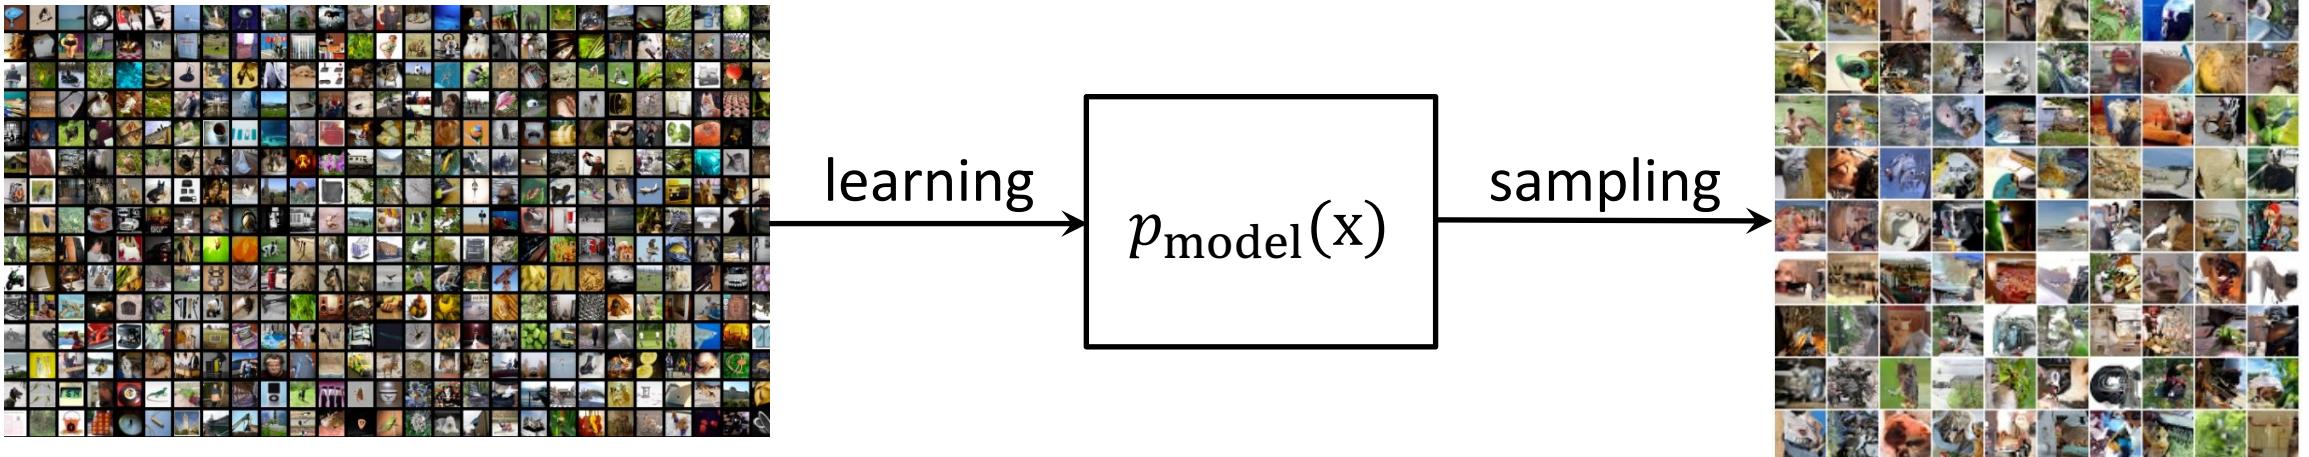
Generative Models

Discriminative model vs. Generative model

Model	Training	Test	Supervision
Discriminative model	Estimate $P(y x)$	$f: x \rightarrow y$	Supervised learning
Generative model	Estimate $P(x)$ or $P(x,y)$	$f: \text{seed} \rightarrow x$ $f: \text{seed}, y \rightarrow x$ $f: \text{seed} \rightarrow x, y$	Unsupervised learning

Generative Models

Given train set, generate new samples from same distribution.



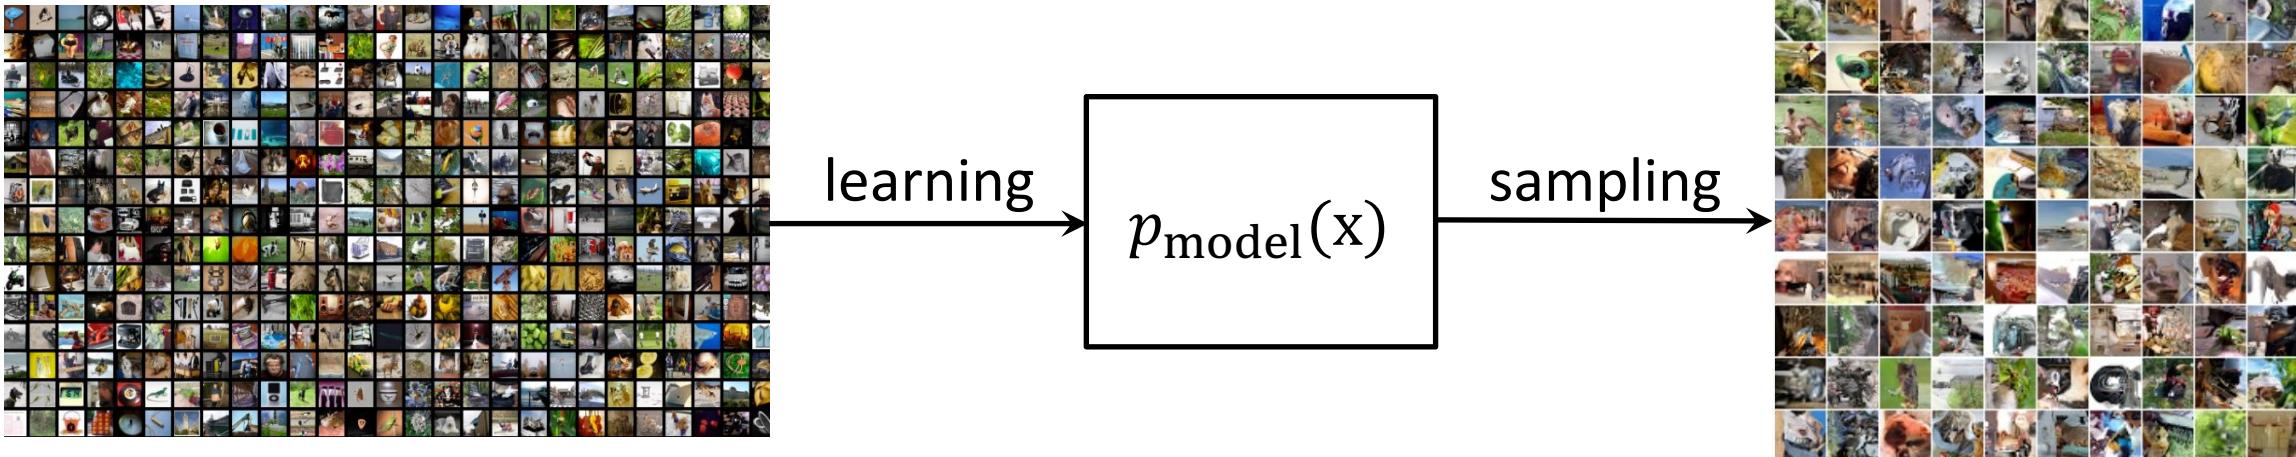
Training data $\sim p_{\text{data}}(x)$

- **Objectives:**

1. Learn $p_{\text{model}}(x)$ that approximates $p_{\text{data}}(x)$.
2. Sampling new x from $p_{\text{model}}(x)$.

Generative Models

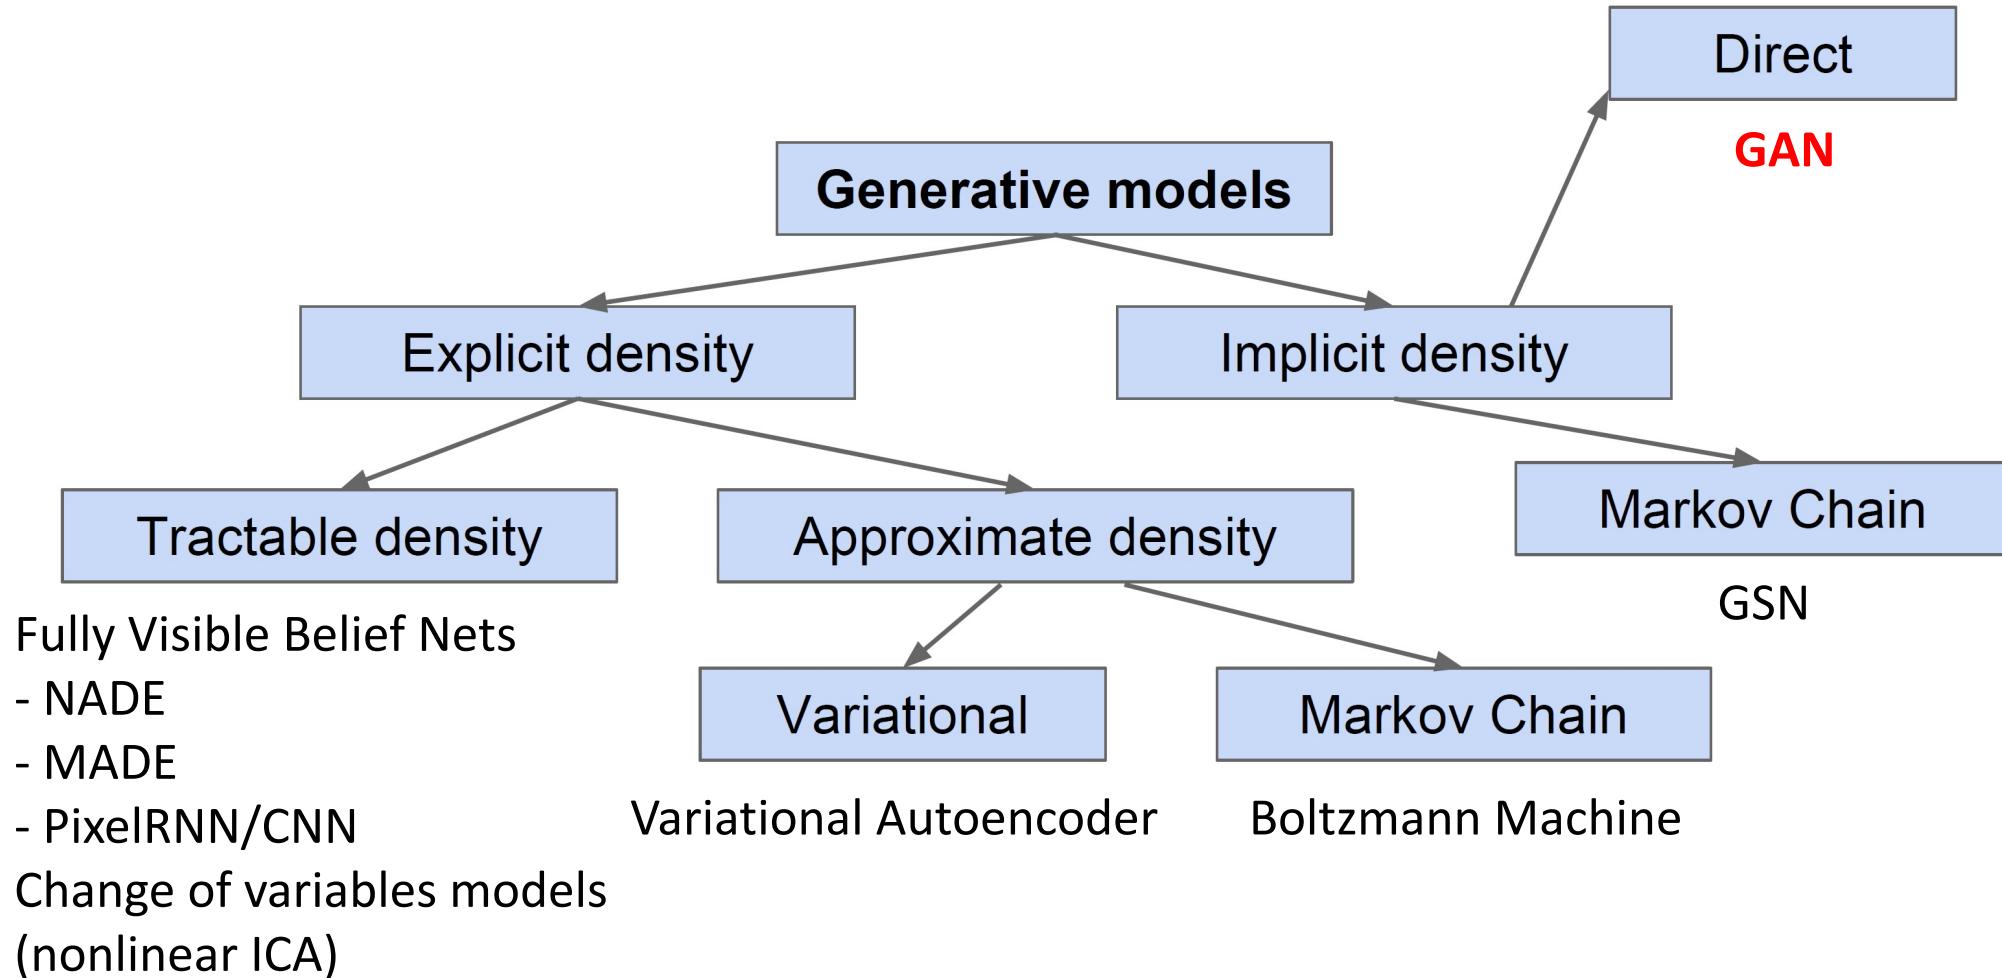
Given train set, generate new samples from same distribution.



Training data $\sim p_{\text{data}}(x)$

- **Formulate as density estimation problems:**
 - **Explicit density estimation:** explicitly define and solve for $p_{\text{model}}(x)$.
 - **Implicit density estimation:** learn model that can sample from $p_{\text{model}}(x)$ without explicitly defining it.

Taxonomy of Generative Models



Generative Adversarial Networks (GANs)

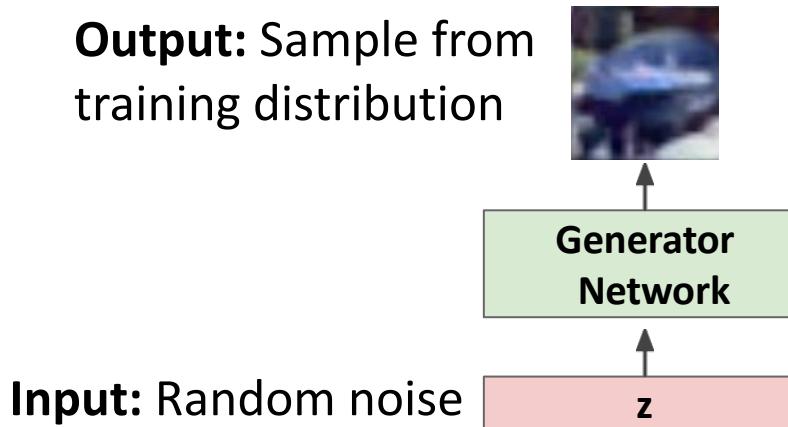
Problem: Want to sample from complex, high-dimensional training distribution. No direct (or explicit) way to do this!

Solution: Sample from a simple distribution we can easily sample from, e.g., random noise. Learn transformation to training distribution.

Q: What can we use to present this complex transformation?

A: A neural network!

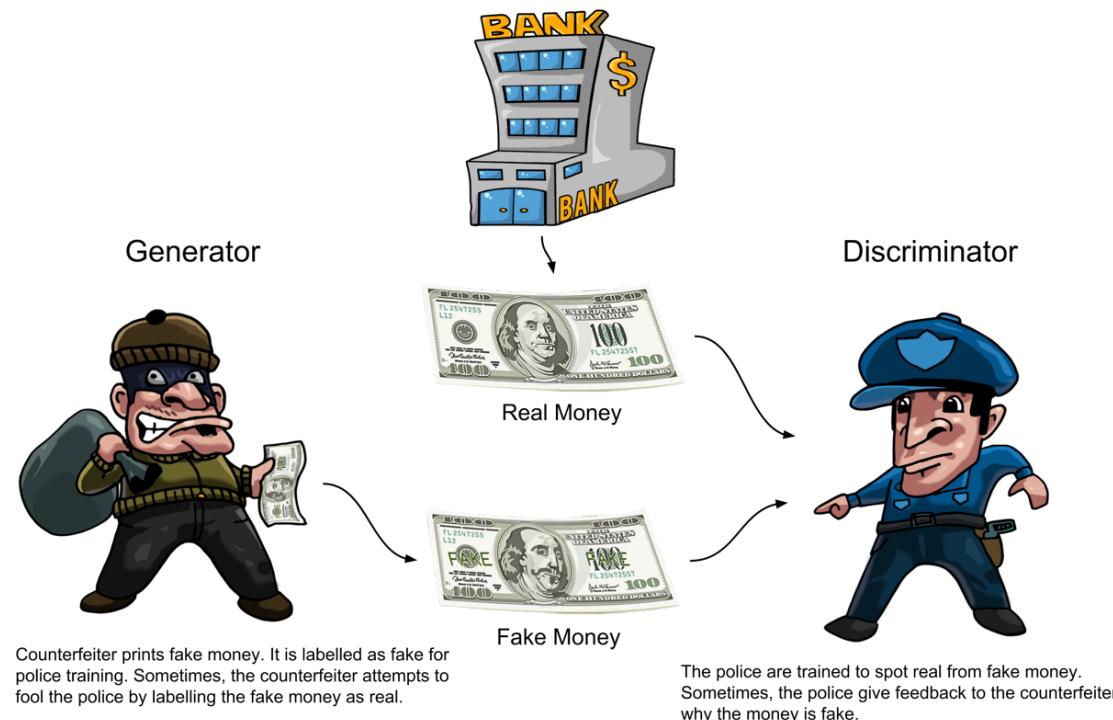
Output: Sample from training distribution



Generative Adversarial Networks (GANs)

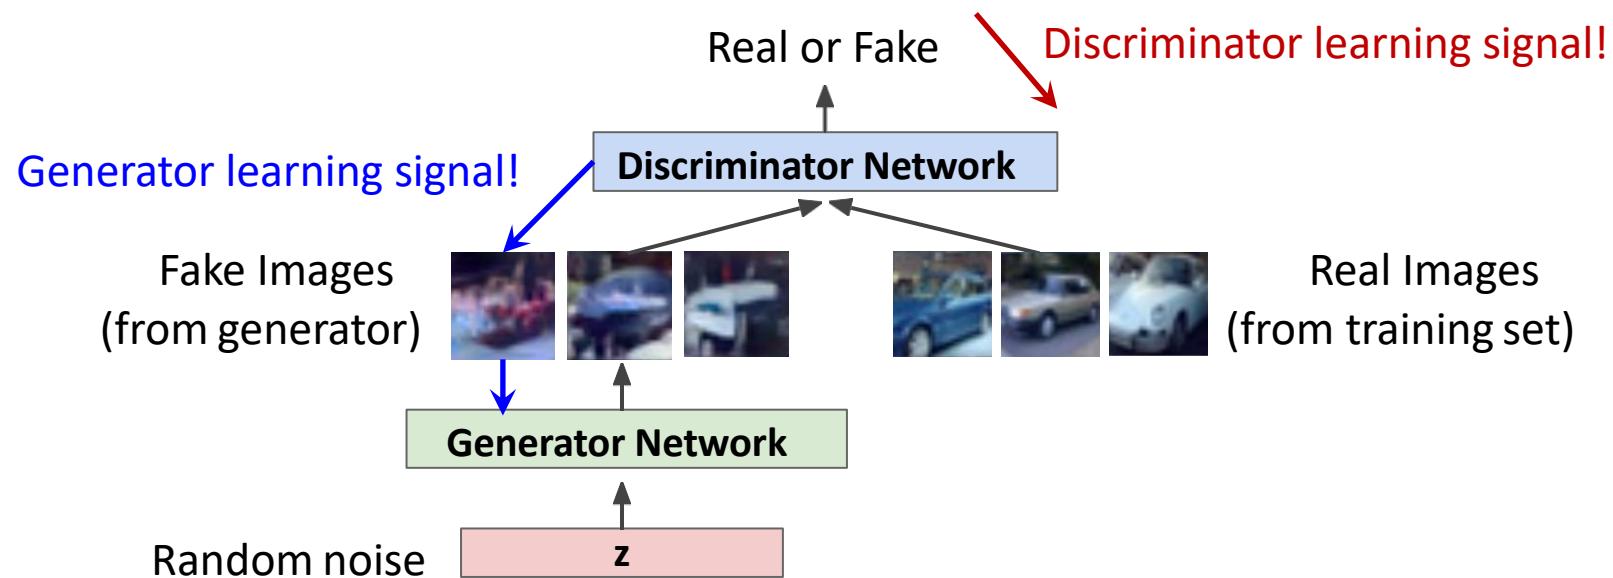
But we don't know which sample z maps to which training image
→ can't learn by reconstructing training images

Solution: Use a discriminator network to tell whether the generator image is within data distribution or not!



Training GANs: Two-player Game

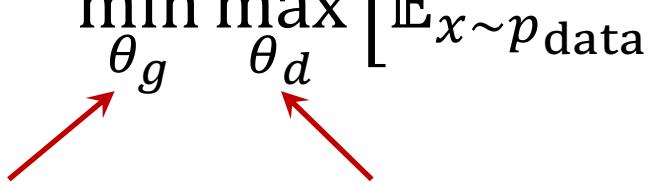
- **Discriminator network:** Try to distinguish between real and fake images
- **Generator network:** Try to fool the discriminator by generating real-looking images.



Training GANs: Two-player Game

- **Discriminator network:** Try to distinguish between real and fake images
- **Generator network:** Try to fool the discriminator by generating real-looking images.
- Train jointly in **minimax game!**
- Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$


Generator Objective Discriminative Objective

Training GANs: Two-player Game

- **Discriminator network:** Try to distinguish between real and fake images
- **Generator network:** Try to fool the discriminator by generating real-looking images.
- Train jointly in **minimax game!**
- Minimax objective function:

Discriminator outputs likelihood in (0,1) of real image

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{\text{data}}} \underbrace{\log D_{\theta_d}(x)}_{\text{Discriminator output for real data } x} + \mathbb{E}_{z \sim p(z)} \underbrace{\log(1 - D_{\theta_d}(G_{\theta_g}(z)))}_{\text{Discriminator output for generated fake data } G(x)} \right]$$

Discriminator
output for real data x

Discriminator output for
generated fake data $G(x)$

Training GANs: Two-player Game

Minimax Objective Function:

Discriminator outputs likelihood in (0,1) of real image

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{\text{data}}} \underbrace{\log D_{\theta_d}(x)}_{\text{Discriminator output for real data } x} + \mathbb{E}_{z \sim p(z)} \underbrace{\log(1 - D_{\theta_d}(G_{\theta_g}(z)))}_{\text{Discriminator output for generated fake data } G(x)} \right]$$

- Discriminator (θ_d) wants to **maximize objective** such that $D(x)$ is close to 1 (real) and $D(G(z))$ is close to 0 (fake).
- Generator (θ_g) wants to **minimize objective** such that $D(G(z))$ is close to 1 (discriminator is fooled into thinking generated $G(z)$ is real)

Training GANs: Two-player Game

Minimax Objective Function:

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

- Alternate between:
 - 1. Gradient ascent** on discriminator
 - 1. Gradient descent** on generator

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

Training GANs: Two-player Game

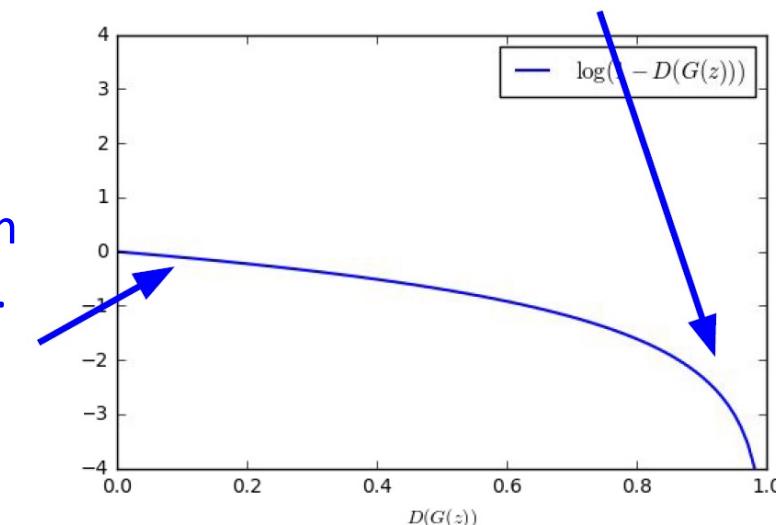
Minimax Objective Function:

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

In practice, optimizing this generator objective does not work well!

When sample is likely fake, want to learn from it to improve generator. But gradient in this region is relatively flat!

Gradient signal dominated by region where sample is already good



Training GANs: Two-player Game

Minimax Objective Function:

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

- Alternate between:
 1. **Gradient ascent** on discriminator

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

- 1. **Gradient ascent** on generator

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log D_{\theta_d}(G_{\theta_g}(z))$$

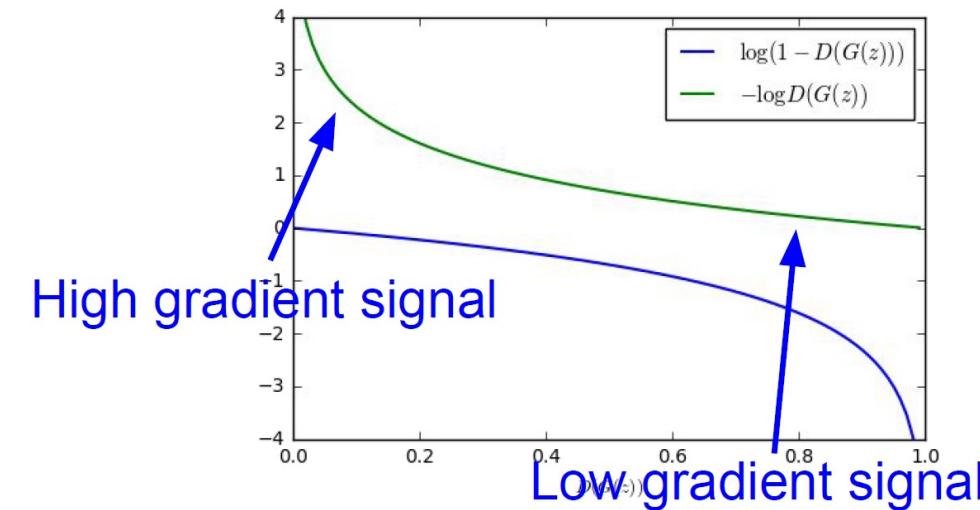
Training GANs: Two-player Game

Minimax Objective Function:

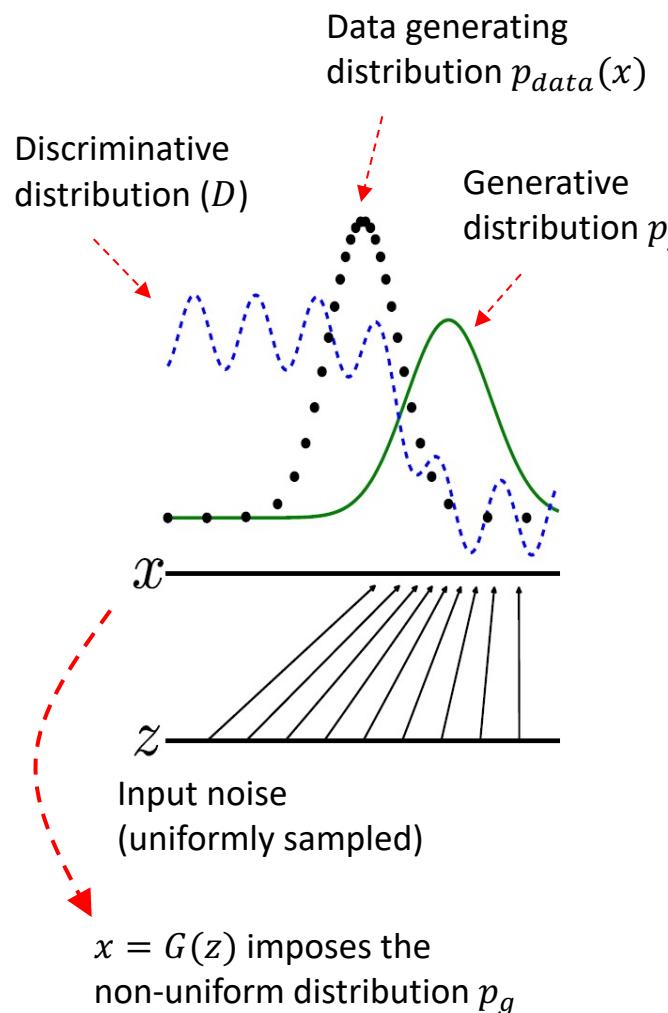
$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Aside: Jointly training two networks is challenging, can be unstable. Choosing objectives with better loss landscapes helps training, is an active area of research.

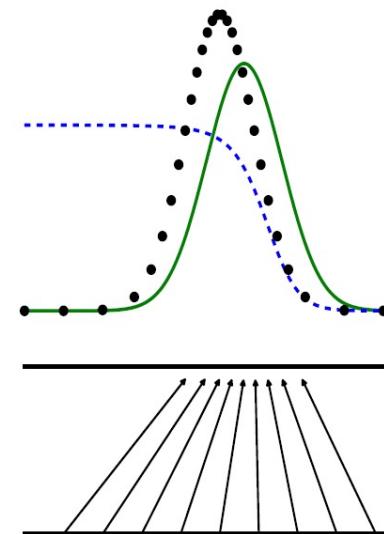
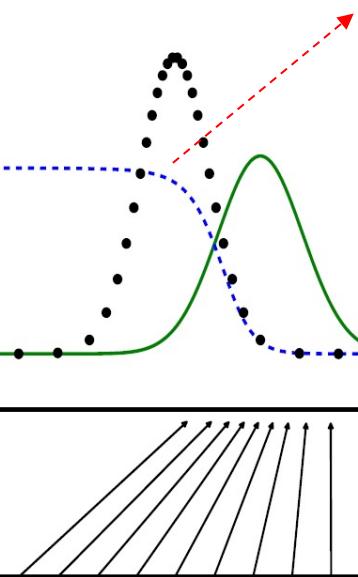
- Minimizing likelihood of discriminator being correct
-> Maximize likelihood of discriminator being wrong.
- Same objective of fooling discriminator, but now higher gradient signal for bad samples
-> Works much better! Standard in practice.



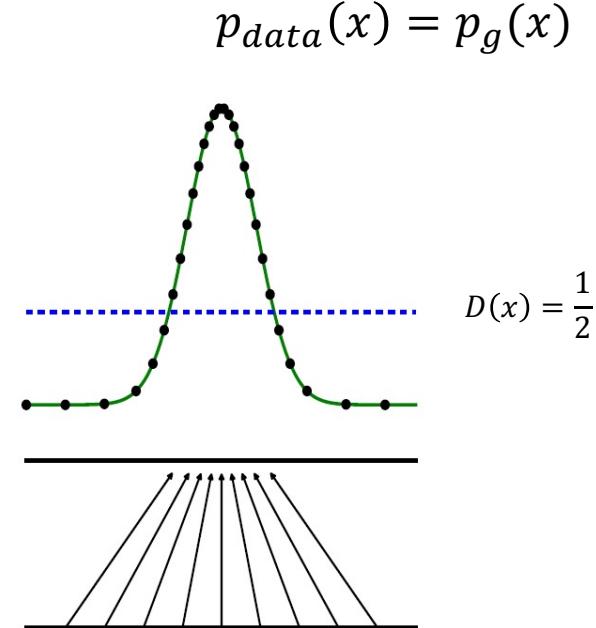
Training GANs: Two-player Game



The discriminator D is trained to discriminate samples from data, converging to $D^*(x) = \frac{p_{data}(x)}{p_{data}(x)+p_g(x)}$

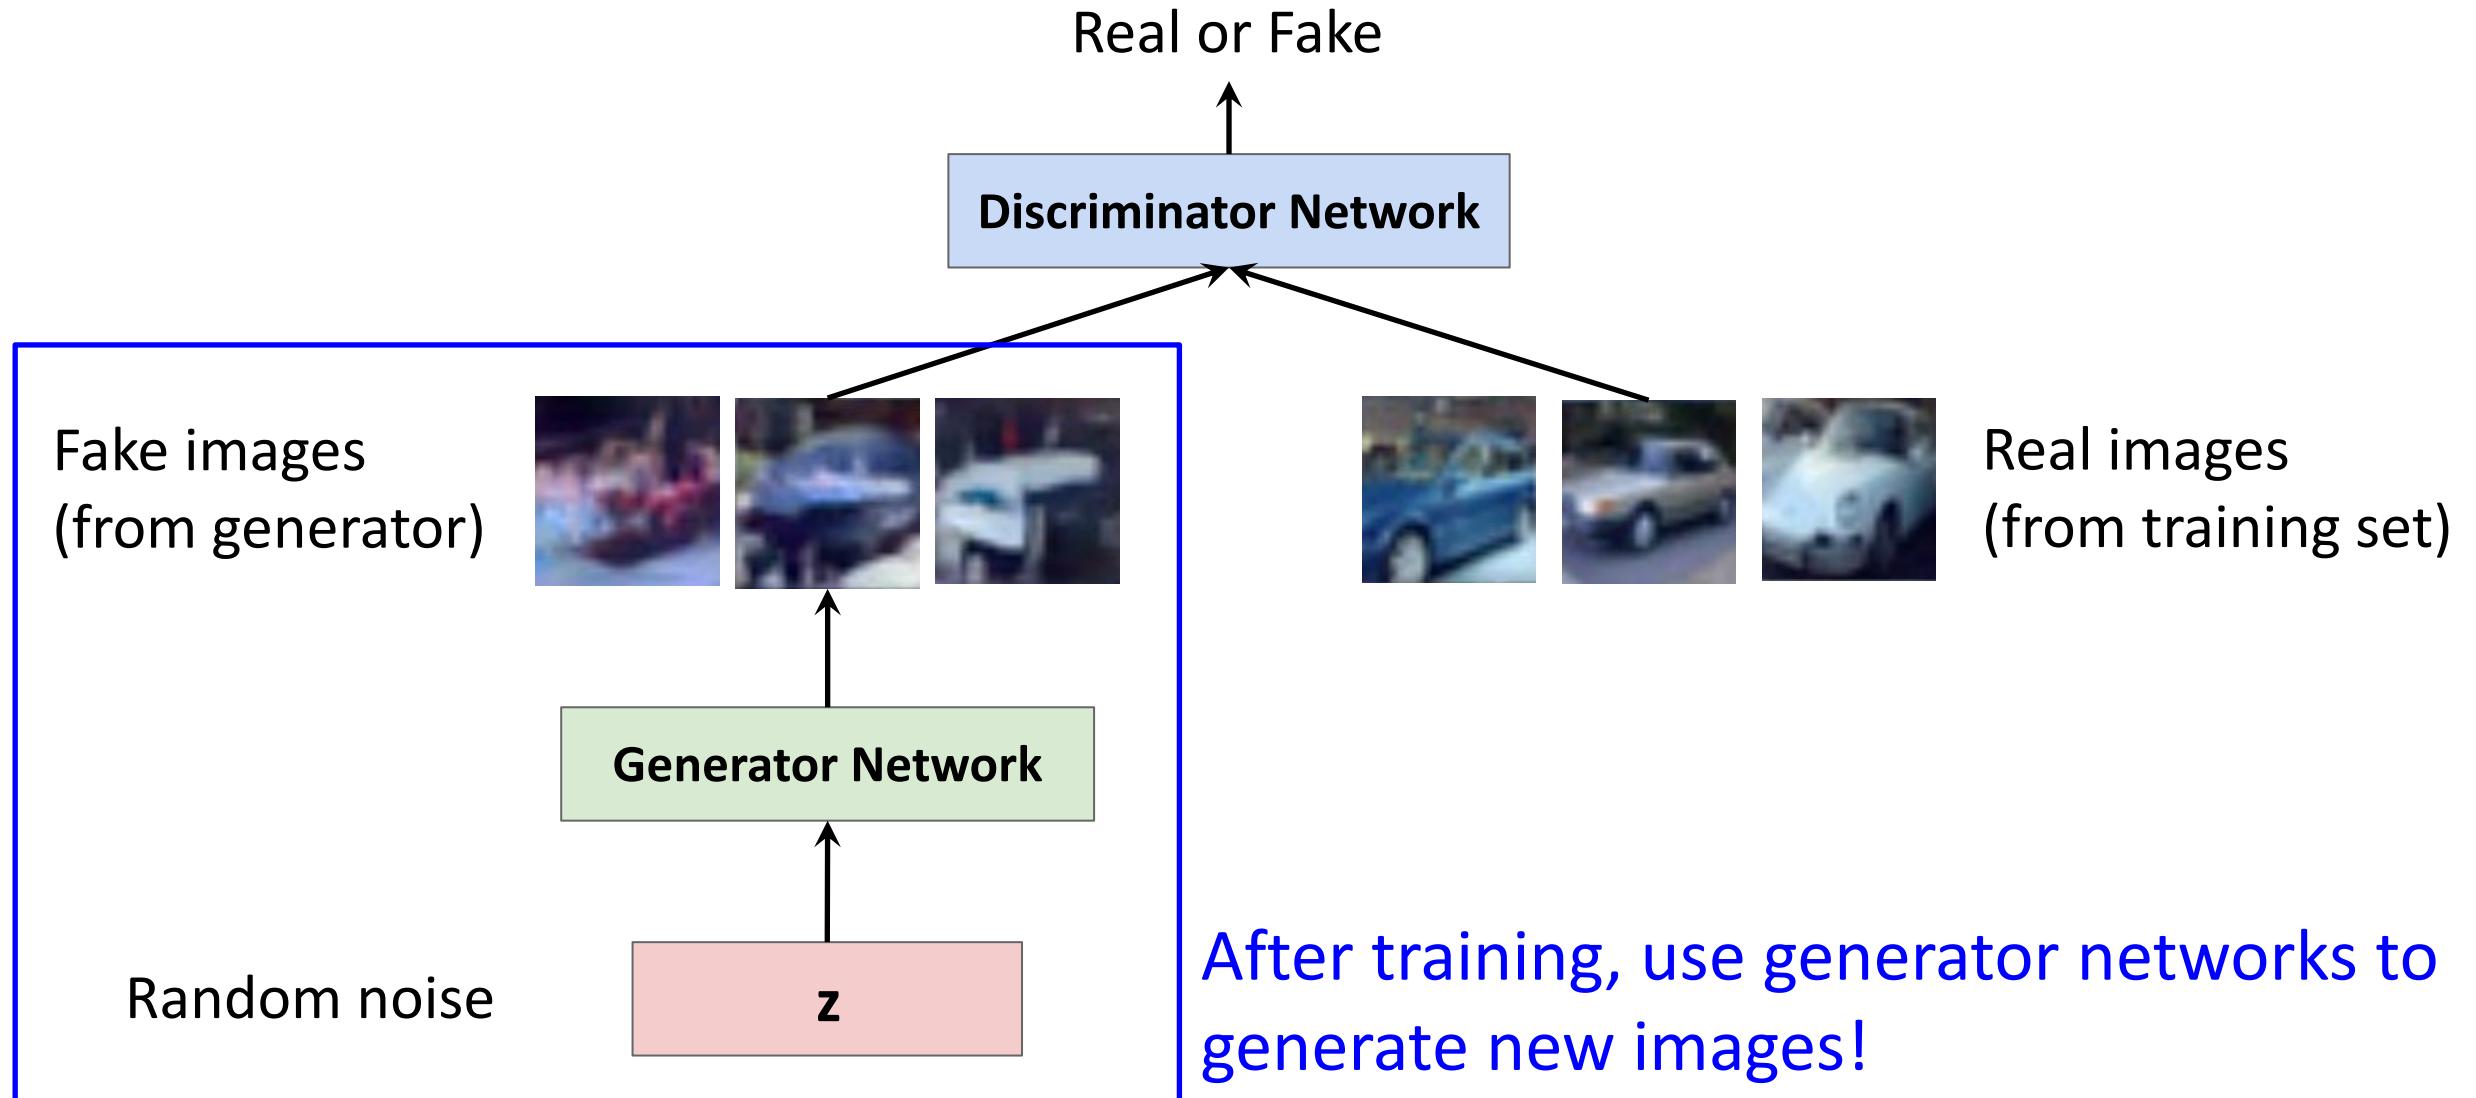


...



Gradient of D guides $G(z)$ to flow to regions that are more likely to be classified as data

Training GANs: Two-player Game



Advanced GANs Model

- DCGAN, Progressive GANs, StyleGAN etc.

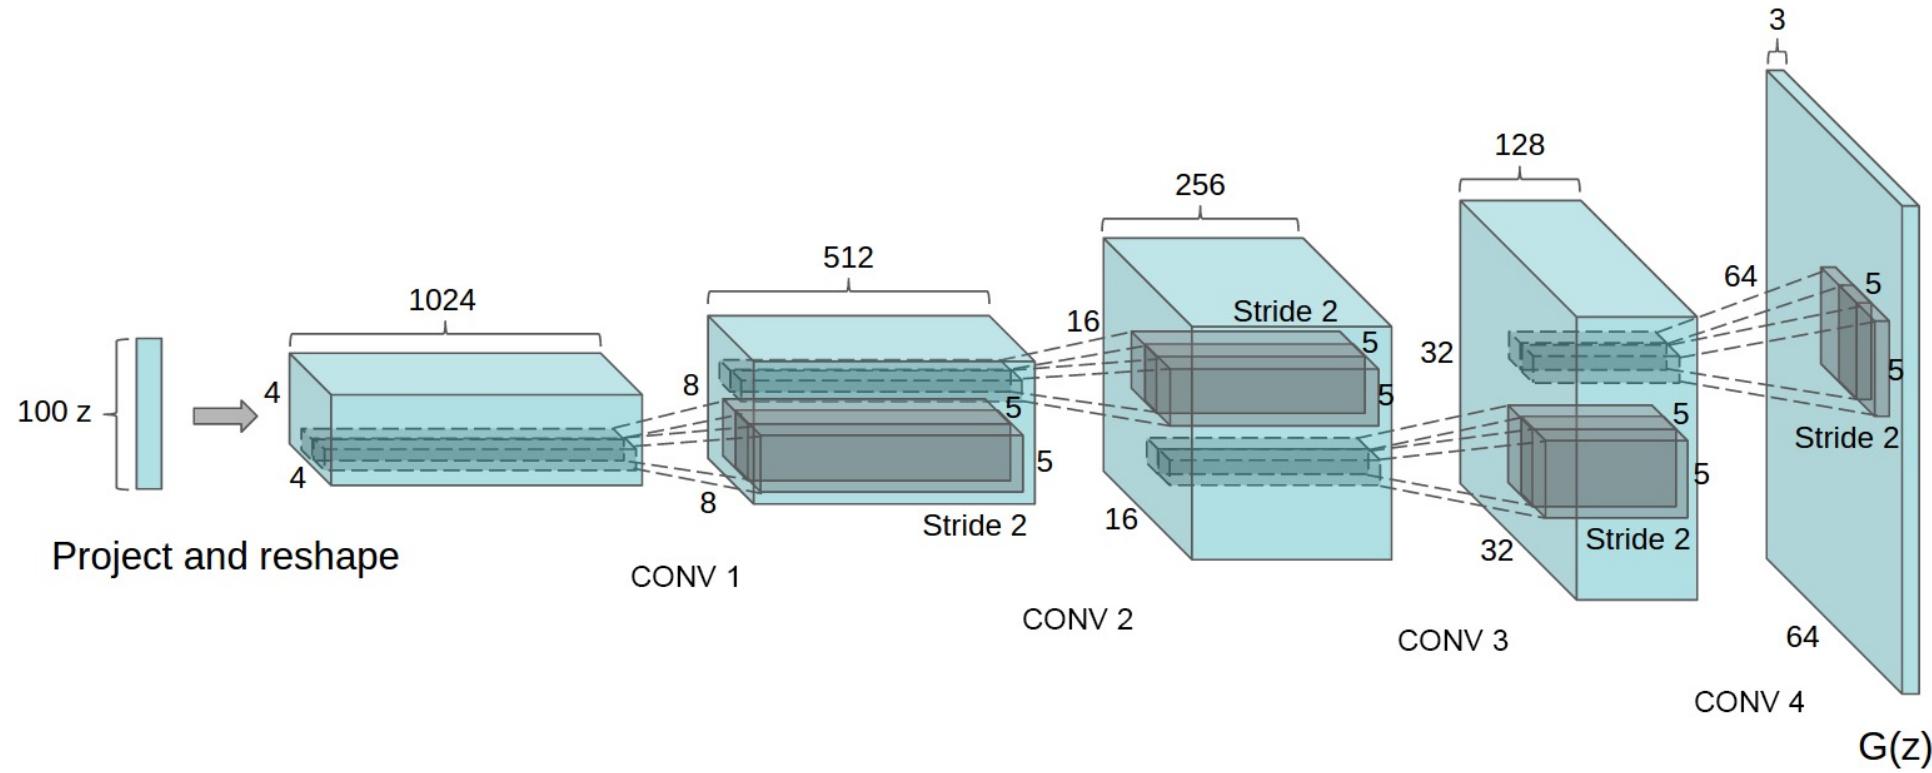
Deep Convolutional GANs (DCGAN)

- Generator is an upsampling network with fractionally-strided convolutions.
- Discriminator is a convolutional network.

Architecture guidelines for stable Deep Convolutional GANs

- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.

Deep Convolutional GANs (DCGAN)



- A 100-D uniform distribution z is projected to a small spatial extent convolutional representation using four fractionally-strided convolutions.
- No fully connected or pooling layers are used.

Deep Convolutional GANs (DCGAN)

Samples from the model look amazing!



Deep Convolutional GANs (DCGAN)

Interpolating between random points in latent space



Deep Convolutional GANs (DCGAN)

Vector arithmetic on latent space z
for visual concepts

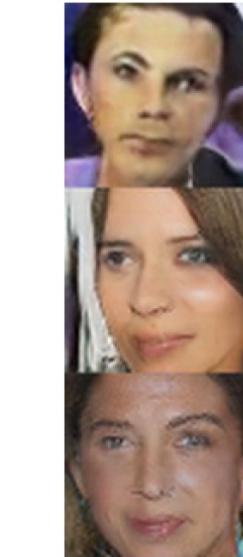
Samples from
the model



Average z
vectors, and
then do
arithmetic to
create Y



smiling
woman



neutral
woman



neutral
man



smiling man

The center sample is produced by feeding Y to the generator.

To demonstrate the interpolation capabilities of the generator, uniform noise was added to Y to produce the 8 other samples.

Deep Convolutional GANs (DCGAN)

Vector arithmetic on latent space z
for visual concepts

Samples from
the model



Average z
vectors, and
then do
arithmetic to
create Y



man
with glasses



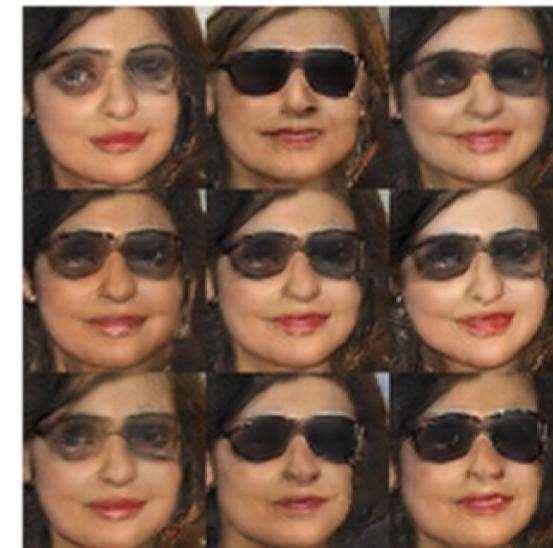
man
without glasses



woman
without glasses

The center sample is produced by feeding Y to the generator.

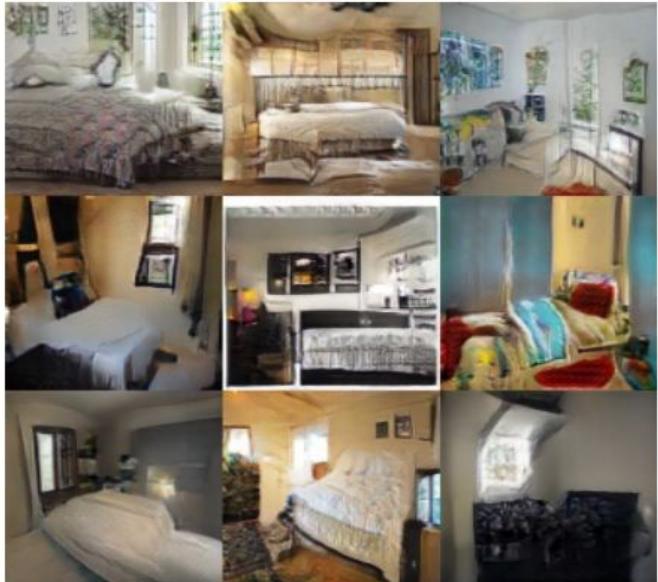
To demonstrate the interpolation capabilities of the generator, uniform noise was added to Y to produce the 8 other samples.



woman with glasses

2017: Explosion of GANs

Better Training and Generation



LSGAN, Zhu, 2017



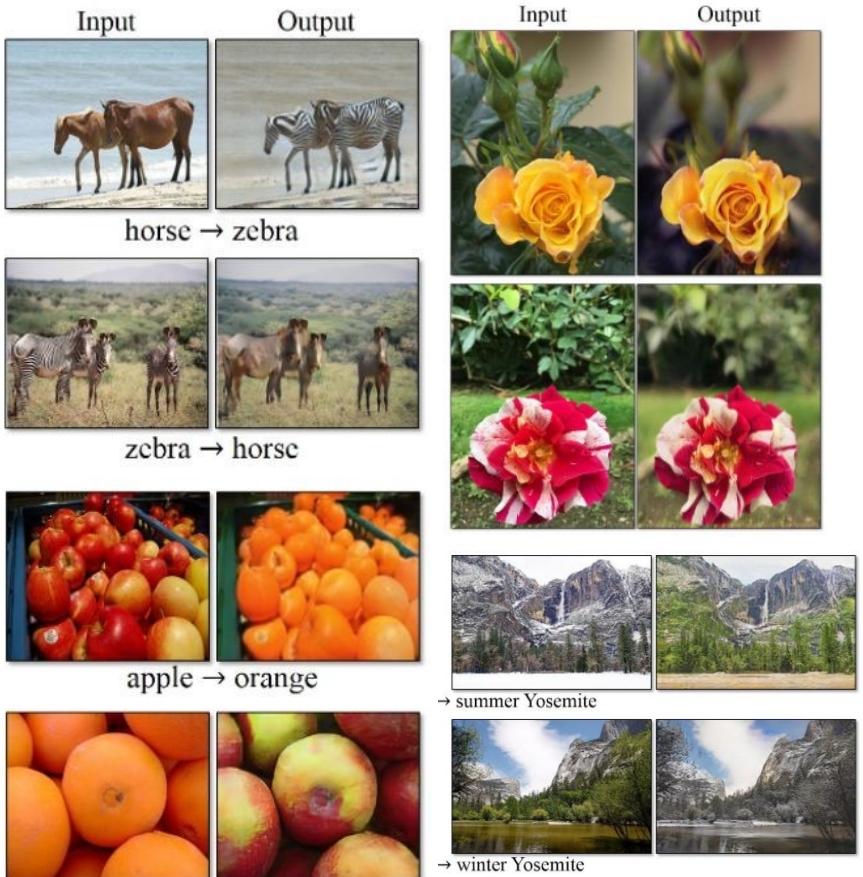
Wasserstein GAN,
Arjovsky 2017,
Improved Wasserstein GAN,
Gulrajani, 2017



Progressive GAN,
Karra, 2018

2017: Explosion of GANs

Source to Target Domain Transfer



CycleGAN, Zhu et al., 2017

this small bird has a pink breast and crown, and black primaries and secondaries.



this magnificent fellow is almost all black with a red crest, and white cheek patch.



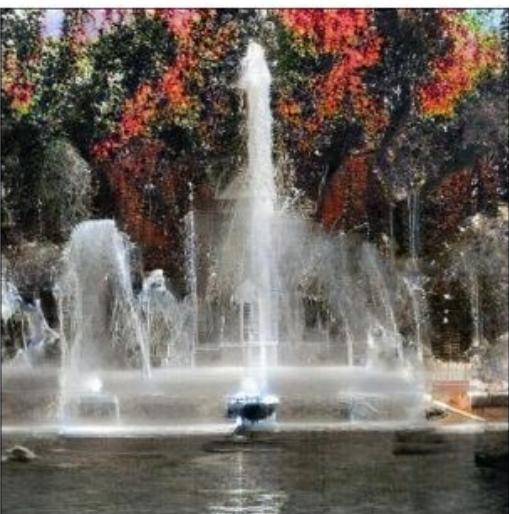
Reed et al., 2017



Pix2Pix, Isola, 2017

[slide courtesy: Stanford, CS231]

2019: BigGAN



Brock et al., 2019

Other Loss Functions

Name	Paper Link	Value Function
GAN	Arxiv	$L_D^{GAN} = E[\log(D(x))] + E[\log(1 - D(G(z)))]$ $L_G^{GAN} = E[\log(D(G(z)))]$
LSGAN	Arxiv	$L_D^{LSGAN} = E[(D(x) - 1)^2] + E[D(G(z))^2]$ $L_G^{LSGAN} = E[(D(G(z)) - 1)^2]$
WGAN	Arxiv	$L_D^{WGAN} = E[D(x)] - E[D(G(z))]$ $L_G^{WGAN} = E[D(G(z))]$ $W_D \leftarrow clip_by_value(W_D, -0.01, 0.01)$
WGAN_GP	Arxiv	$L_D^{WGAN_GP} = L_D^{WGAN} + \lambda E[V D(\alpha x - (1 - \alpha G(z))) - 1]^2$ $L_G^{WGAN_GP} = L_G^{WGAN}$
DRAGAN	Arxiv	$L_D^{DRAGAN} = L_D^{GAN} + \lambda E[V D(\alpha x - (1 - \alpha x_p)) - 1]^2$ $L_G^{DRAGAN} = L_G^{GAN}$
CGAN	Arxiv	$L_D^{CGAN} = E[\log(D(x, c))] + E[\log(1 - D(G(z), c))]$ $L_G^{CGAN} = E[\log(D(G(z), c))]$
infoGAN	Arxiv	$L_{D,Q}^{InfoGAN} = L_D^{GAN} - \lambda L_I(c, c')$ $L_G^{InfoGAN} = L_G^{GAN} - \lambda L_I(c, c')$
ACGAN	Arxiv	$L_{D,Q}^{ACGAN} = L_D^{GAN} + E[P(class = c x)] + E[P(class = c G(z))]$ $L_G^{ACGAN} = L_G^{GAN} + E[P(class = c G(z))]$
EBGAN	Arxiv	$L_D^{EBGAN} = D_{AE}(x) + \max(0, m - D_{AE}(G(z)))$ $L_G^{EBGAN} = D_{AE}(G(z)) + \lambda \cdot PT$
BEGAN	Arxiv	$L_D^{BEGAN} = D_{AE}(x) - k_t D_{AE}(G(z))$ $L_G^{BEGAN} = D_{AE}(G(z))$ $k_{t+1} = k_t + \lambda(\gamma D_{AE}(x) - D_{AE}(G(z)))$

Method	Local convergence (a.c. case)	Local convergence (general case)
unregularized (Goodfellow et al., 2014)	✓	✗
WGAN (Arjovsky et al., 2017)	✗	✗
WGAN-GP (Gulrajani et al., 2017)	✗	✗
DRAGAN (Kodali et al., 2017)	✓	✗
Instance noise (Sønderby et al., 2016)	✓	✓
ConOpt (Mescheder et al., 2017)	✓	✓
Gradient penalties (Roth et al., 2017)	✓	✓
Gradient penalty on real data only	✓	✓
Gradient penalty on fake data only	✓	✓

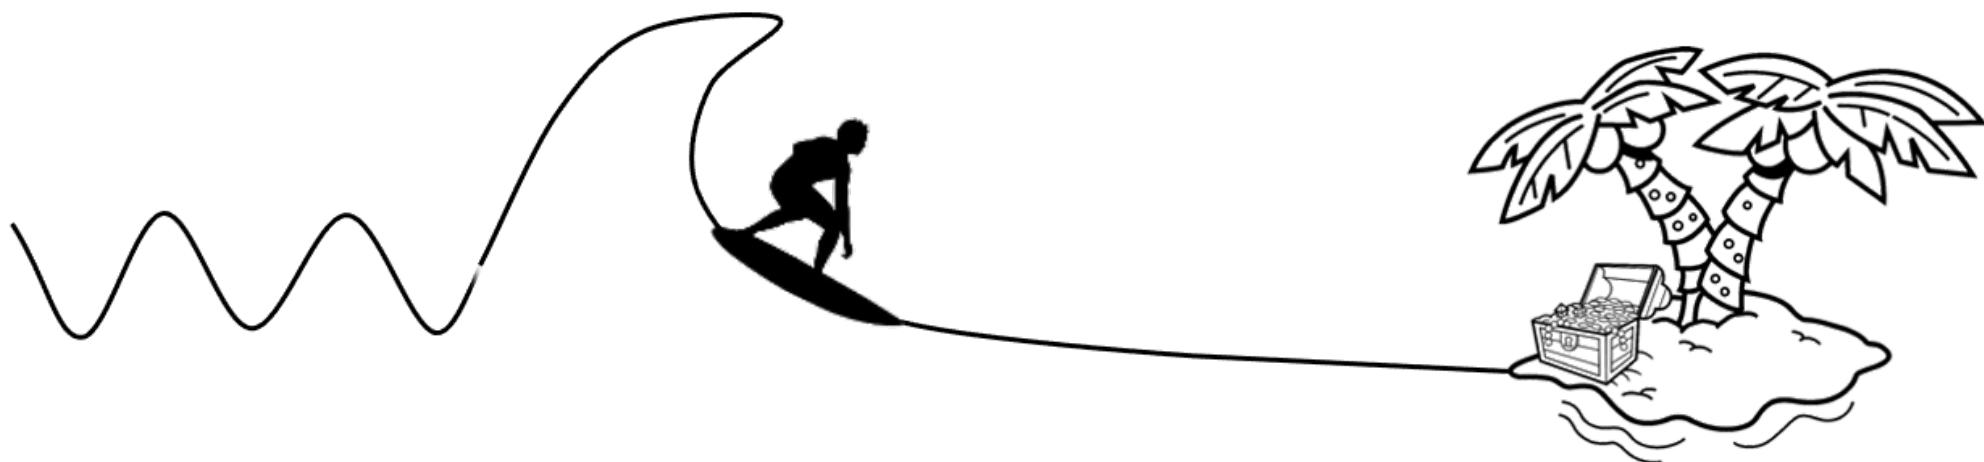
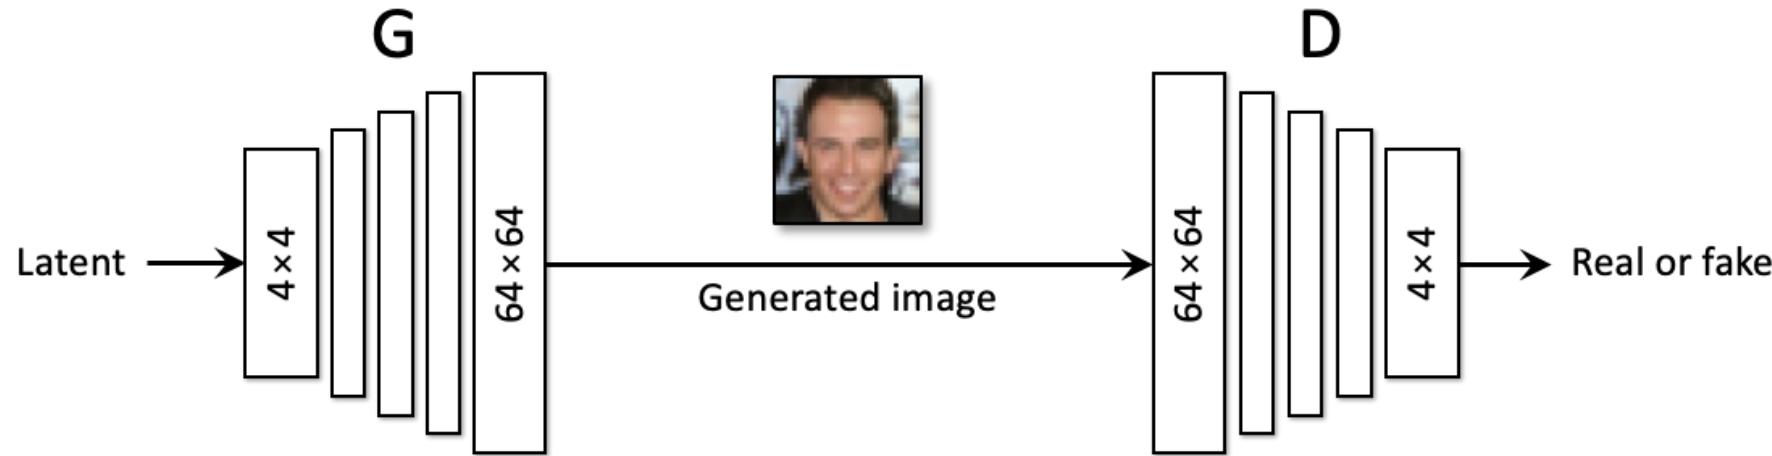
Table 1. Convergence properties of different GAN training algorithms for general GAN-architectures. Here, we distinguish between the case where both the data and generator distributions are absolutely continuous (a.c.) and the general case where they may lie on lower dimensional manifolds.

<https://github.com/hwalsuklee/tensorflow-generative-model-collections>

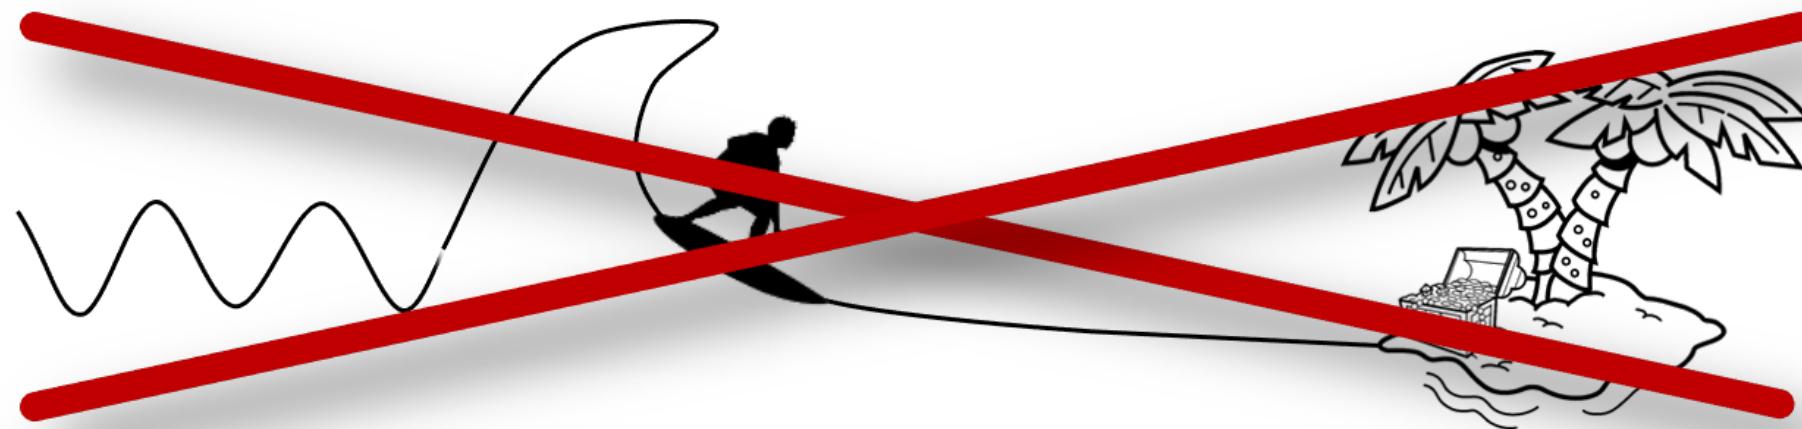
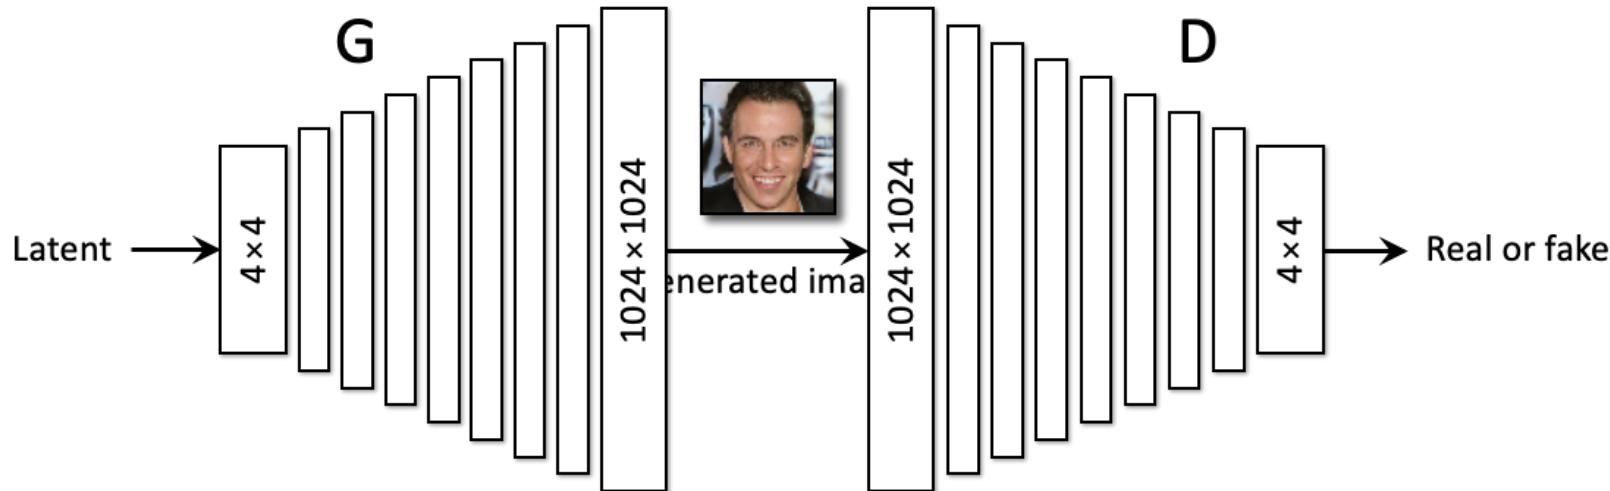
https://github.com/LMescheder/GAN_stability

Which Training Methods for GANs do actually Converge?

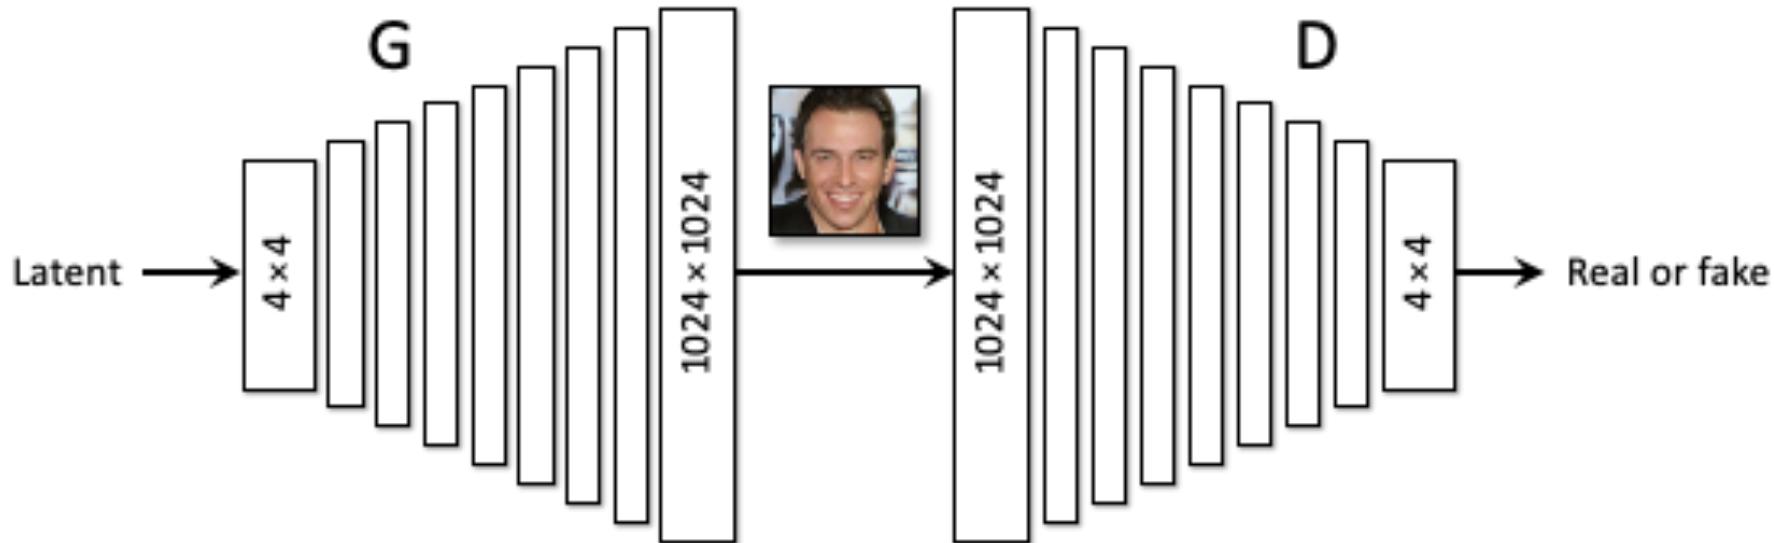
Progressive Generation



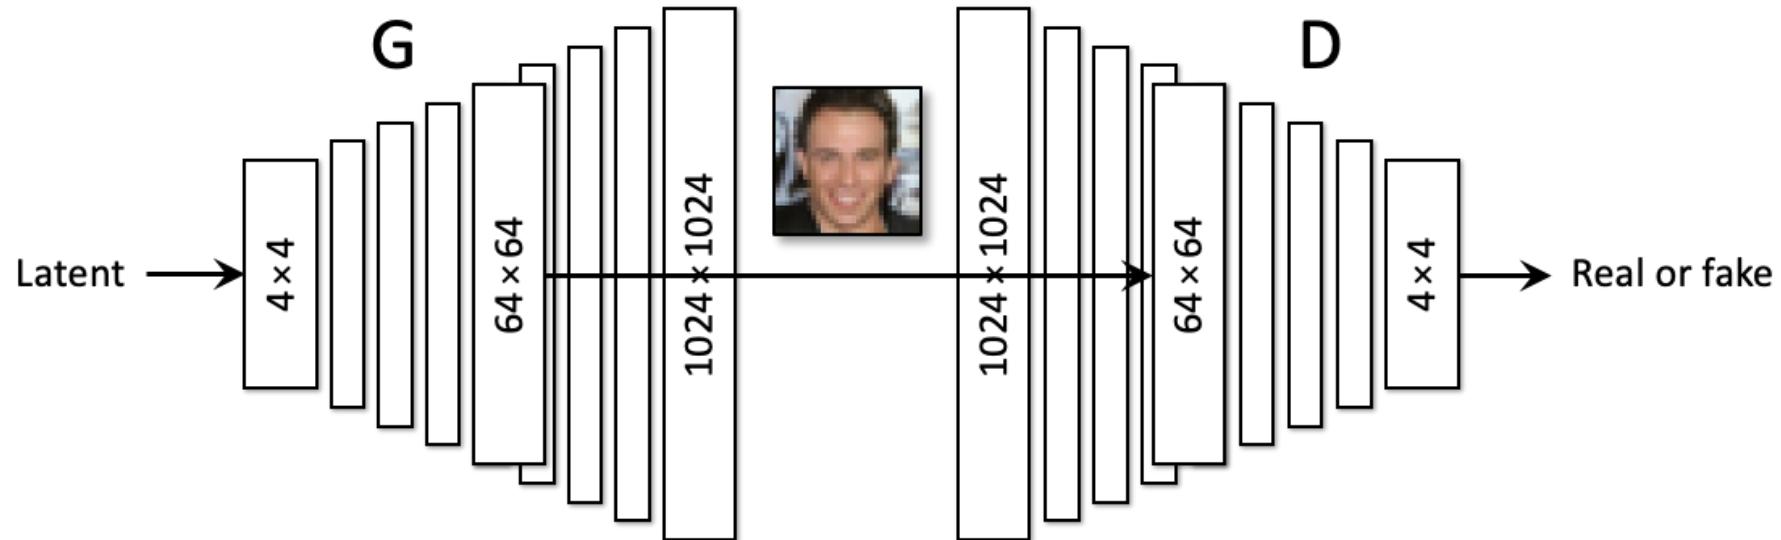
Progressive Generation



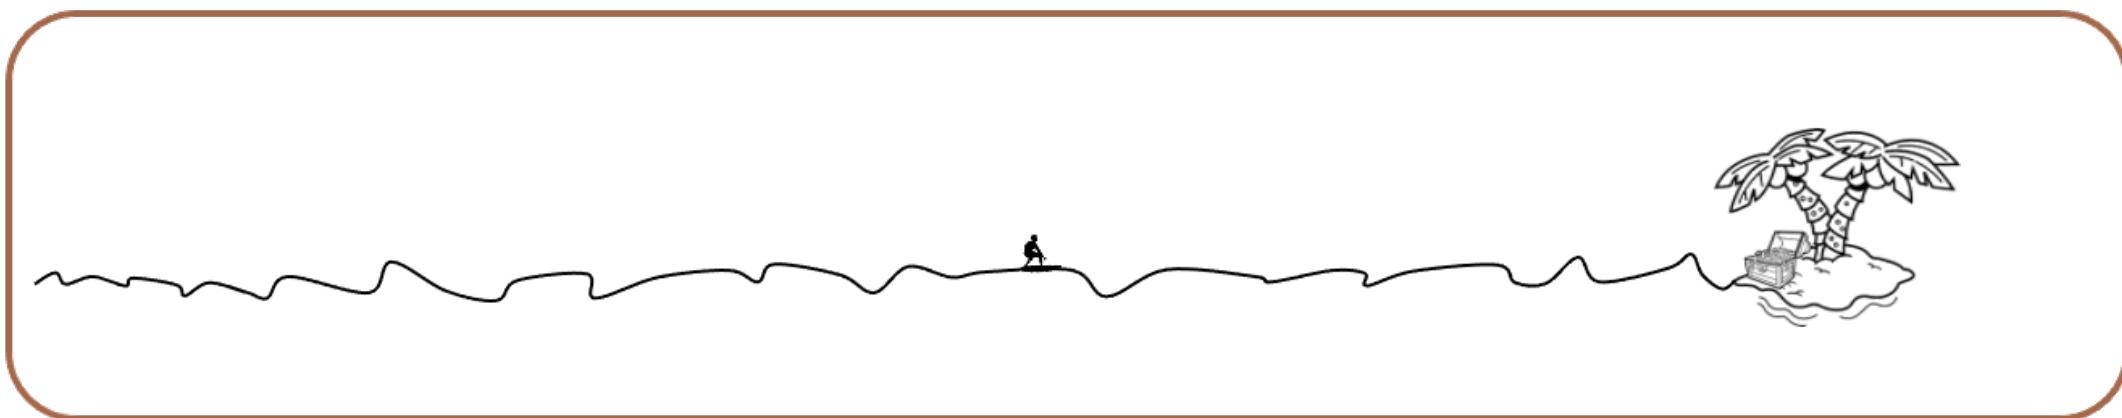
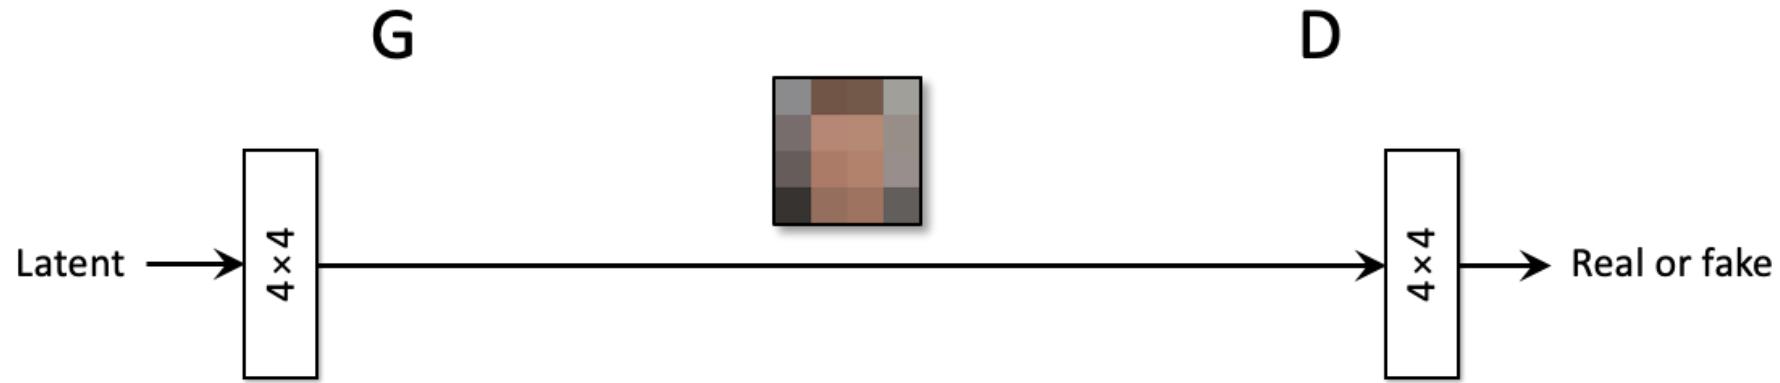
Progressive Generation



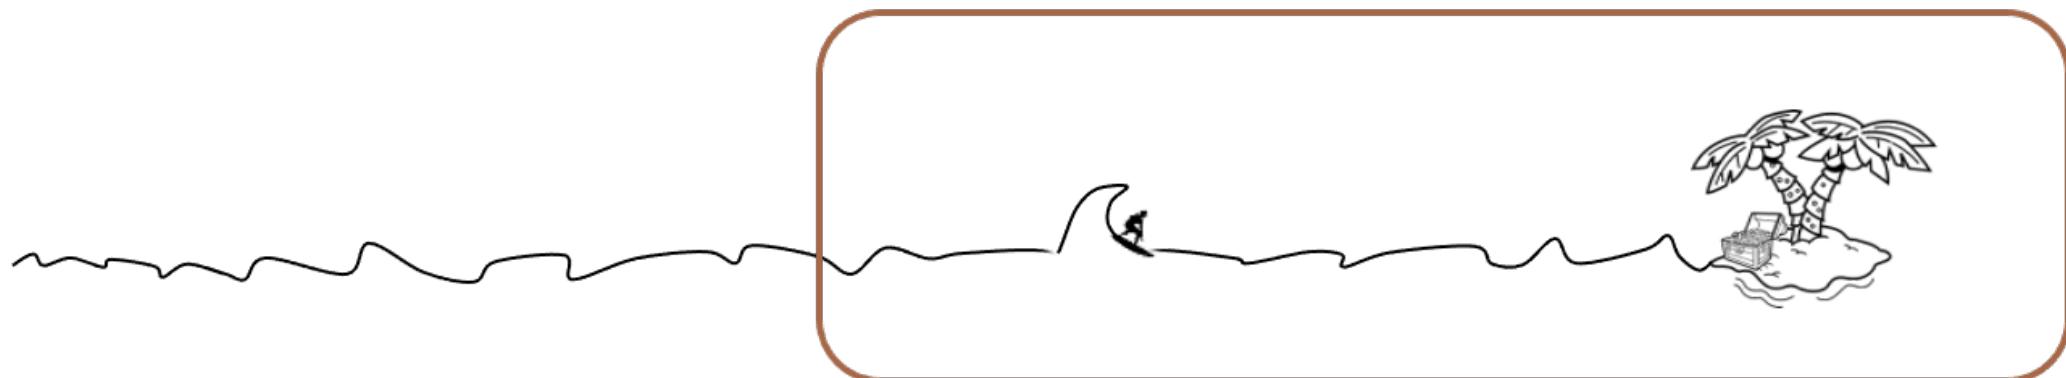
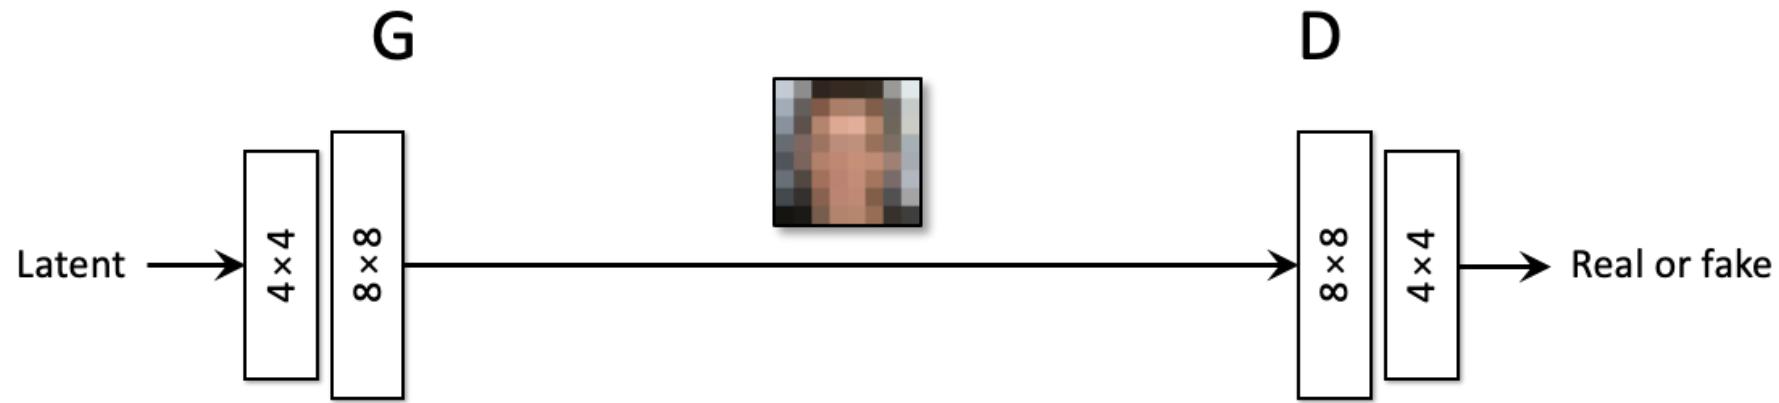
Progressive Generation



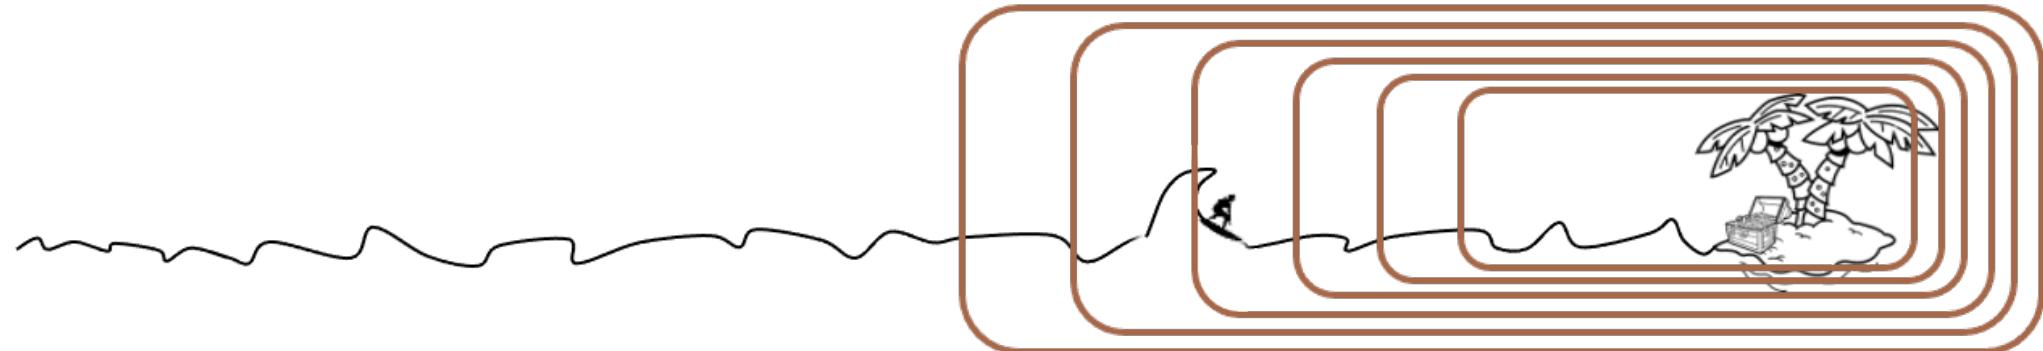
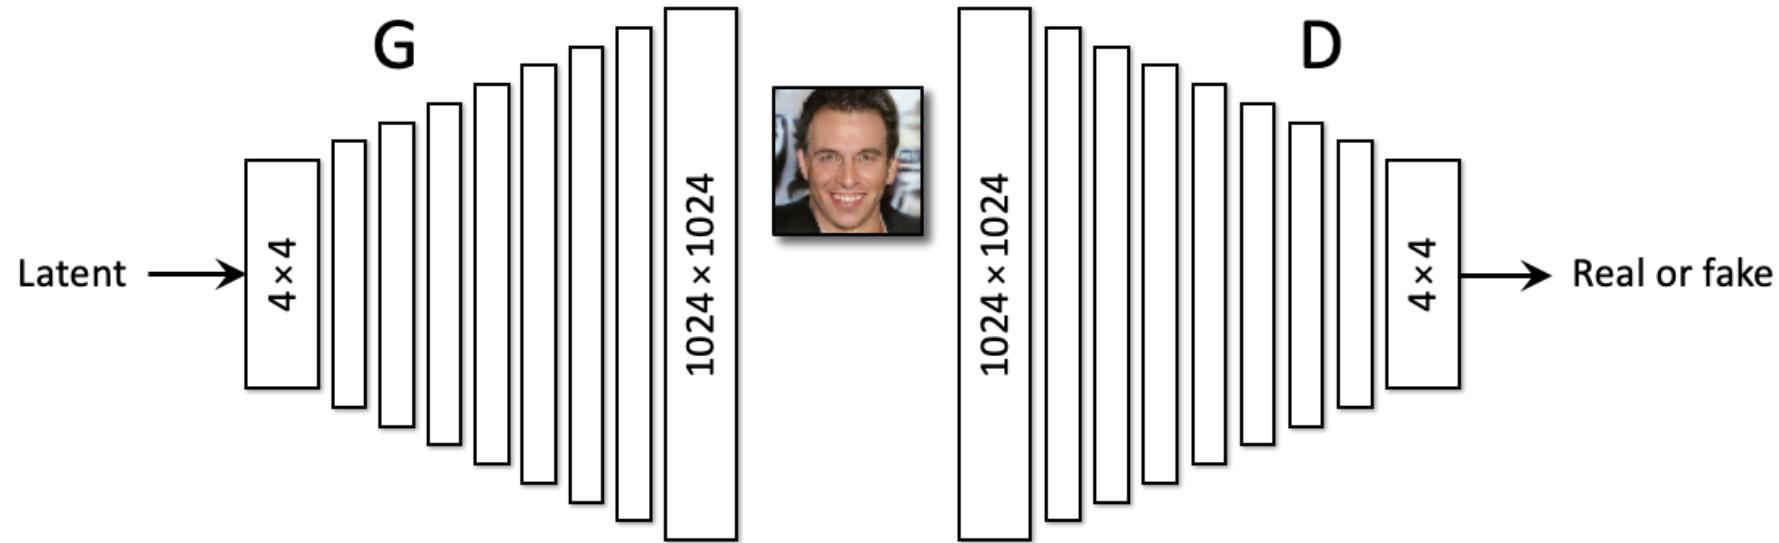
Progressive Generation



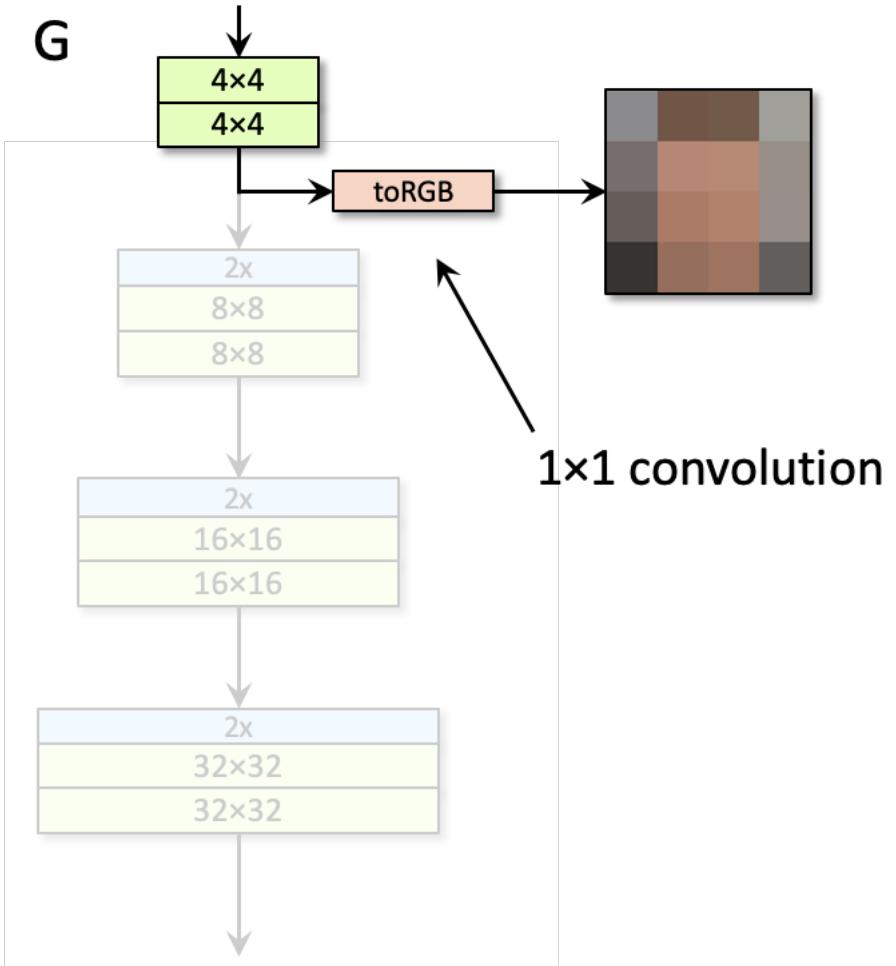
Progressive Generation



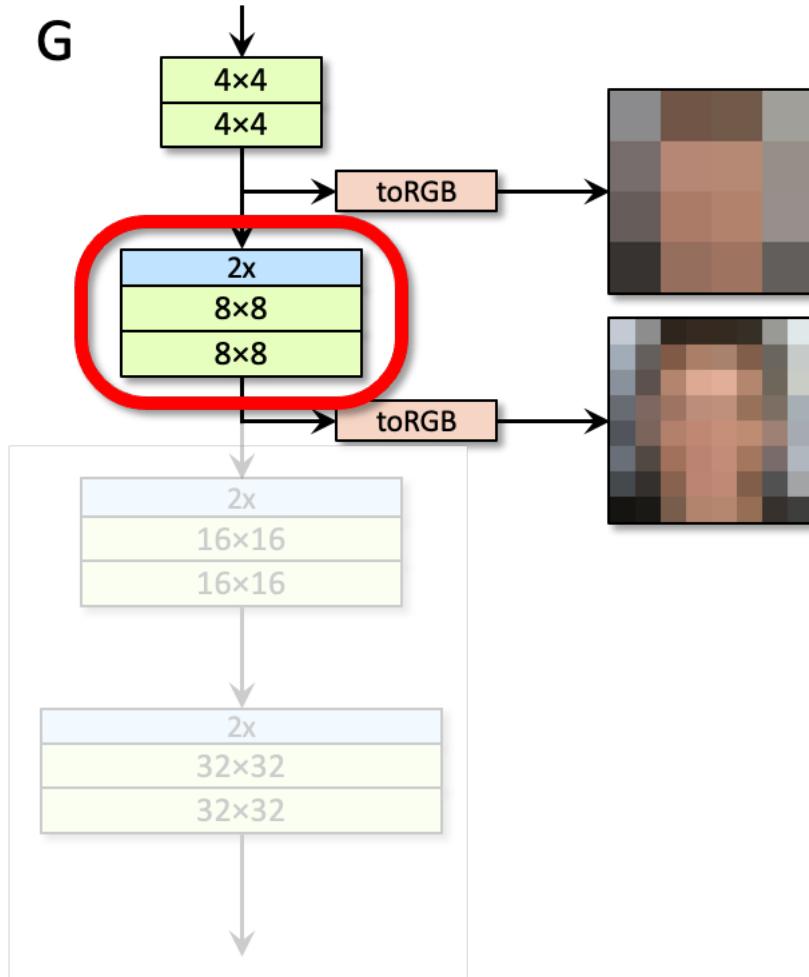
Progressive Generation



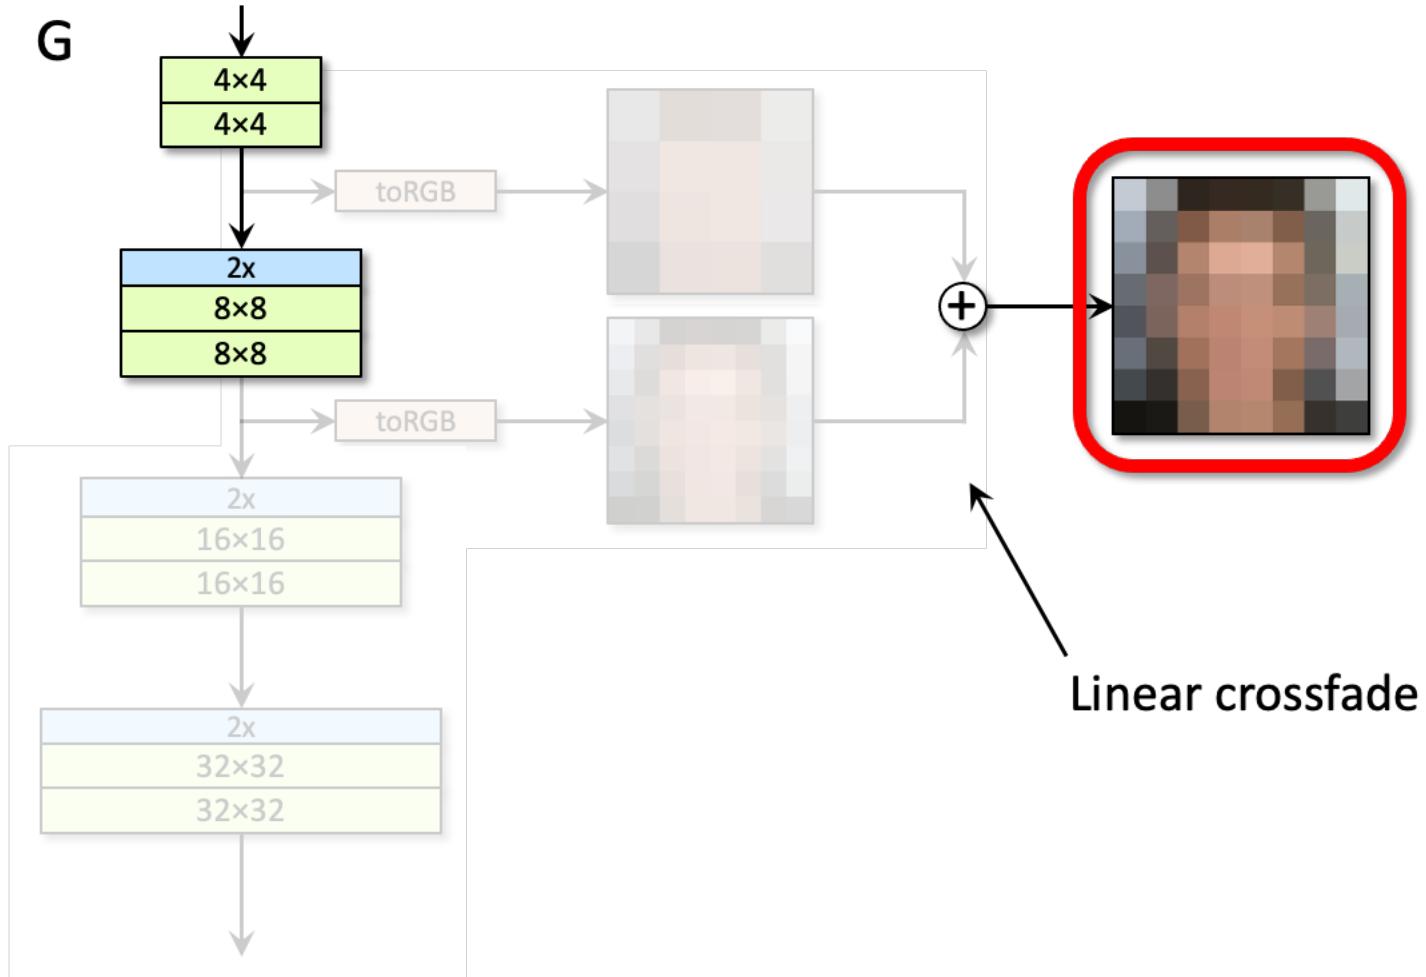
Progressive Generation



Progressive Generation



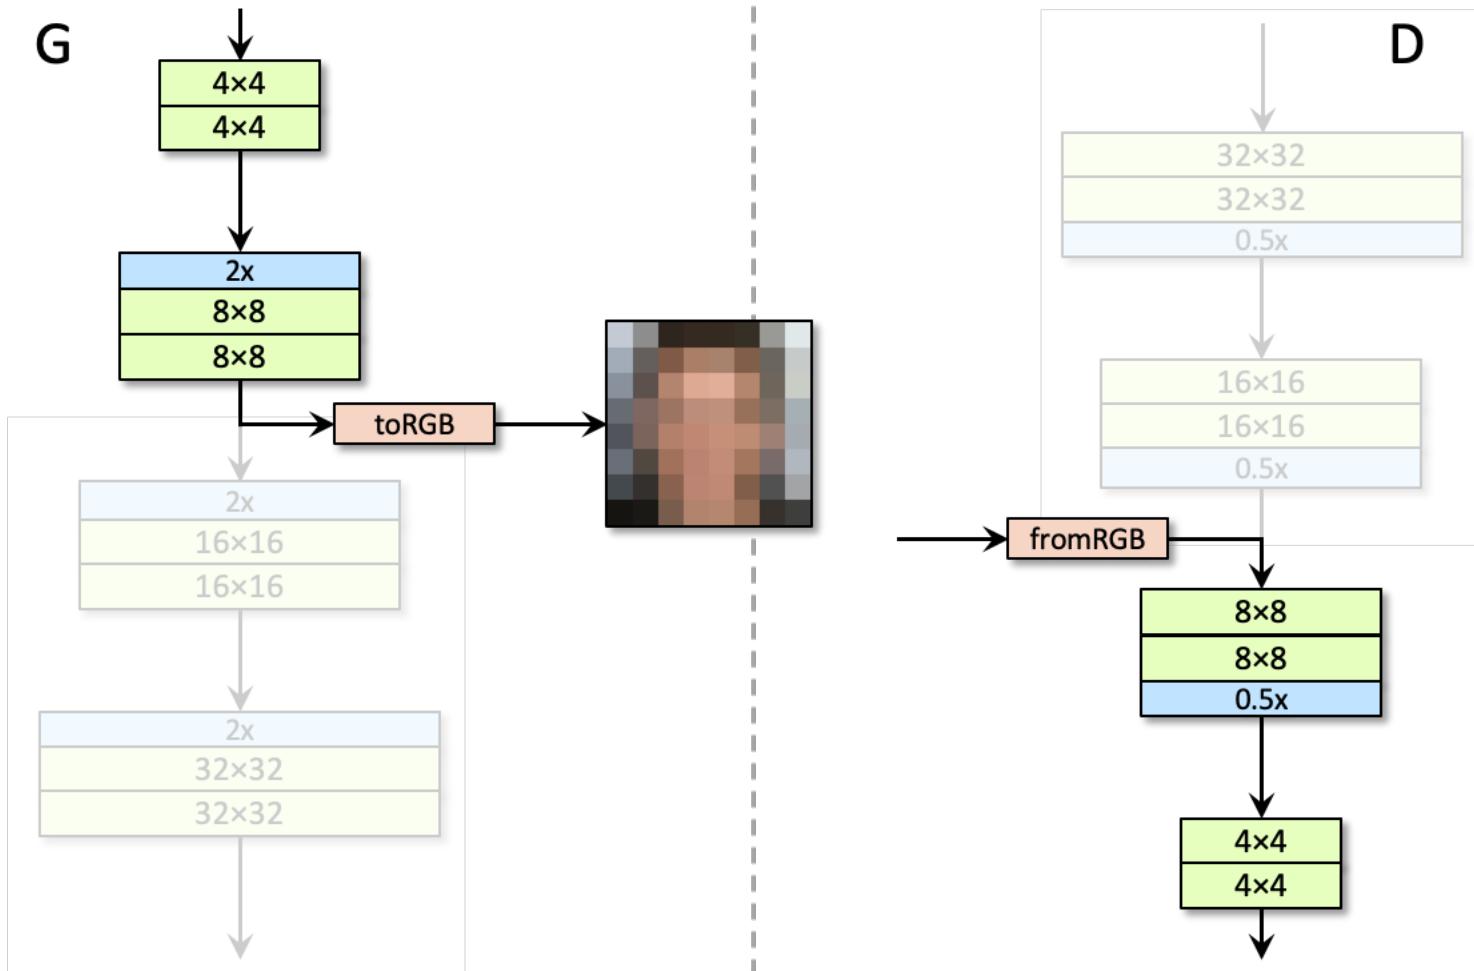
Progressive Generation



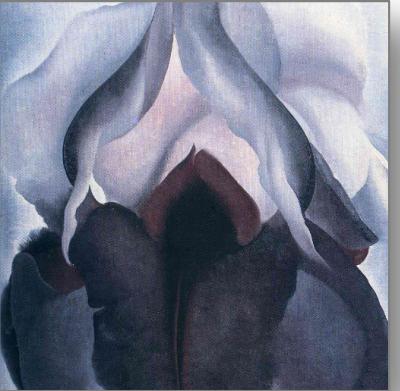
Linear crossfade

Progressive Generation

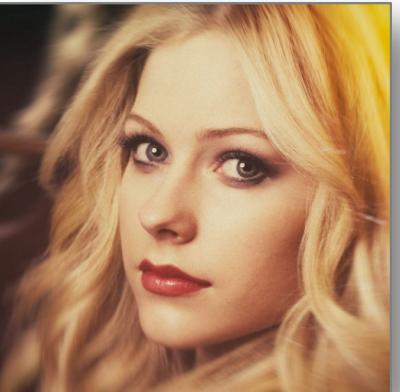
Progressive generation



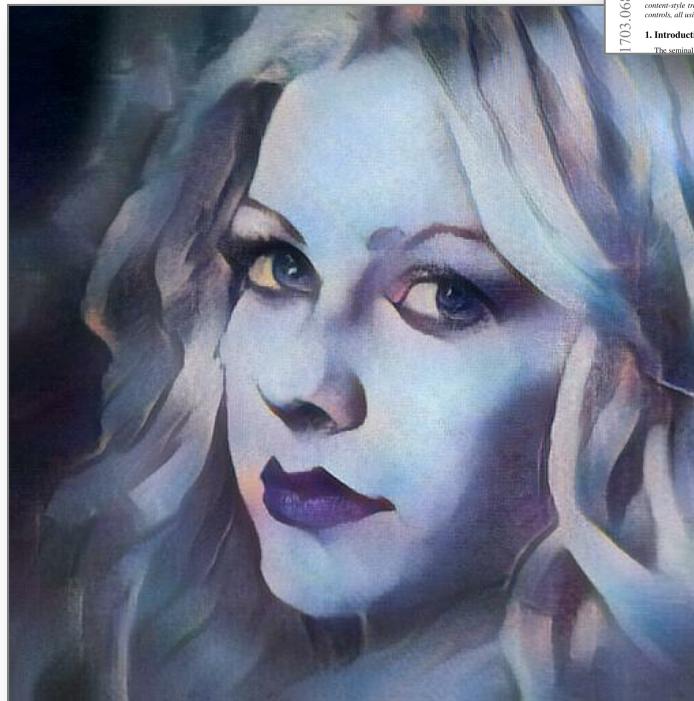
AdaIN in Style Transfer



Style



Content



Result

Arbitrary Style Transfer in Real-time with Adaptive Instance Normalization

Xun Huang Serge Belongie
Department of Computer Science & Cornell Tech, Cornell University
{xh256,sjb34}@cornell.edu

Abstract

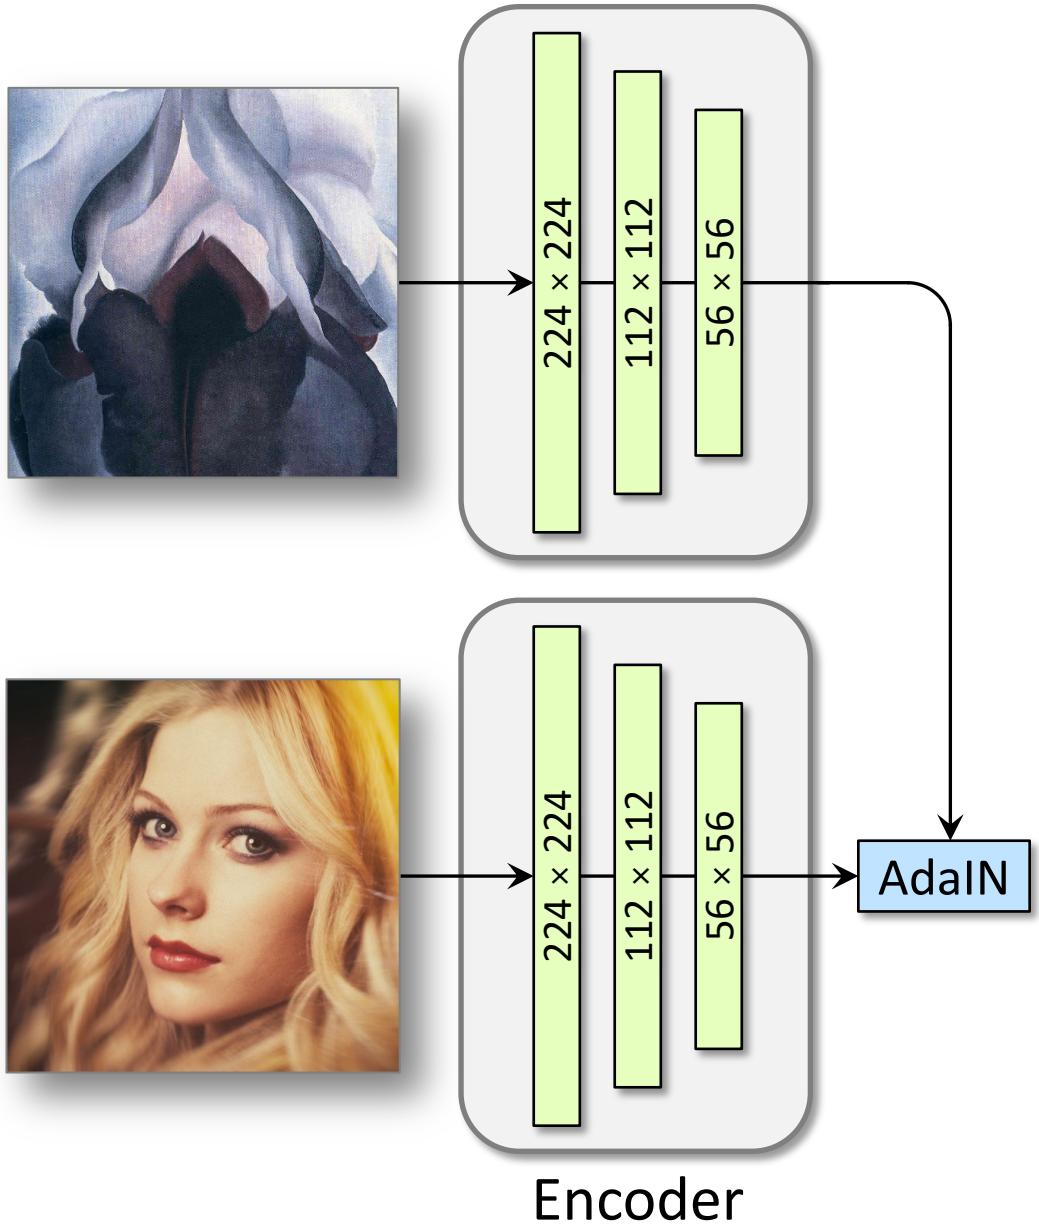
Gatys *et al.* recently introduced a neural algorithm that renders a content image in the style of another image, achieving so-called style transfer. However, their framework requires a slow iterative optimization process, which makes it impractical for real-time applications. Fast style transfer methods have been proposed, but they only work up neural style transfer. Unfortunately, the speed improvements come at the cost of the ability to handle a fixed set of styles and cannot adapt to arbitrary new styles. In this paper, we present a simple yet effective approach that for the first time achieves arbitrary style transfer in real-time. At the heart of our method is a novel adaptive instance normalization (AdaIN) layer that aligns the mean and variance of the content input with those of the style input. Through this layer, AdaIN can directly map the content input to the content of the former and the style later by transferring feature statistics. A decoder network is then learned to generate the final image. Our method is able to handle arbitrary styles and content-style trade-off, style interpolation, color & spatial controls, all using a single feed-forward neural network.

1. Introduction

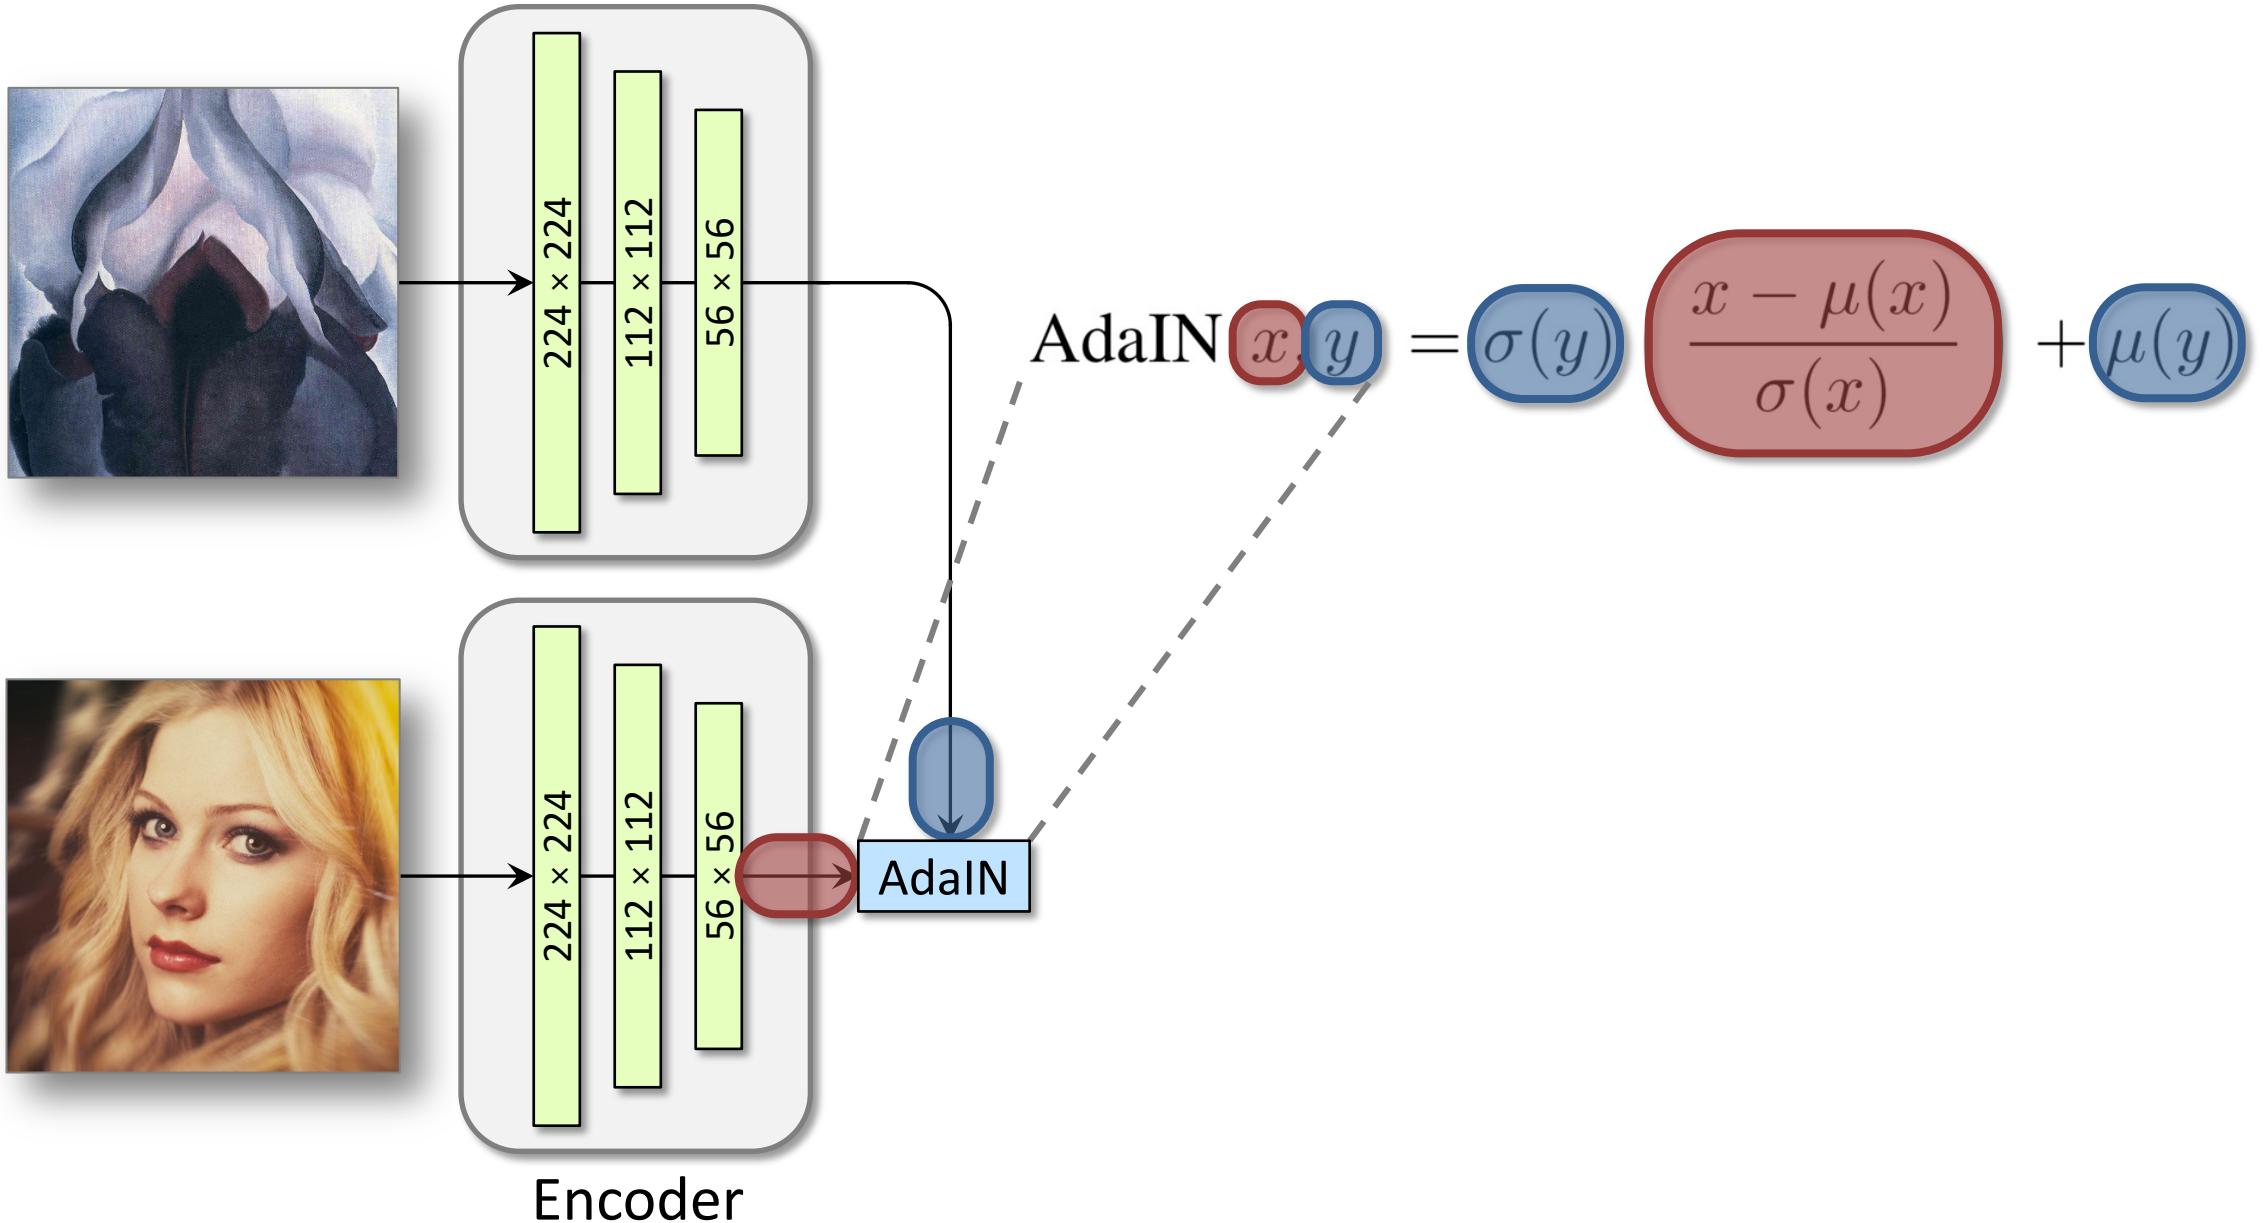
The seminal work of Gatys *et al.* [16] showed that deep

1703.06868v2 [CC-BY-NC-ND] 30 Jul 2017

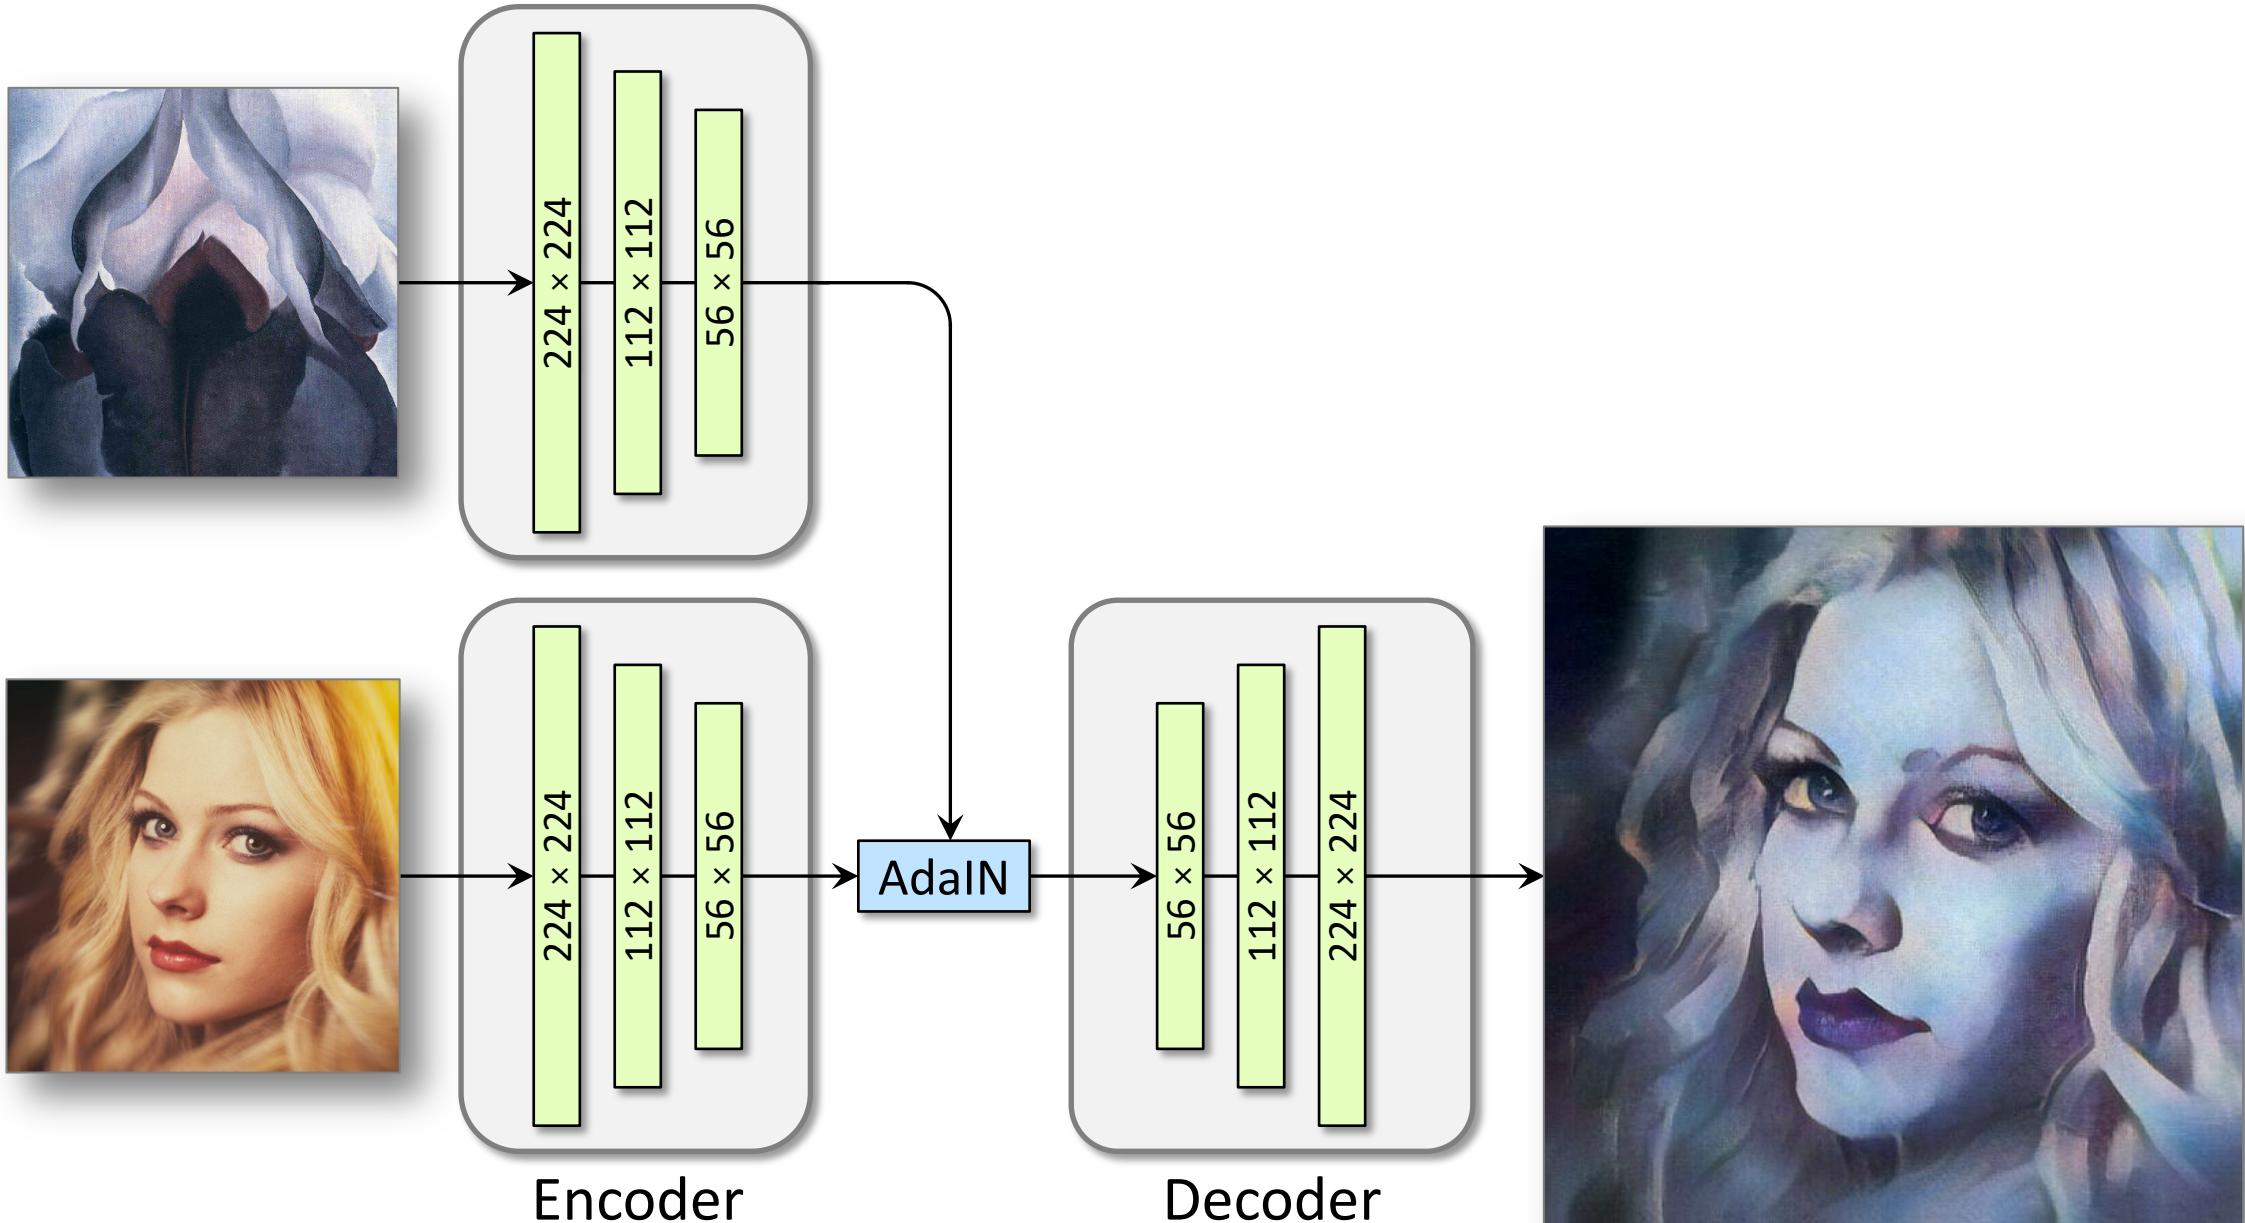
StyleGAN



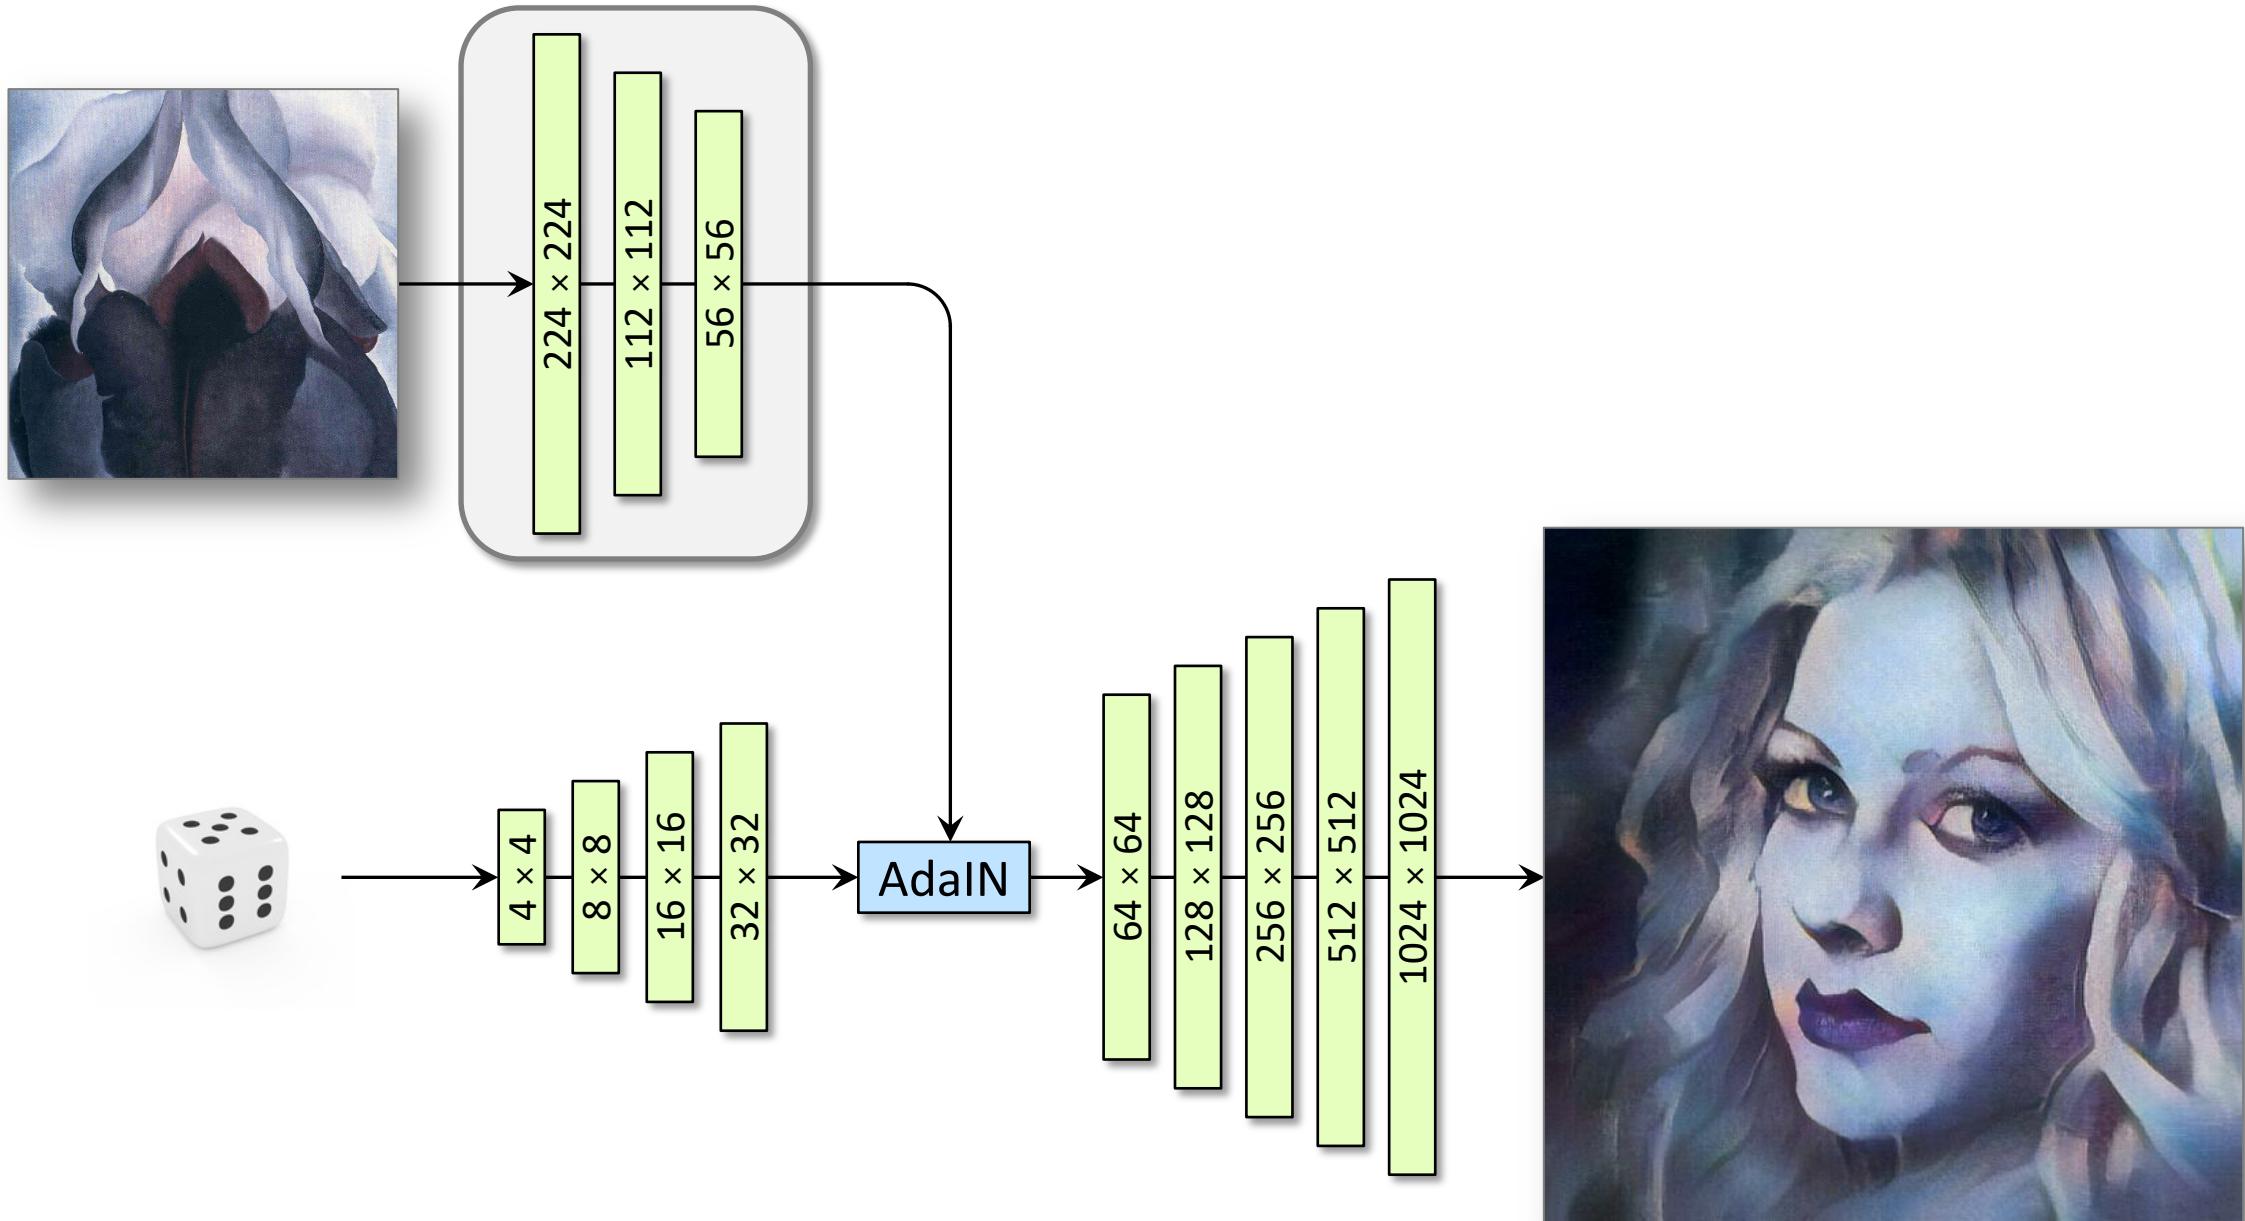
StyleGAN



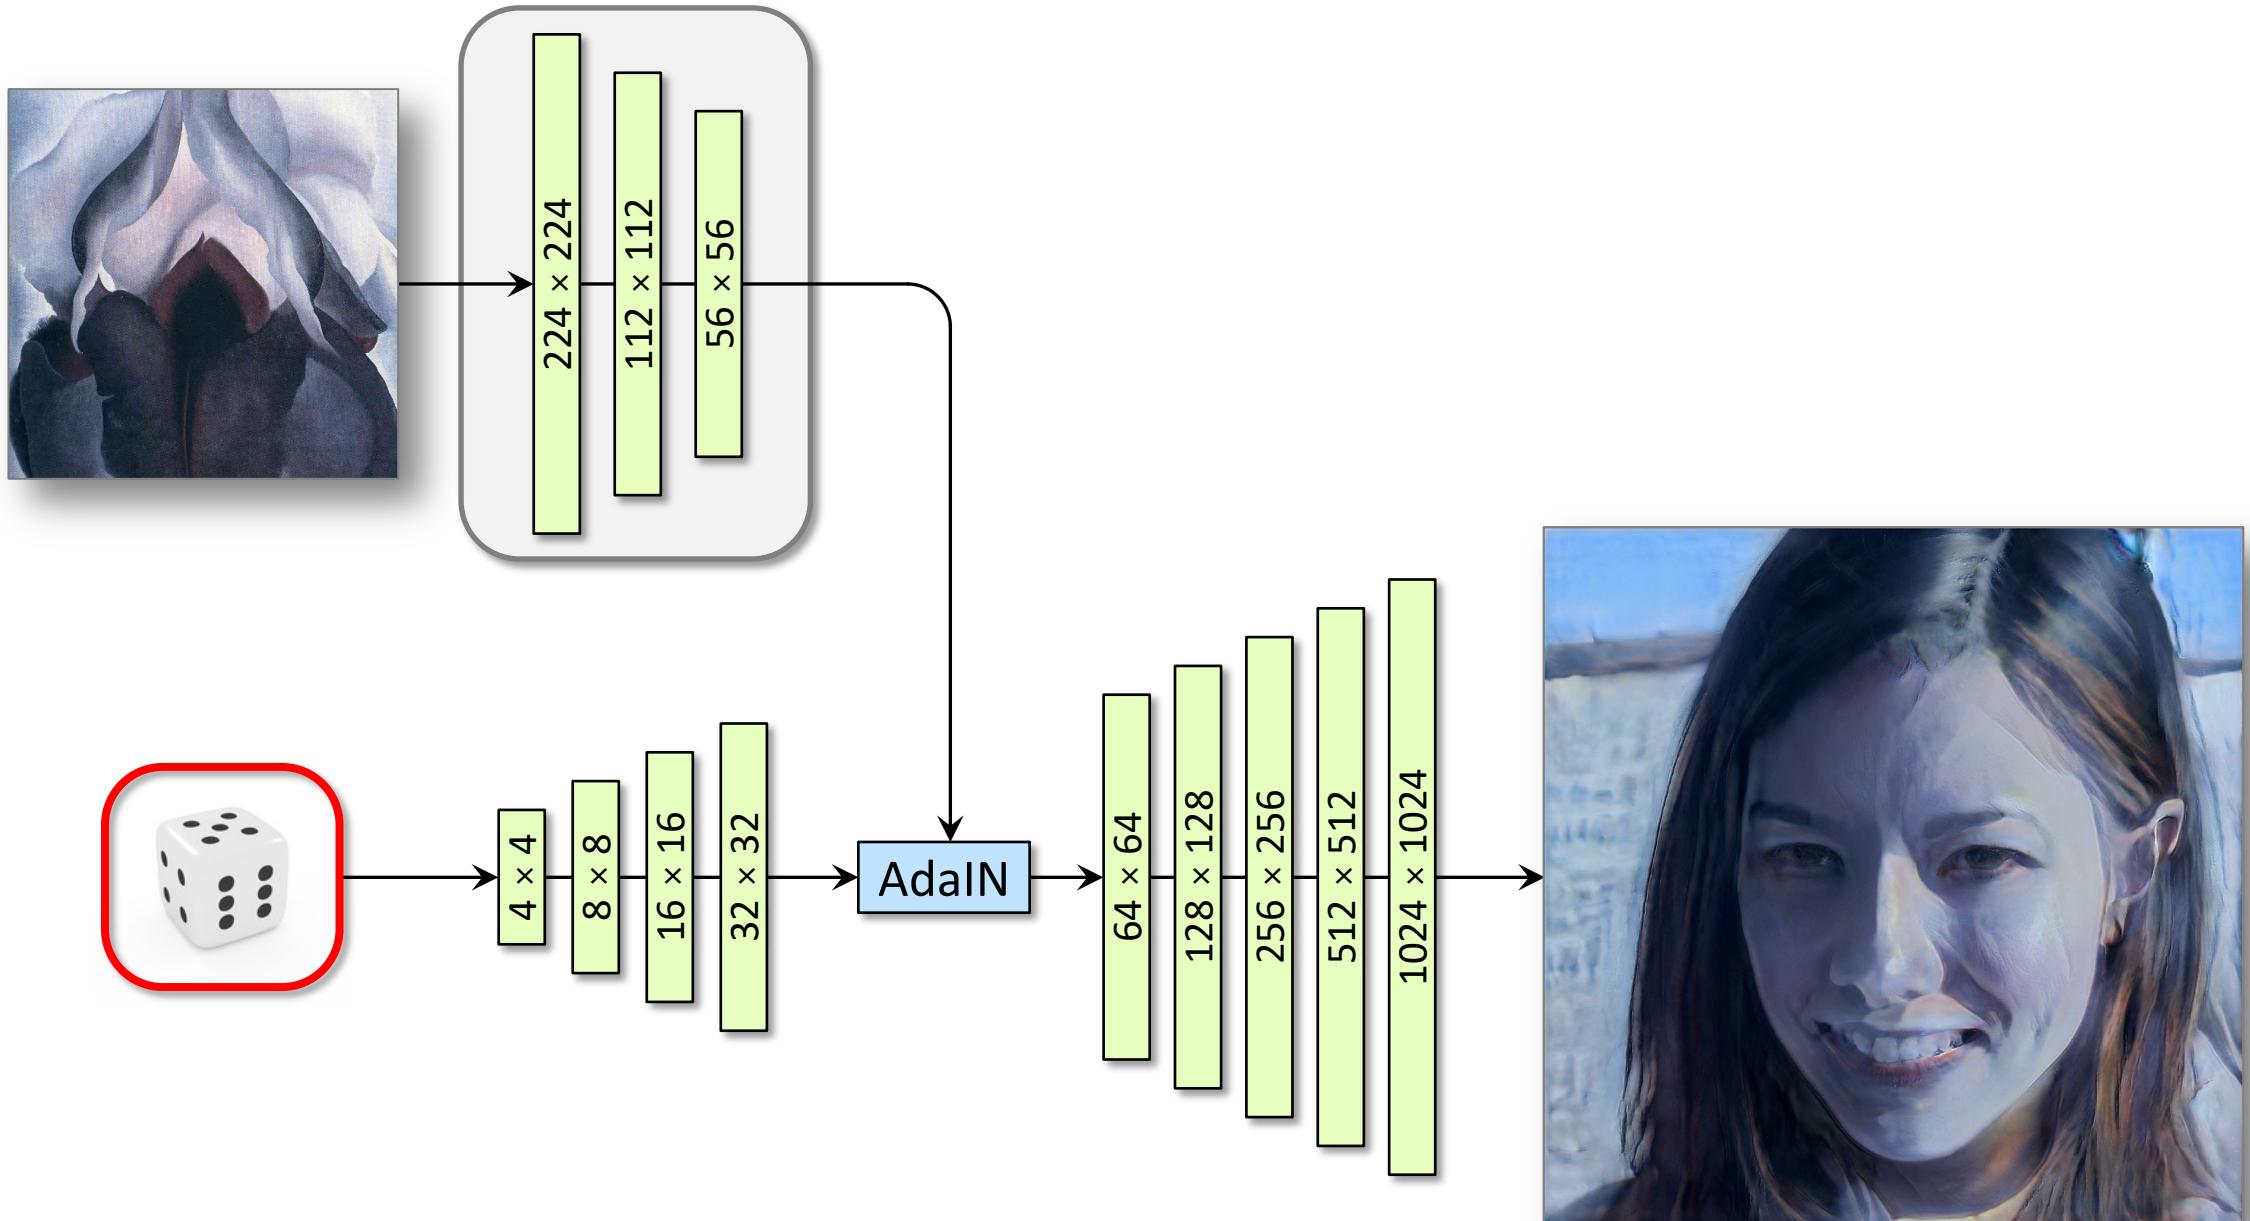
StyleGAN



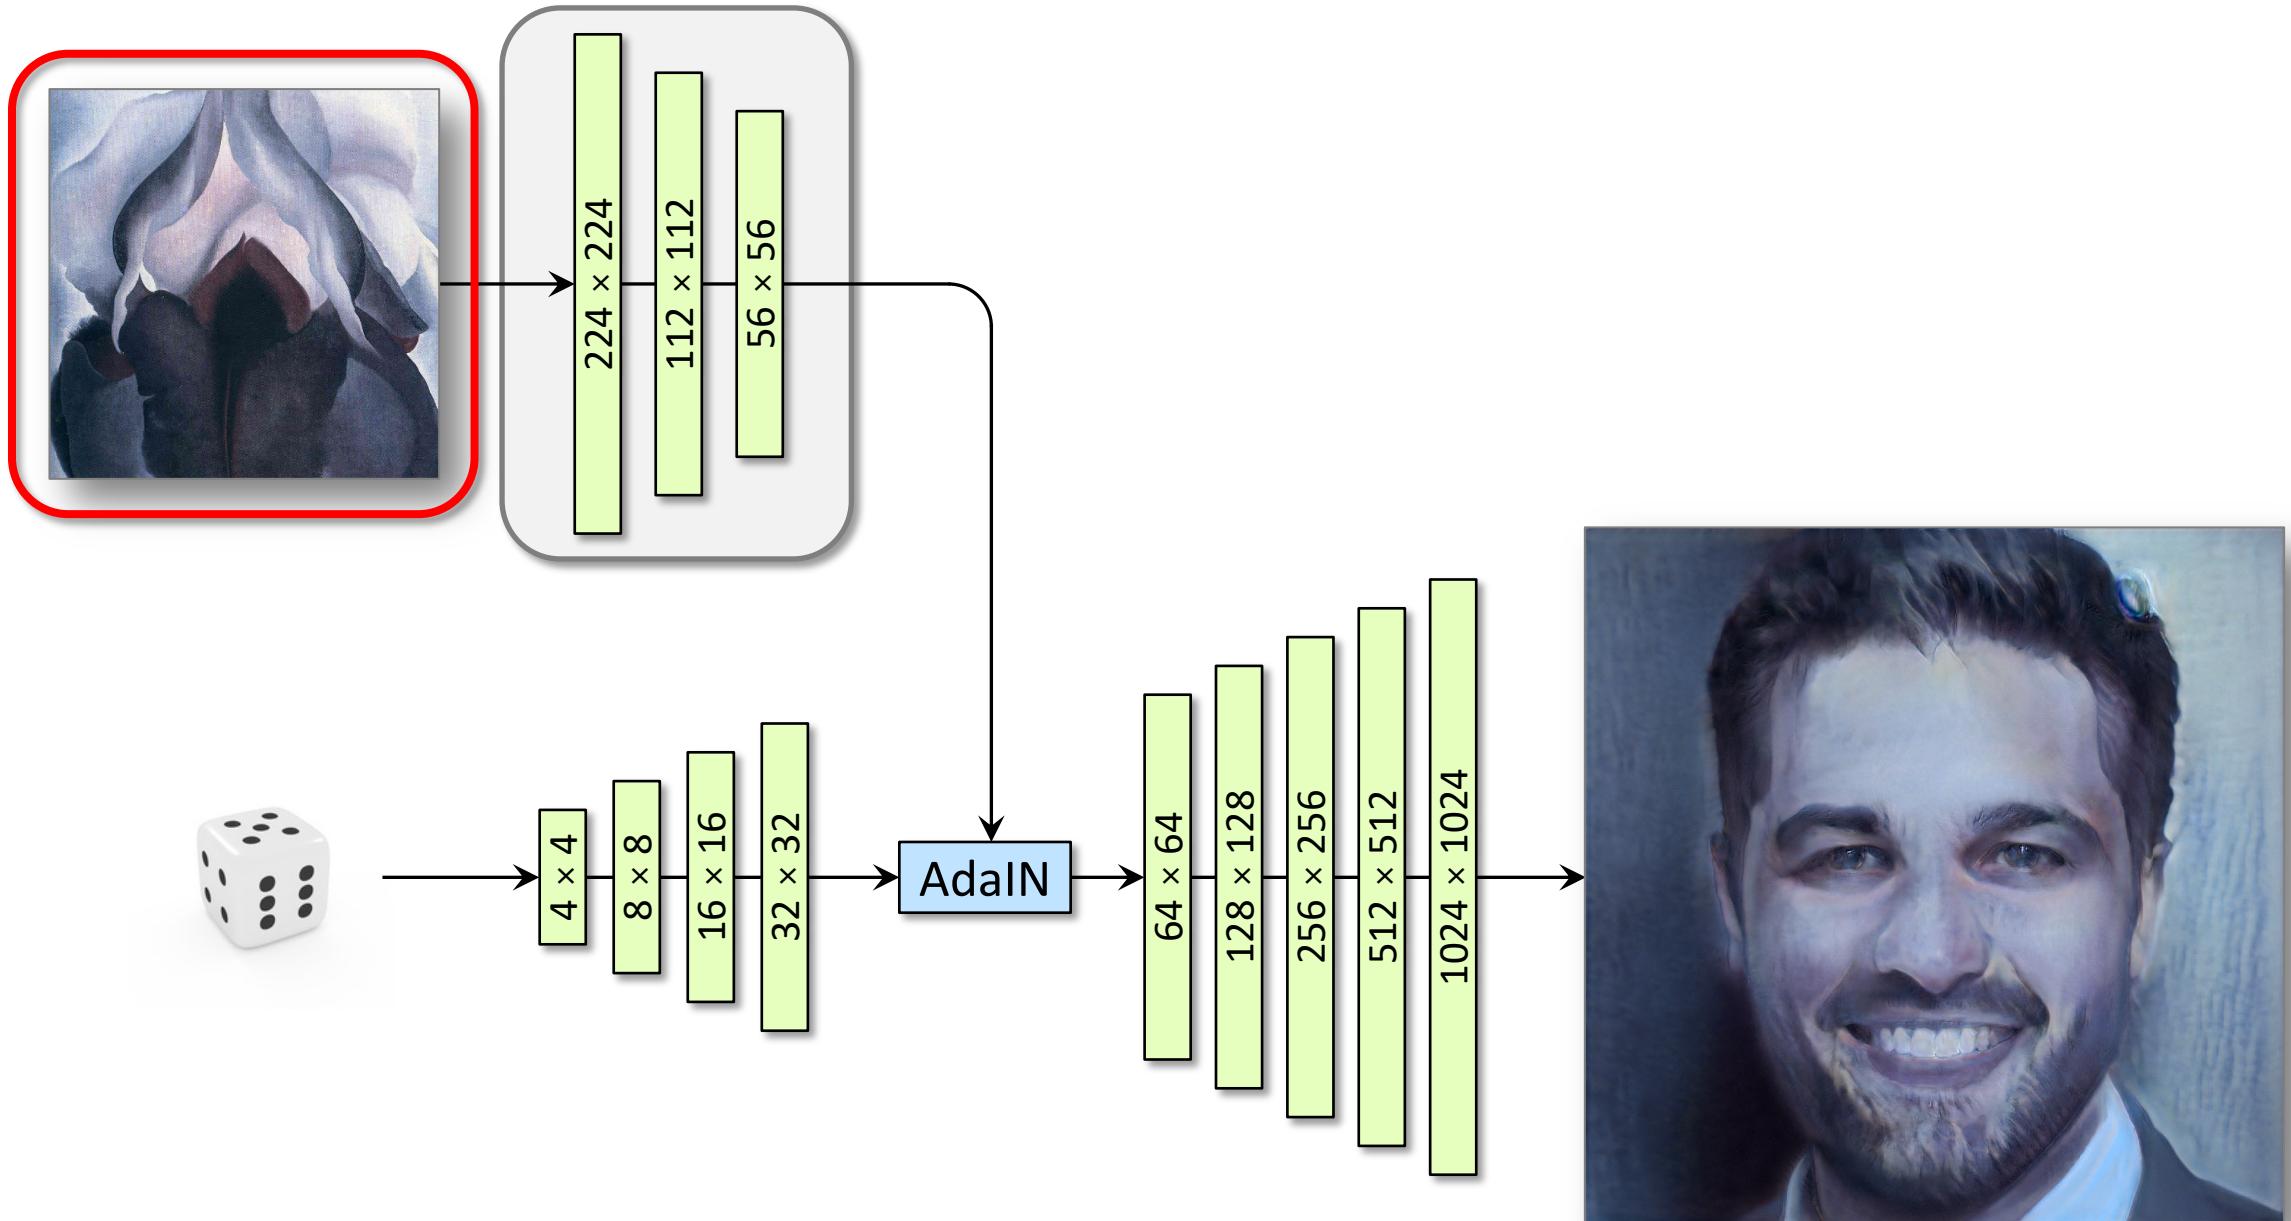
StyleGAN



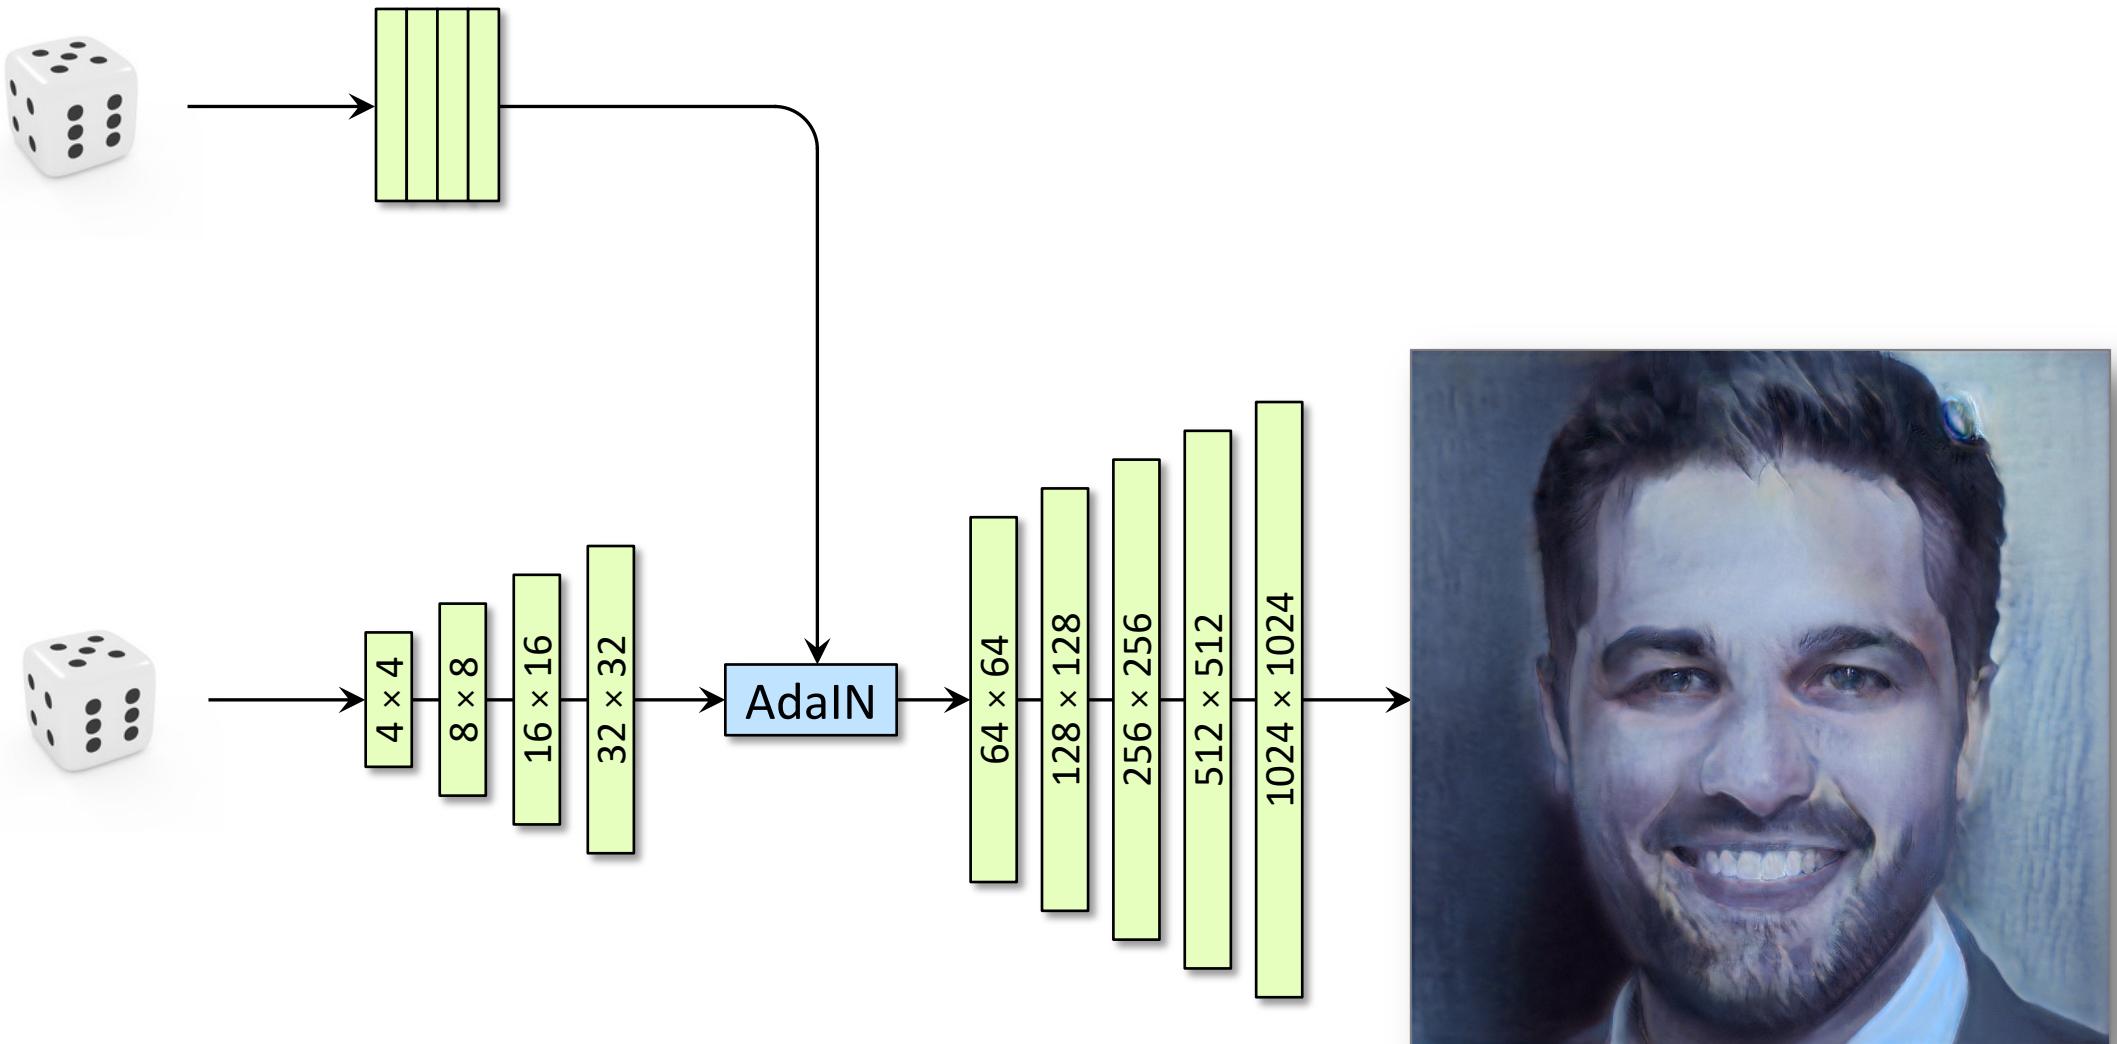
StyleGAN



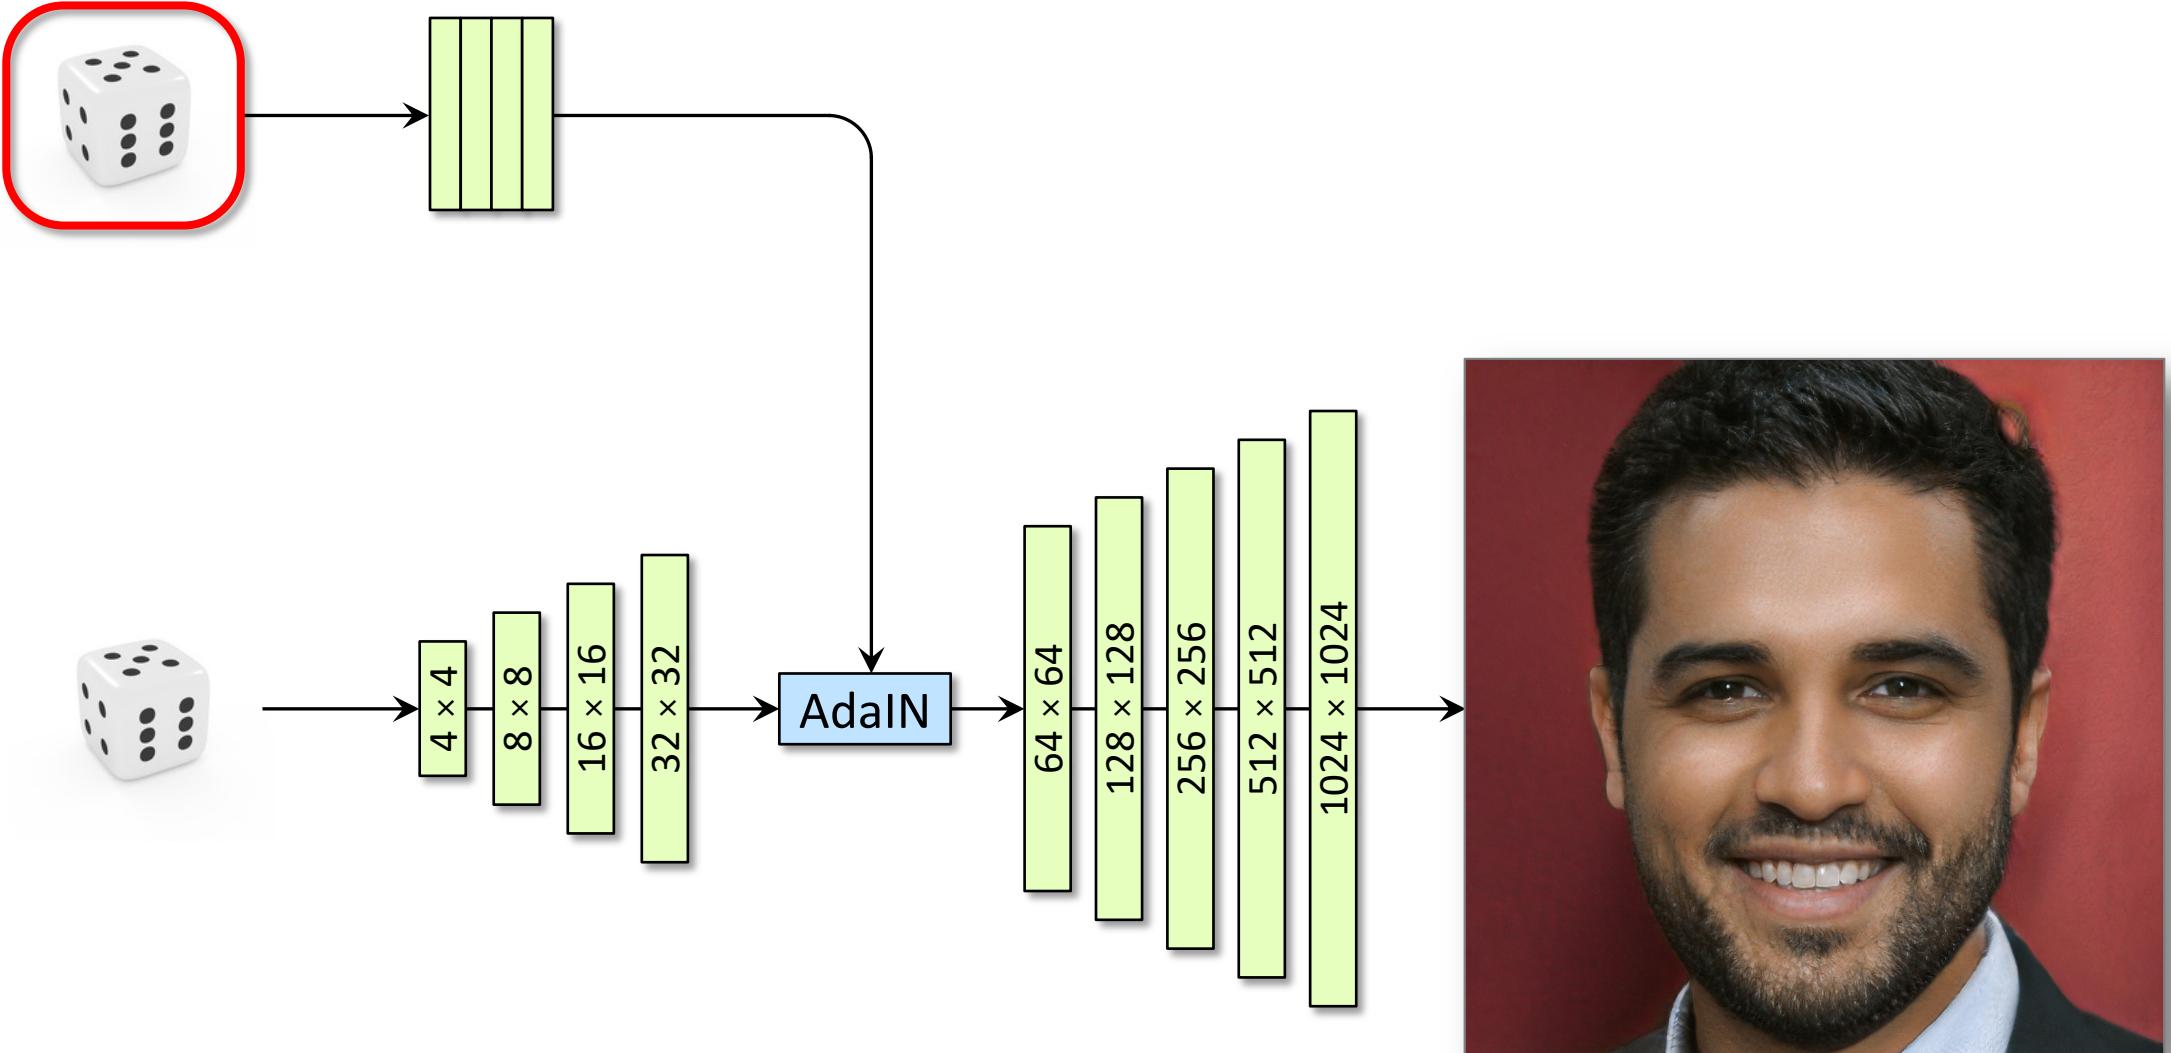
StyleGAN



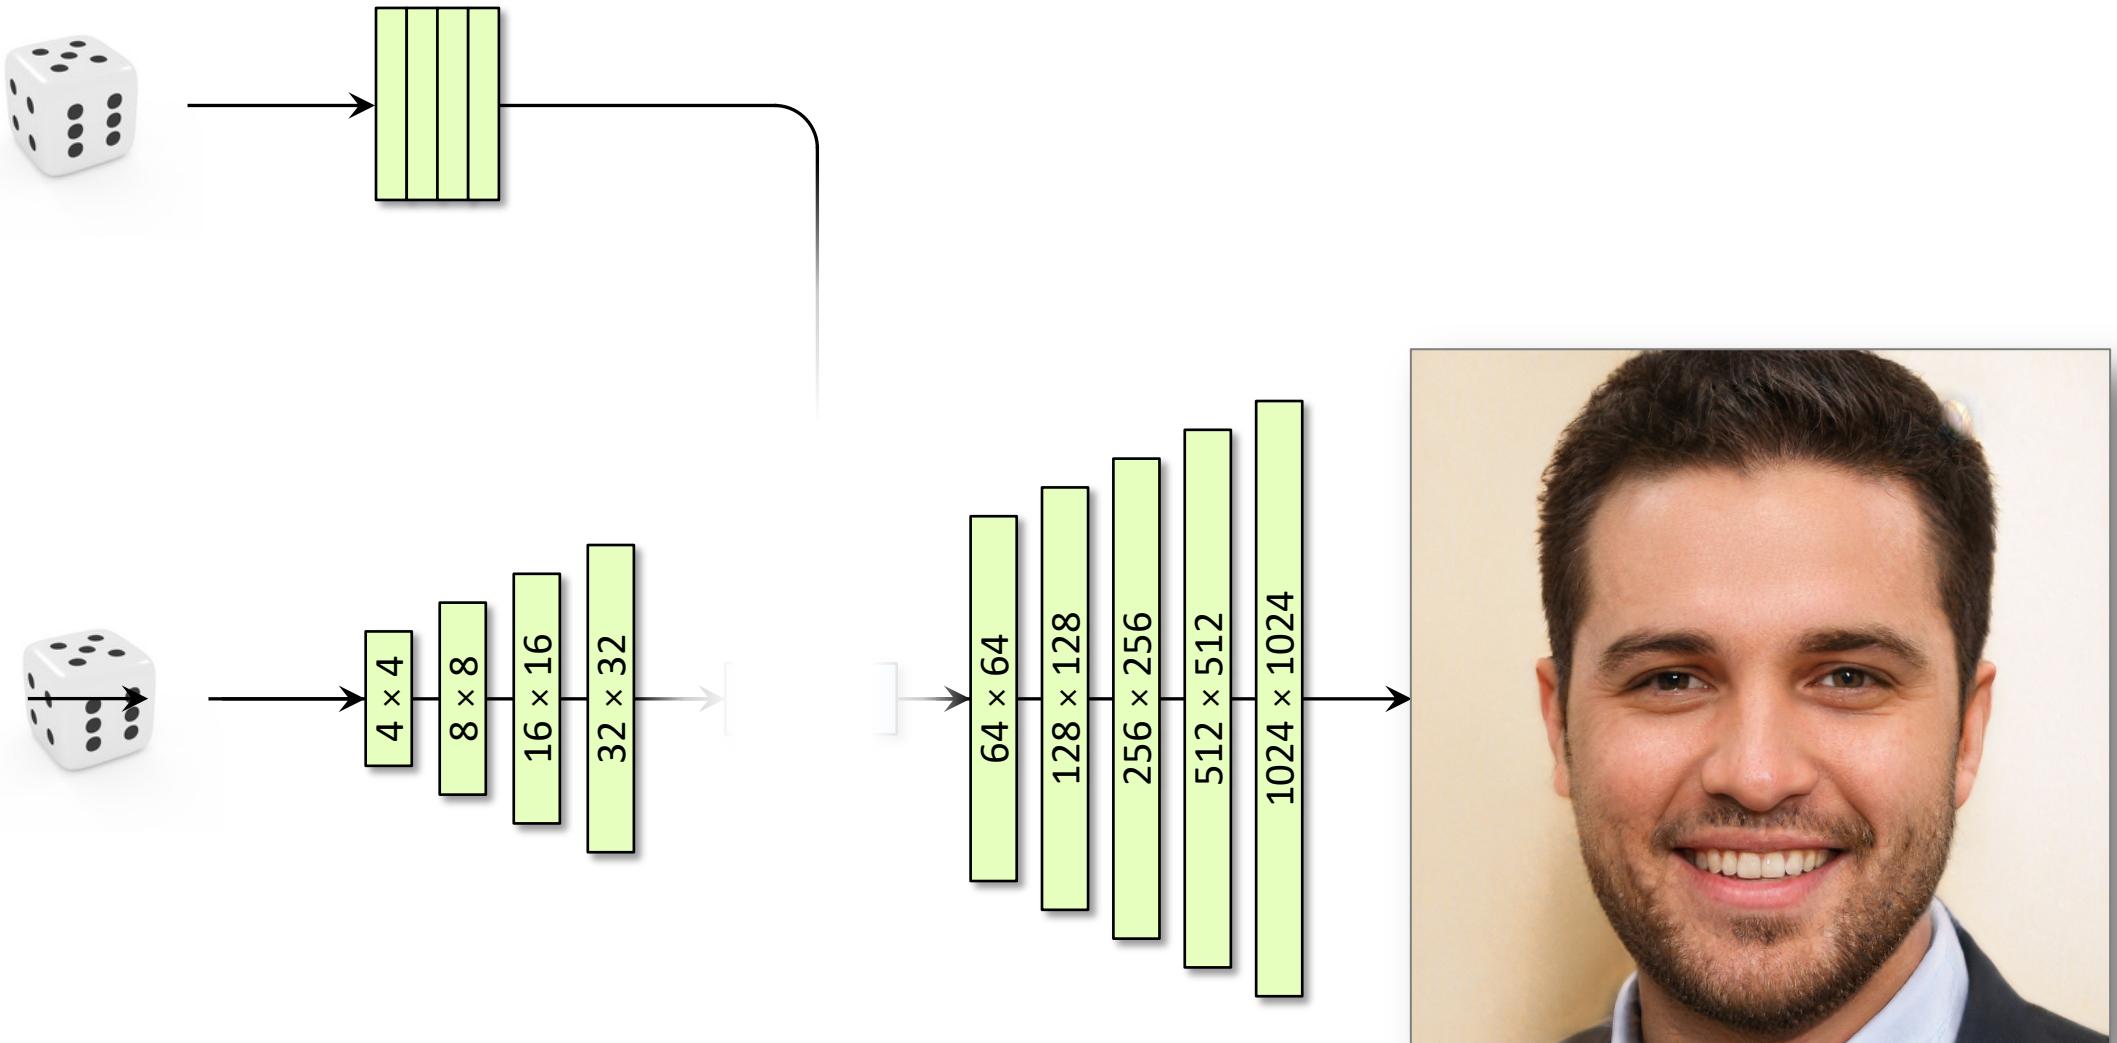
StyleGAN



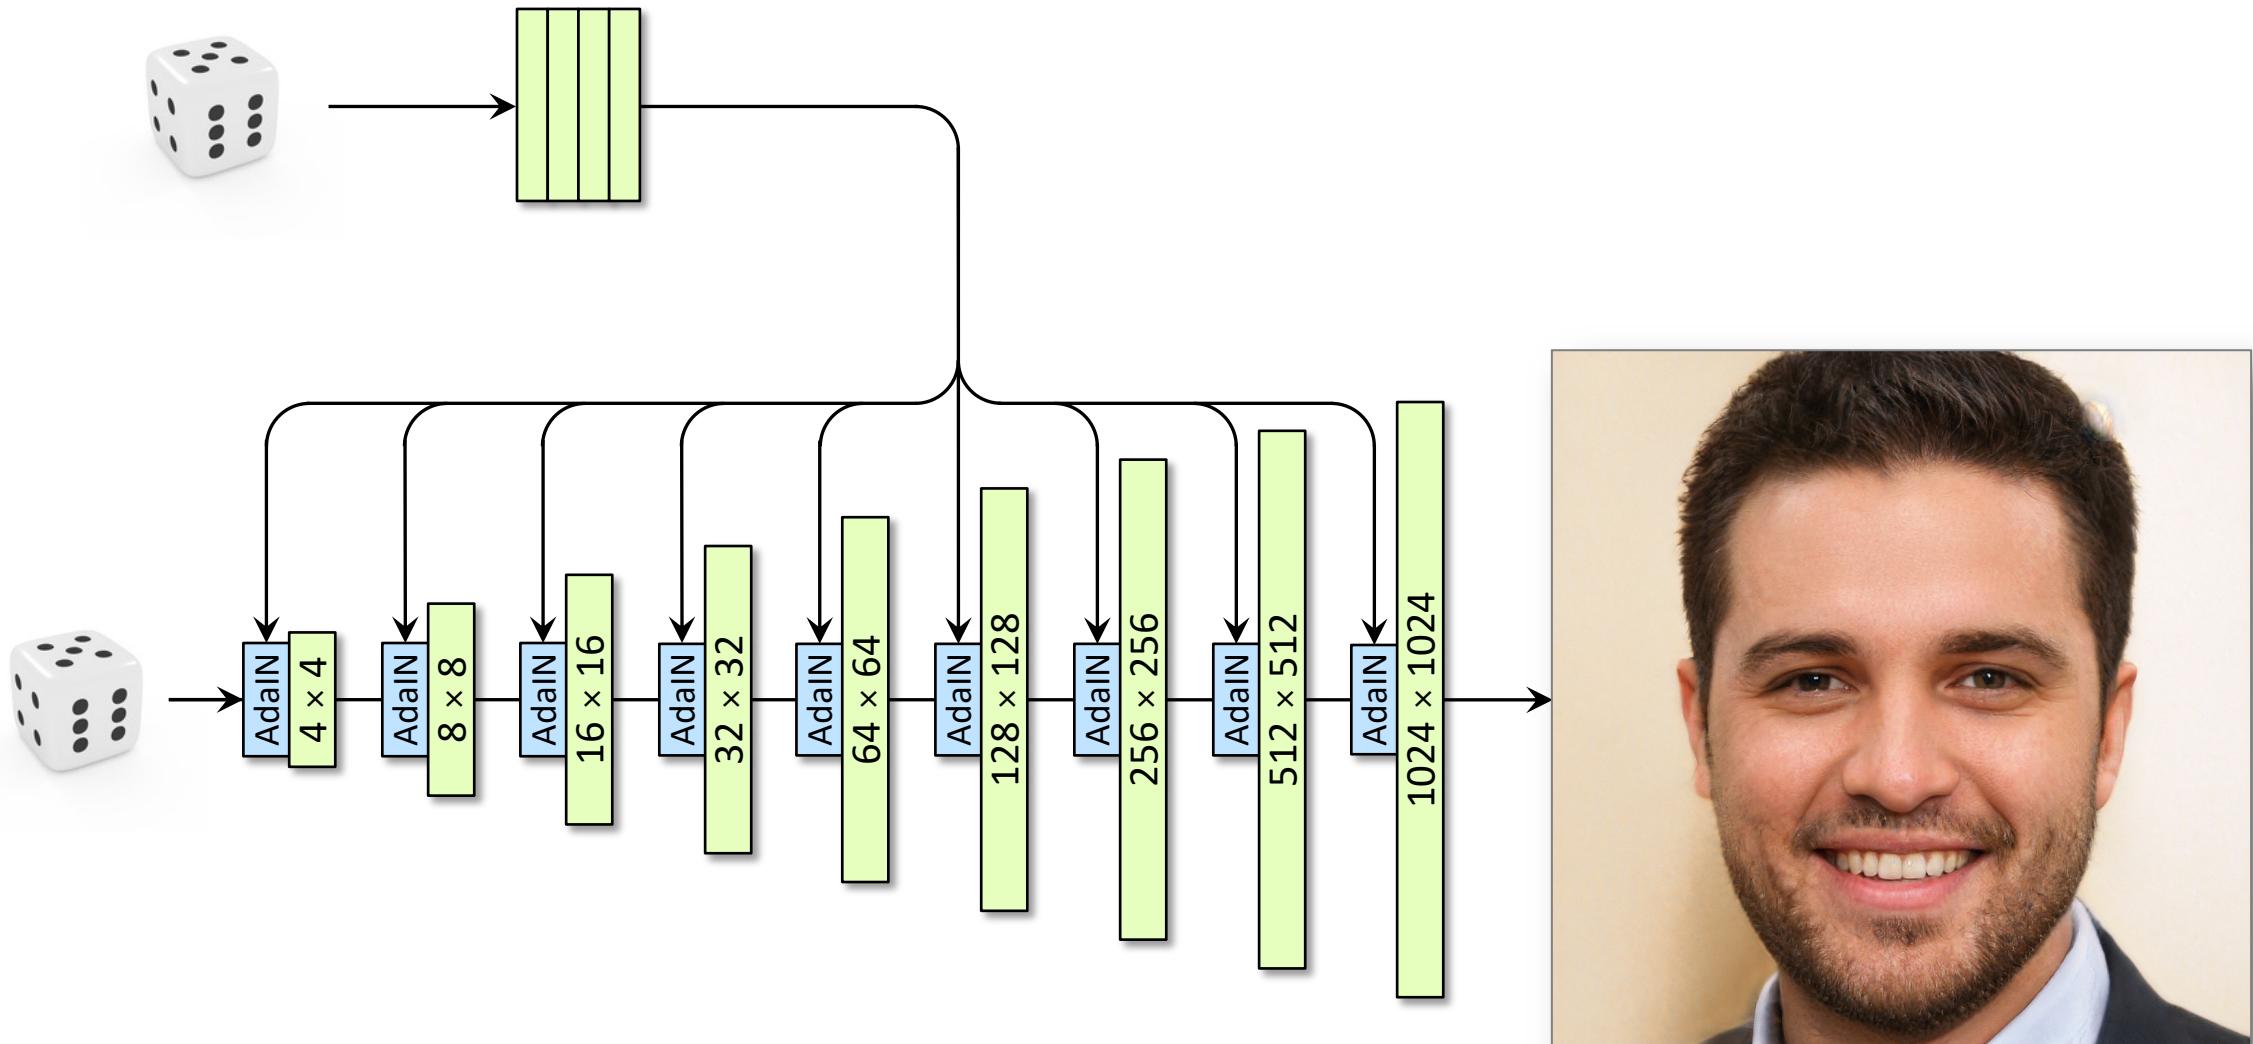
StyleGAN



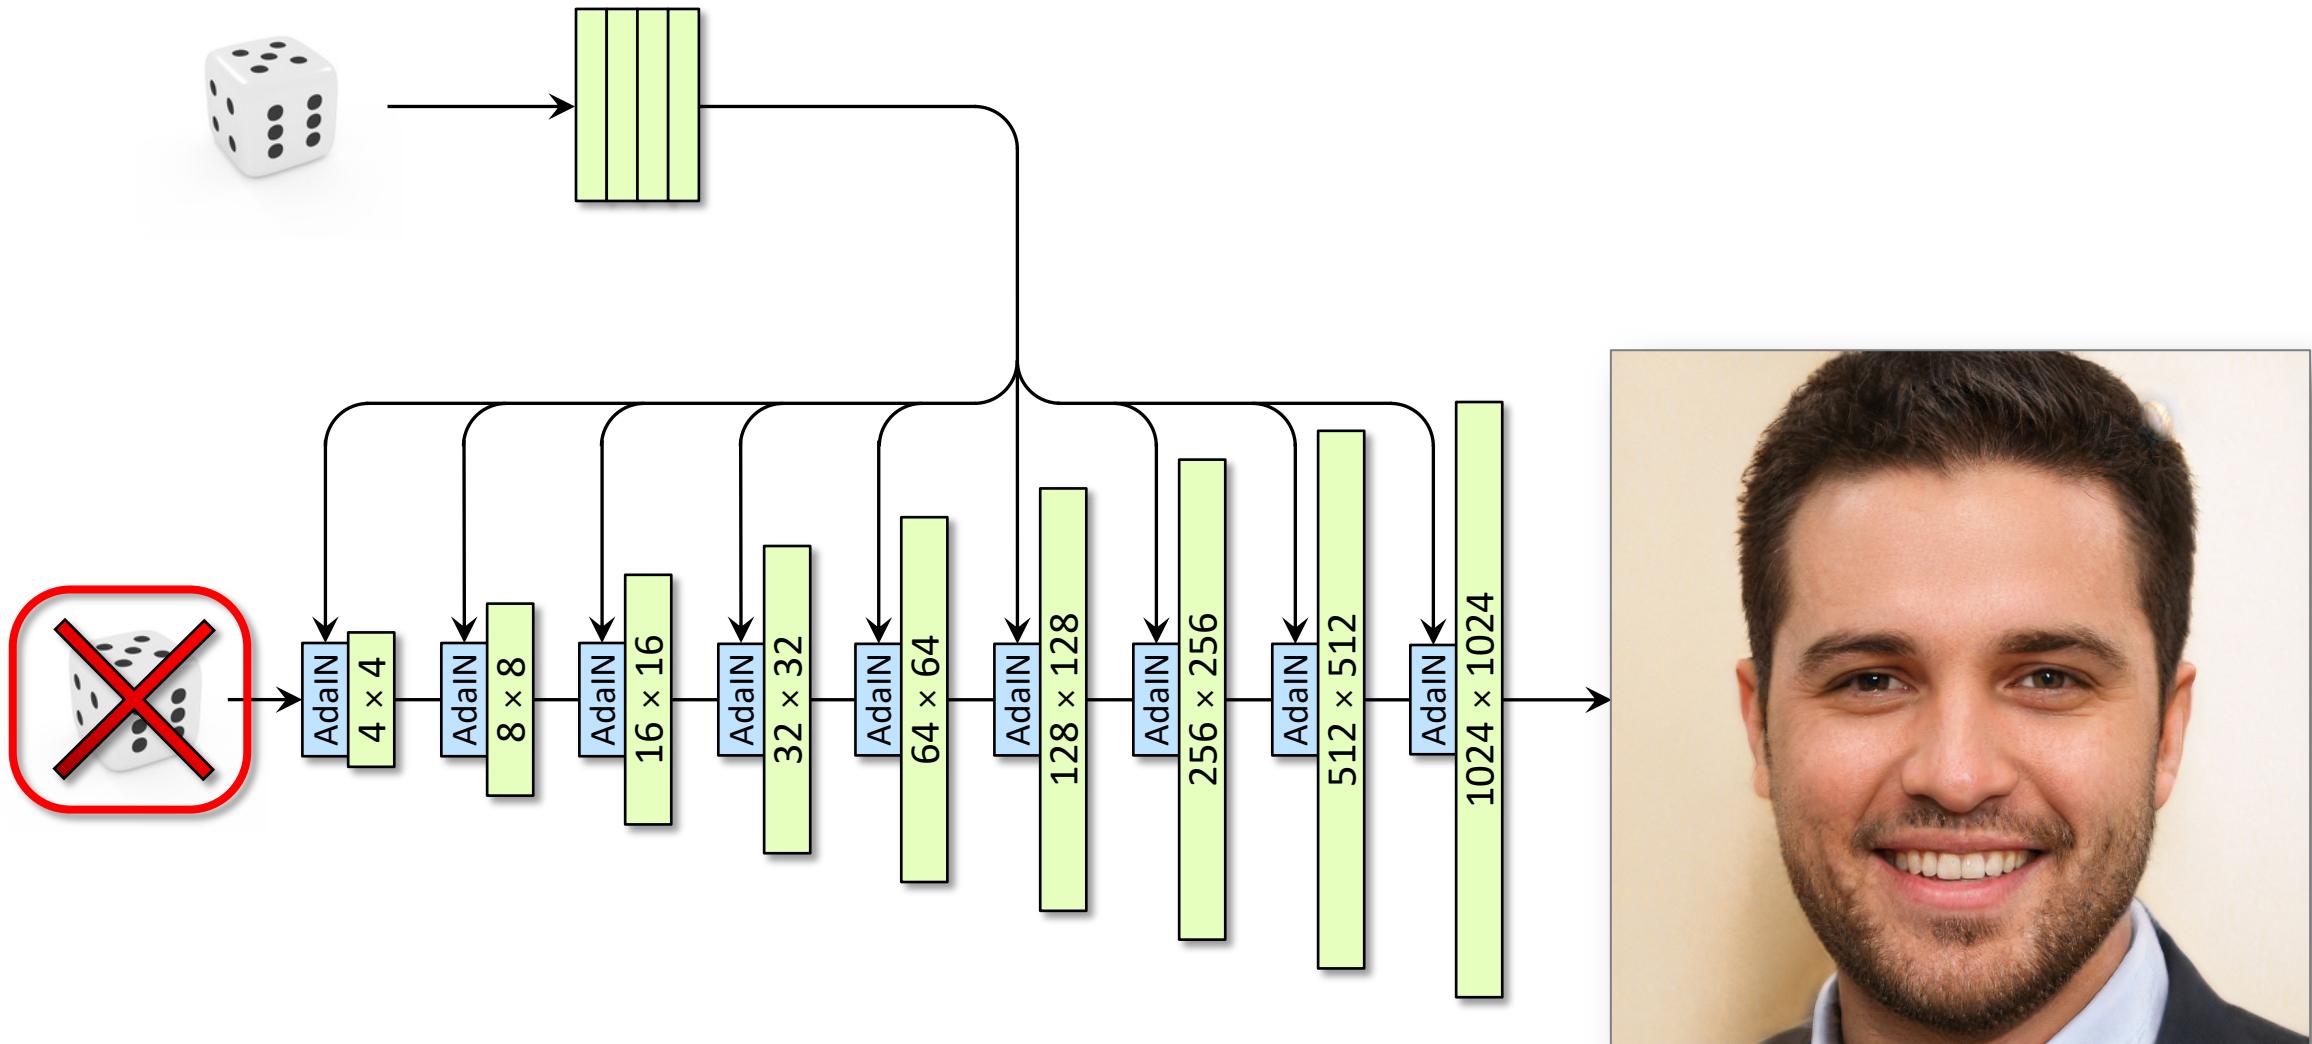
StyleGAN



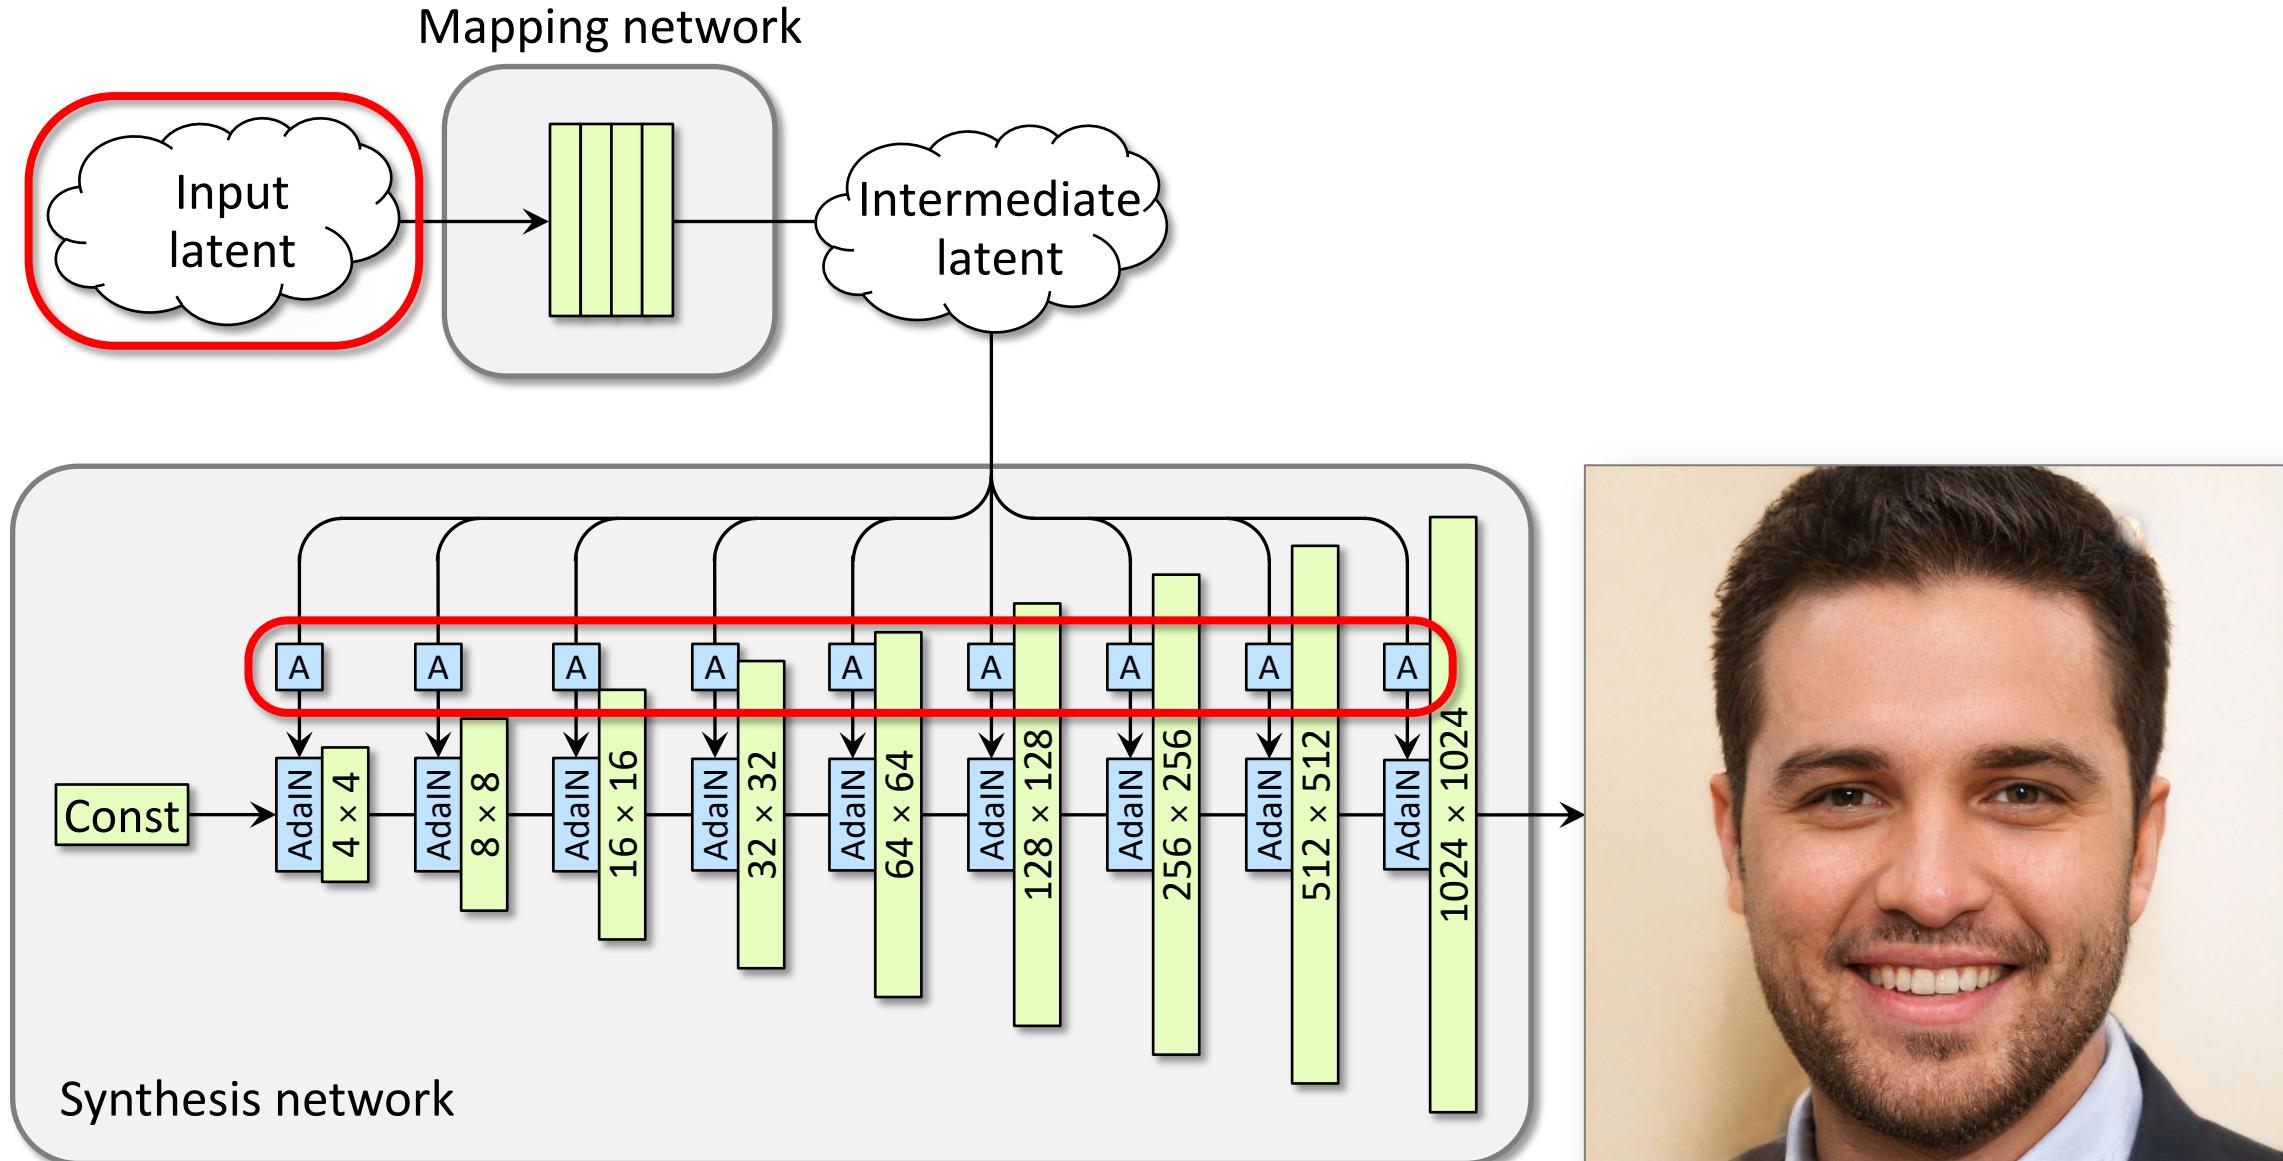
StyleGAN



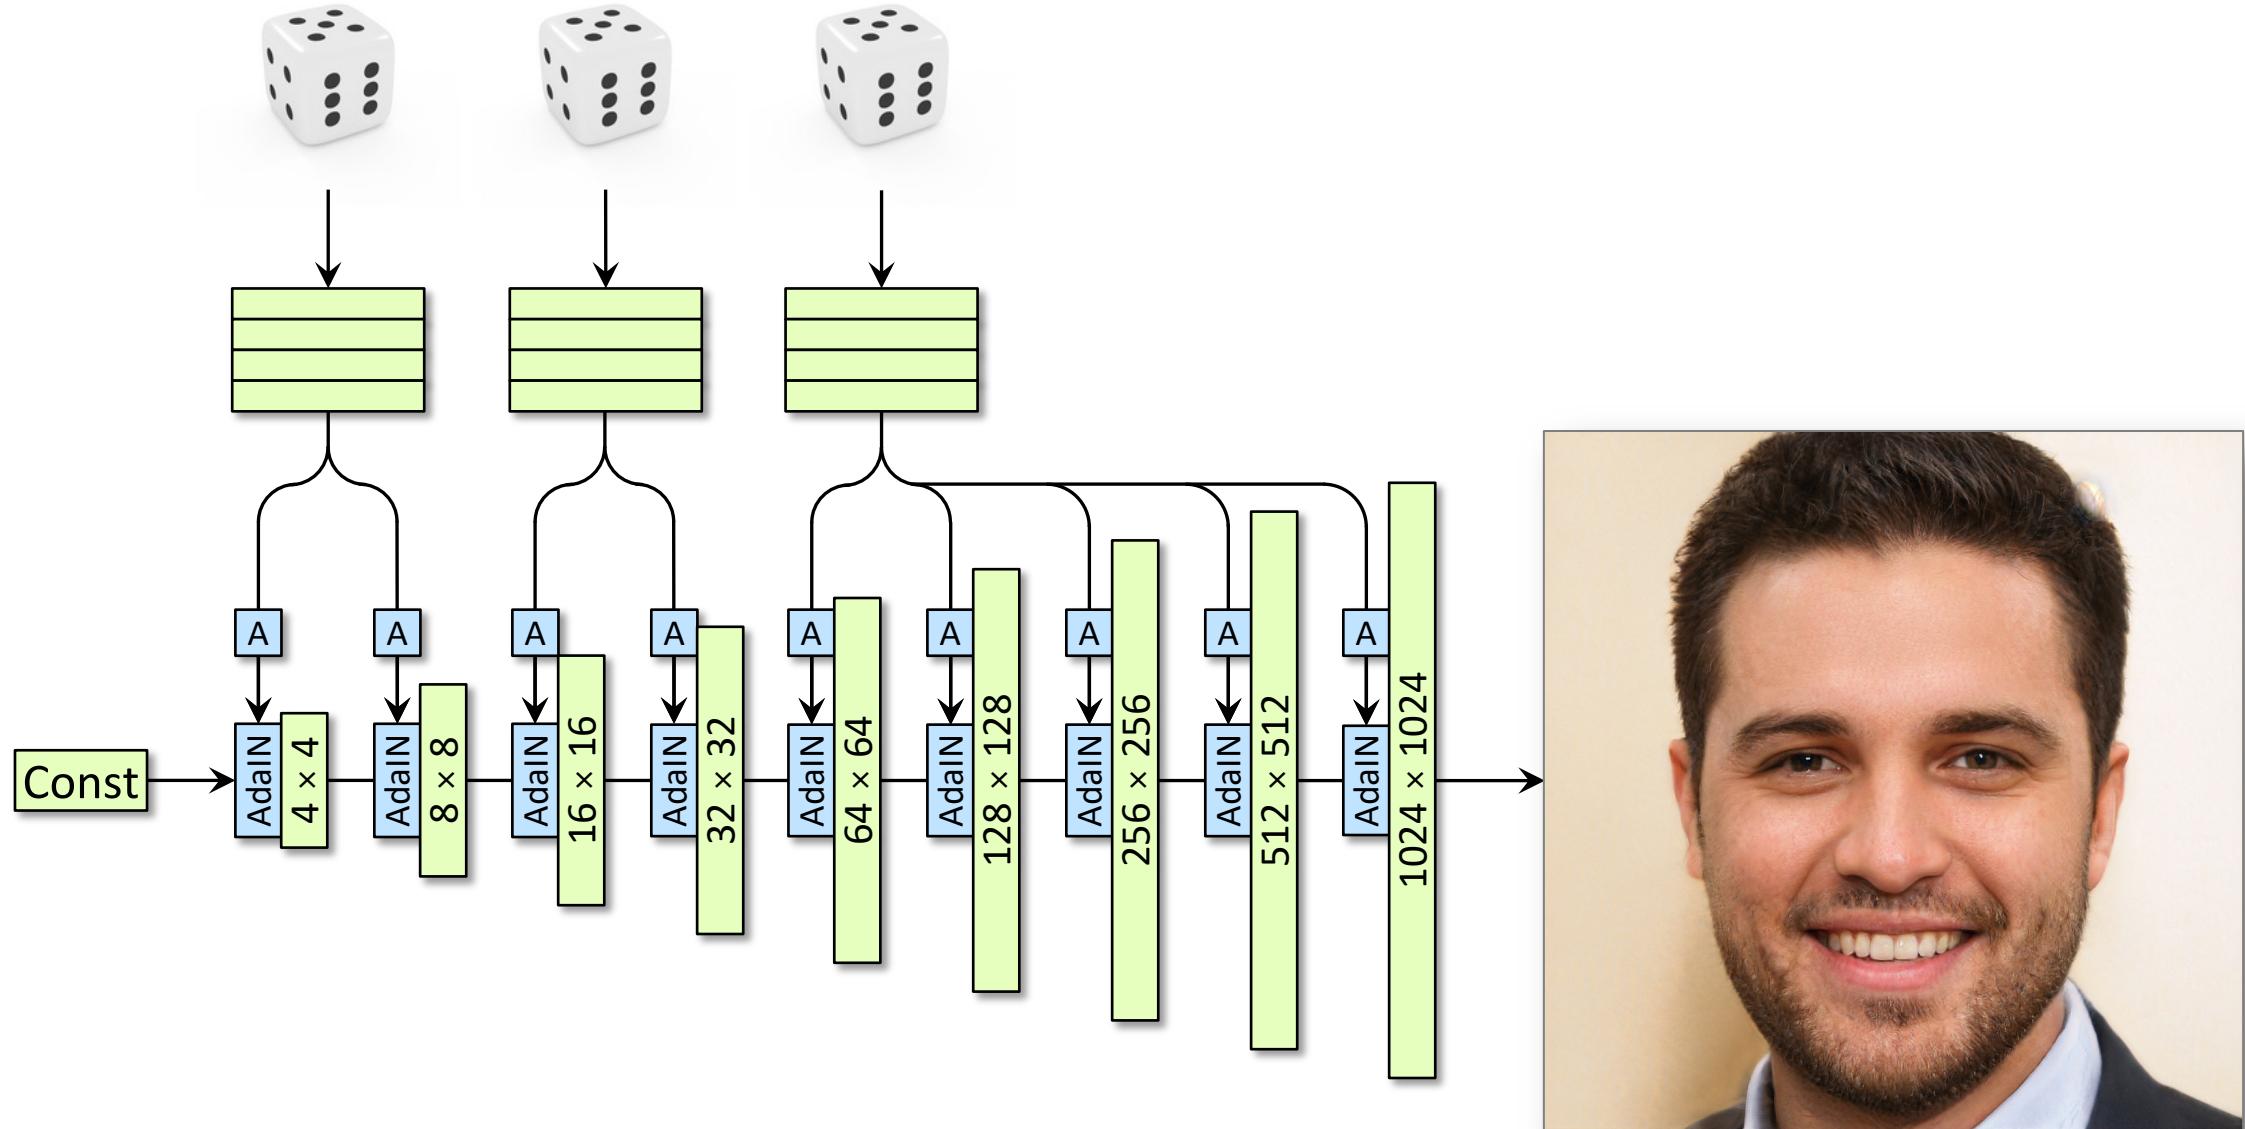
StyleGAN



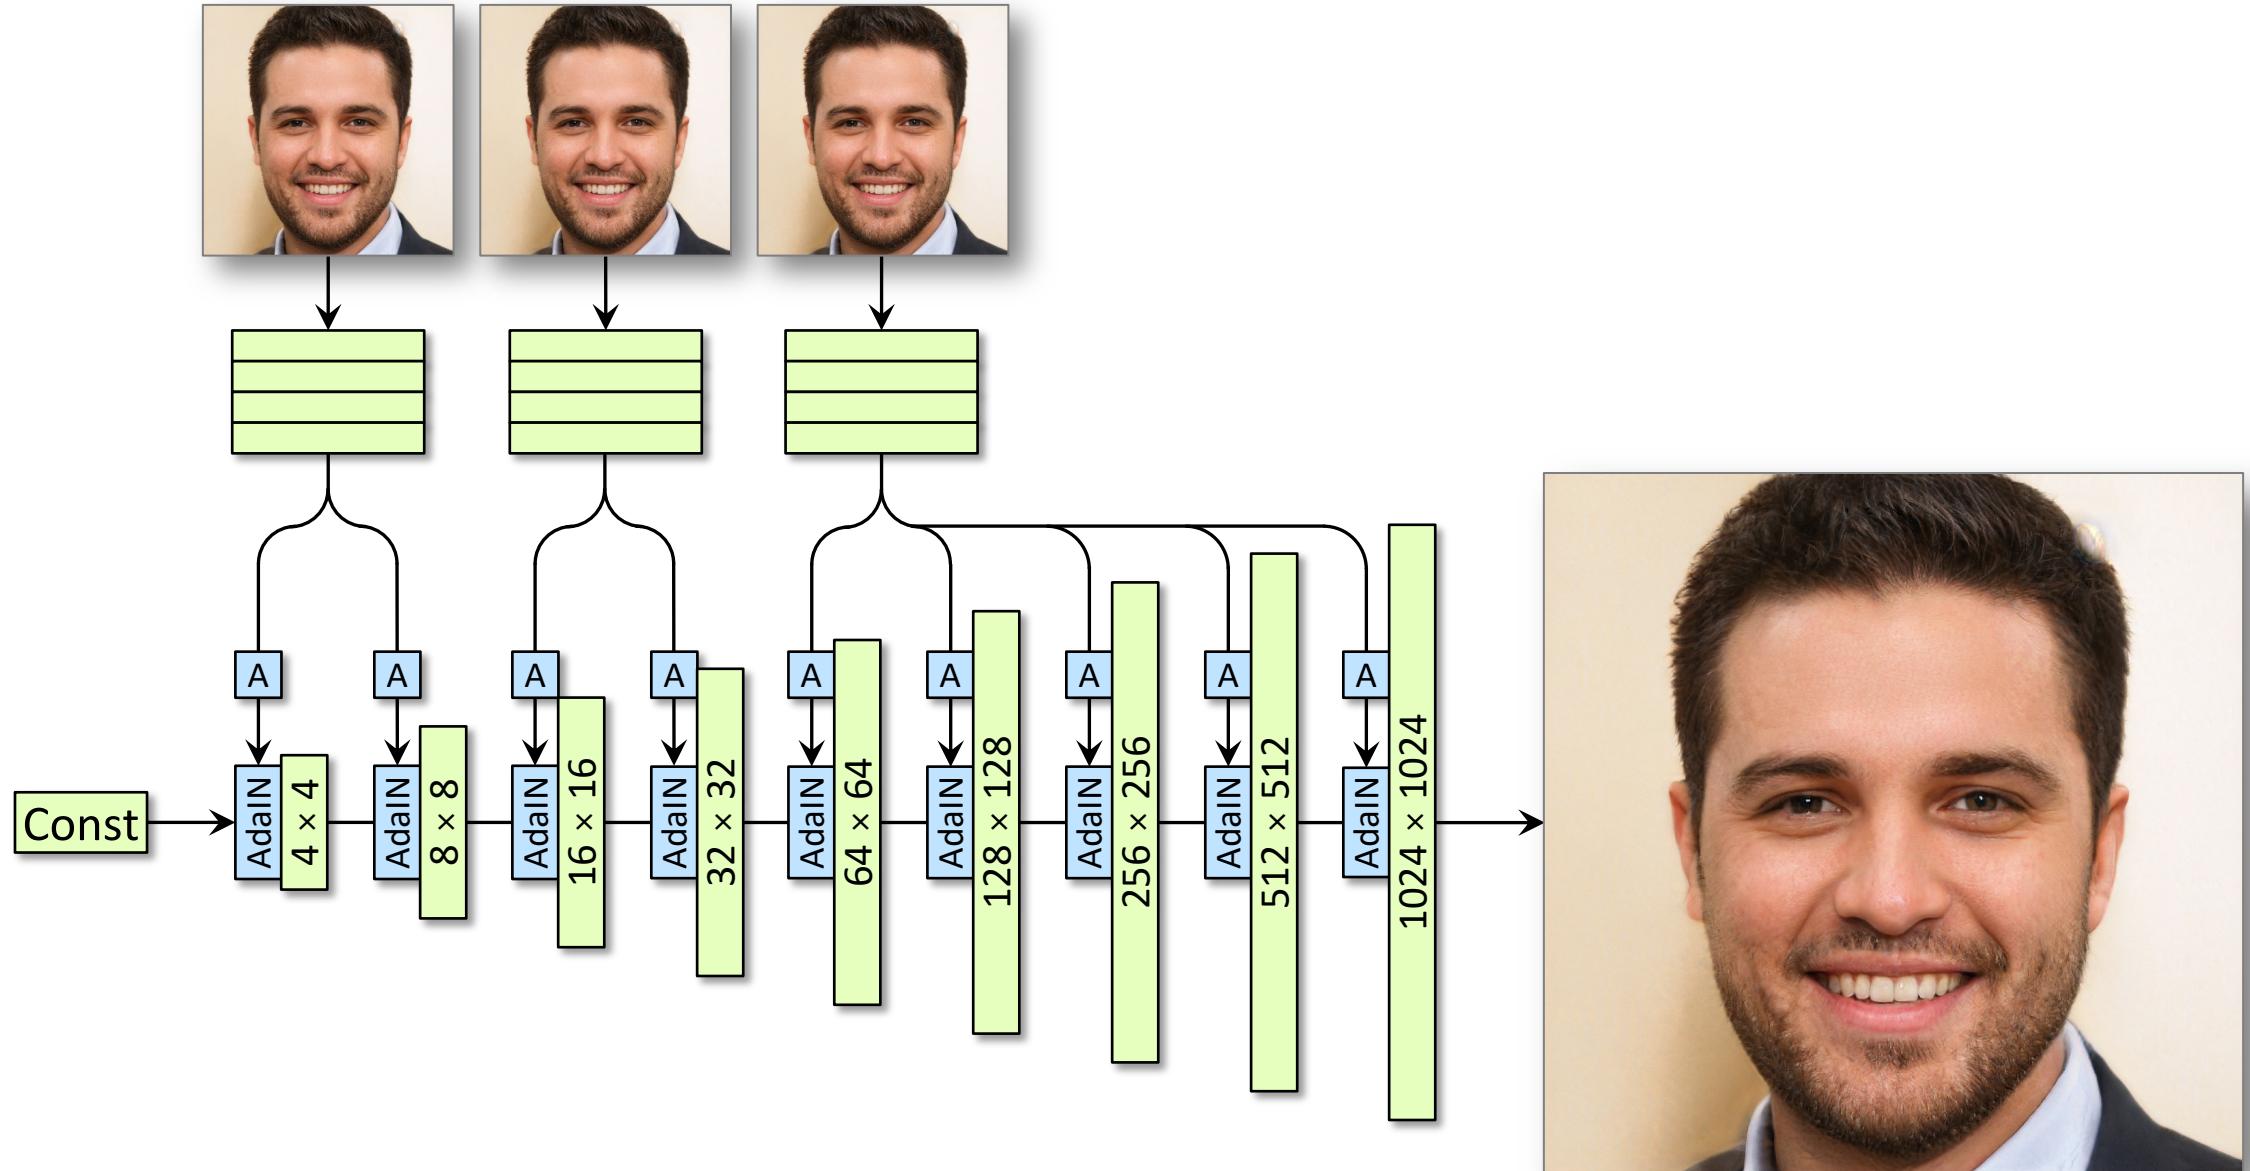
StyleGAN



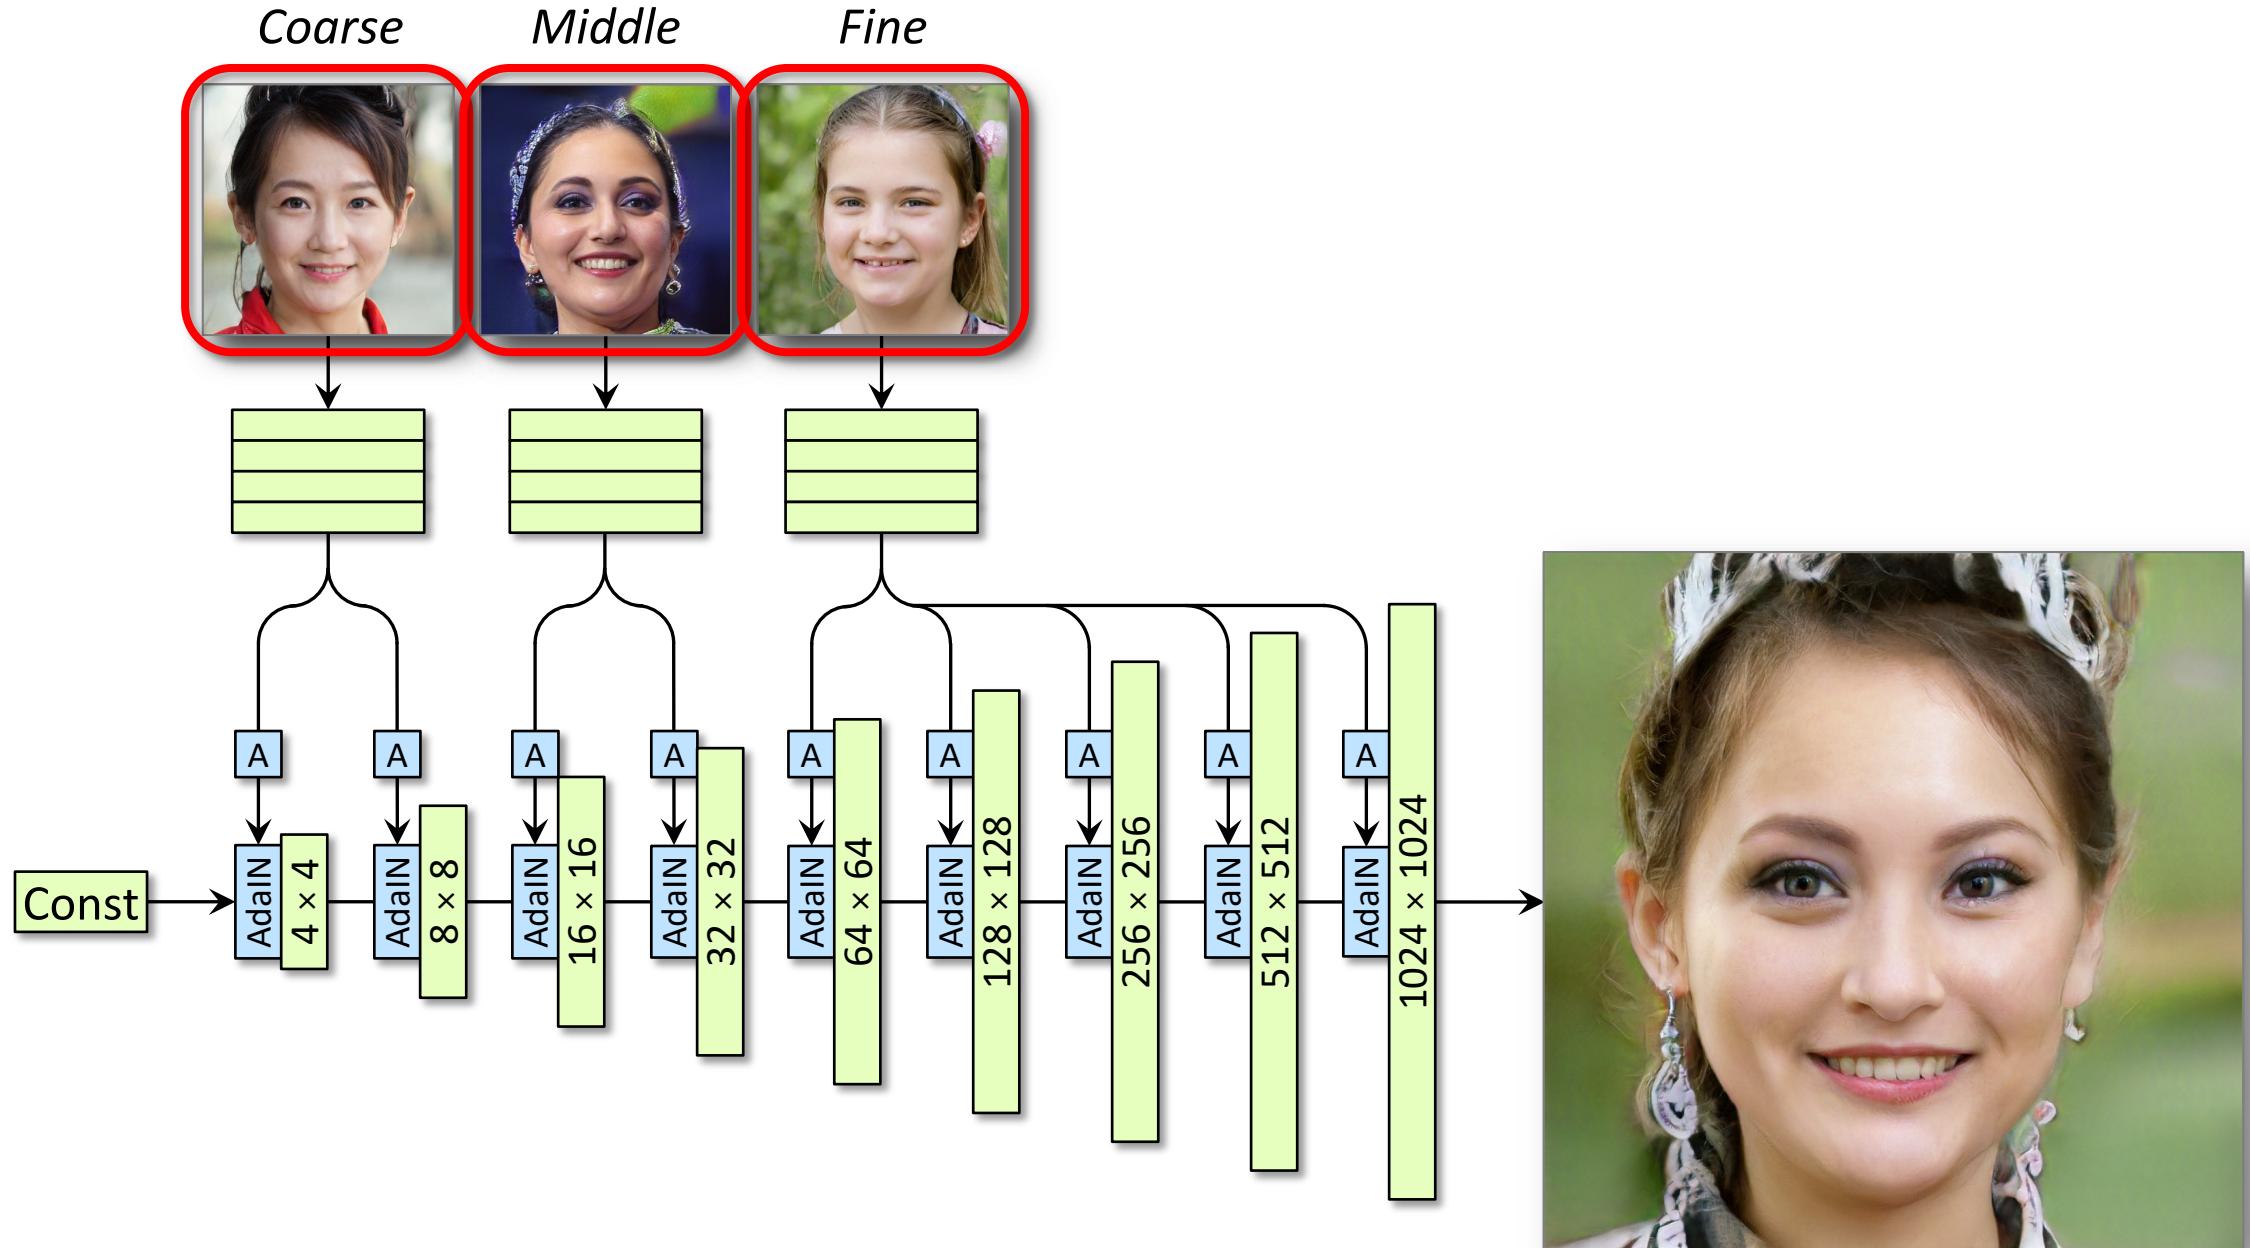
StyleGAN



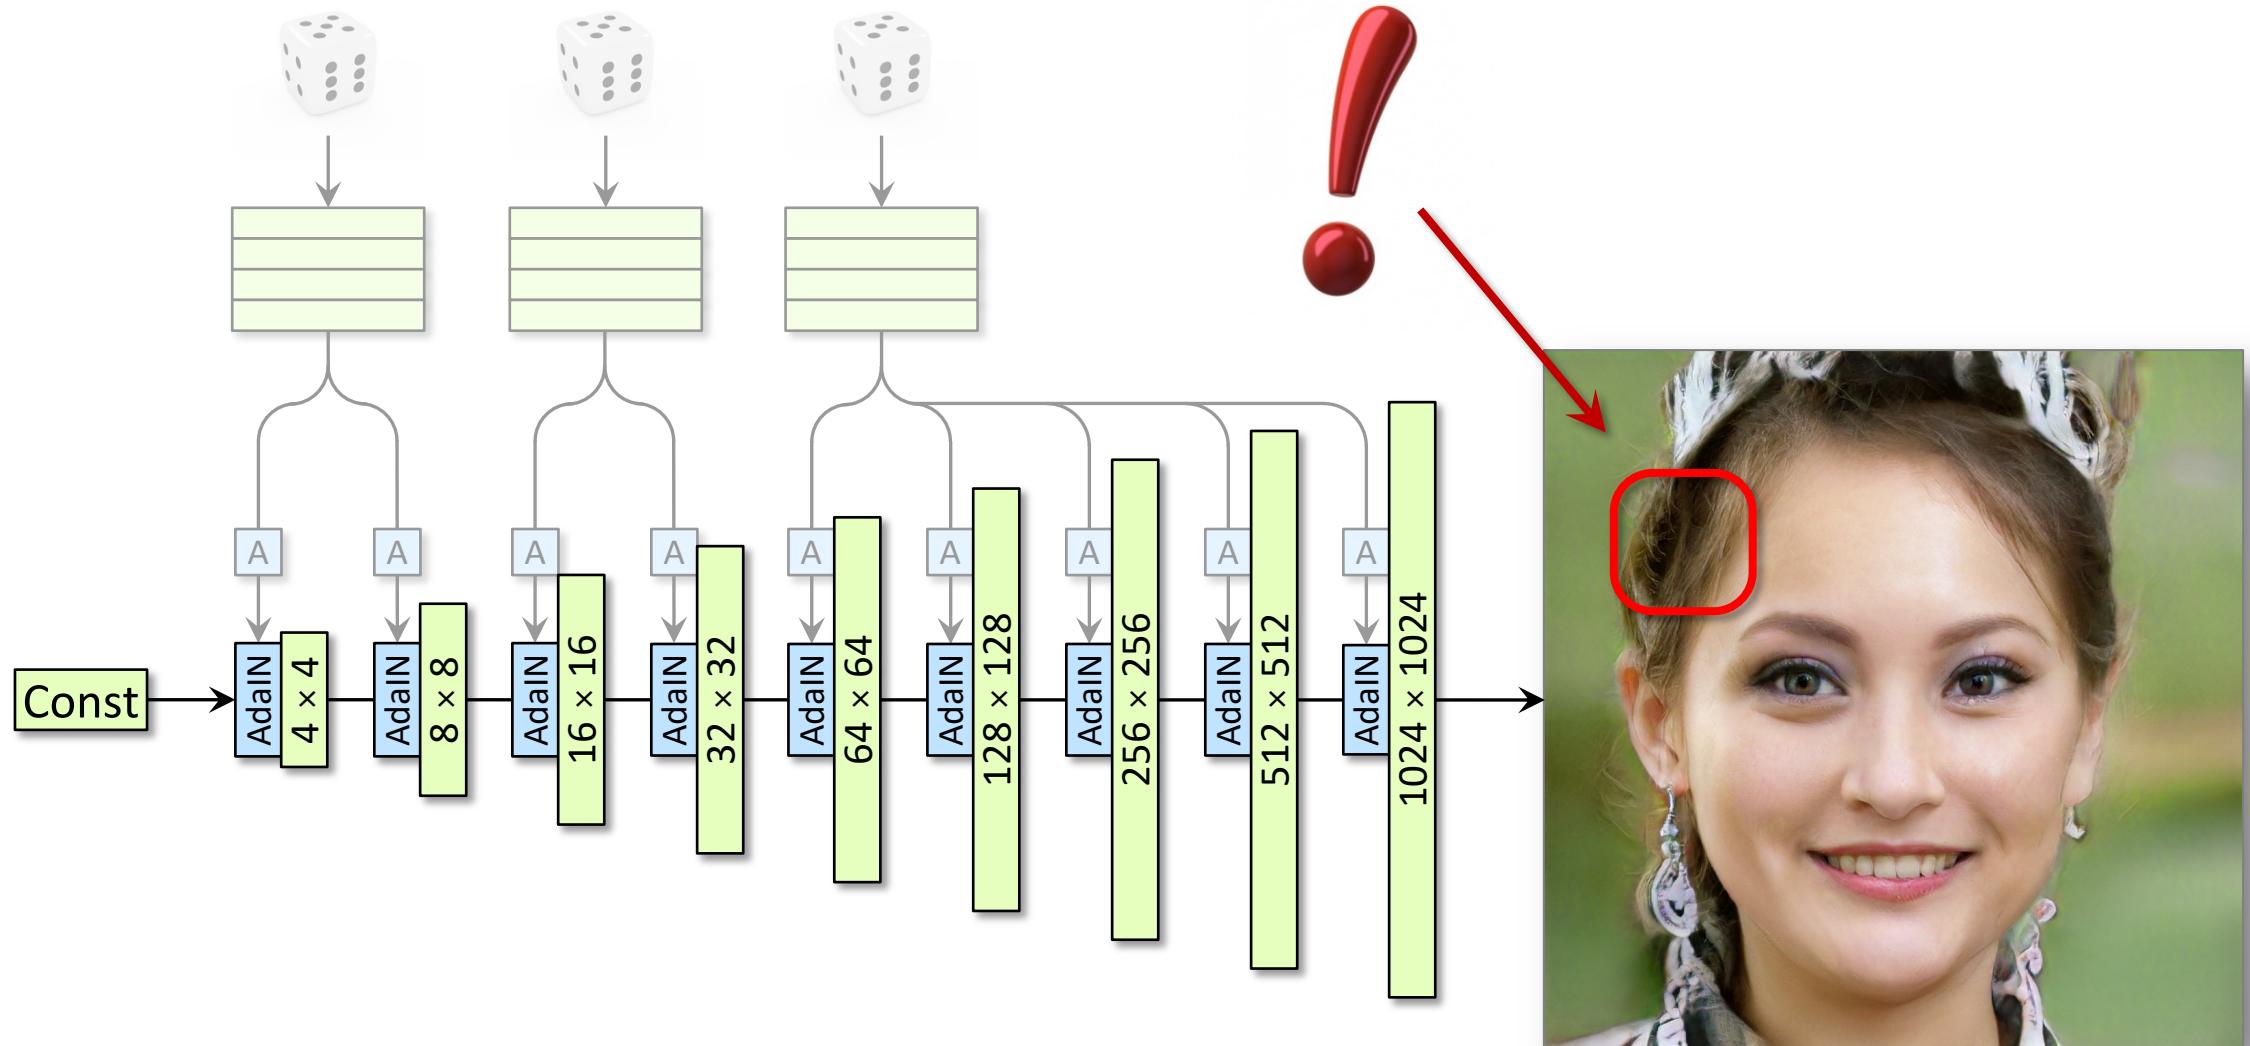
StyleGAN



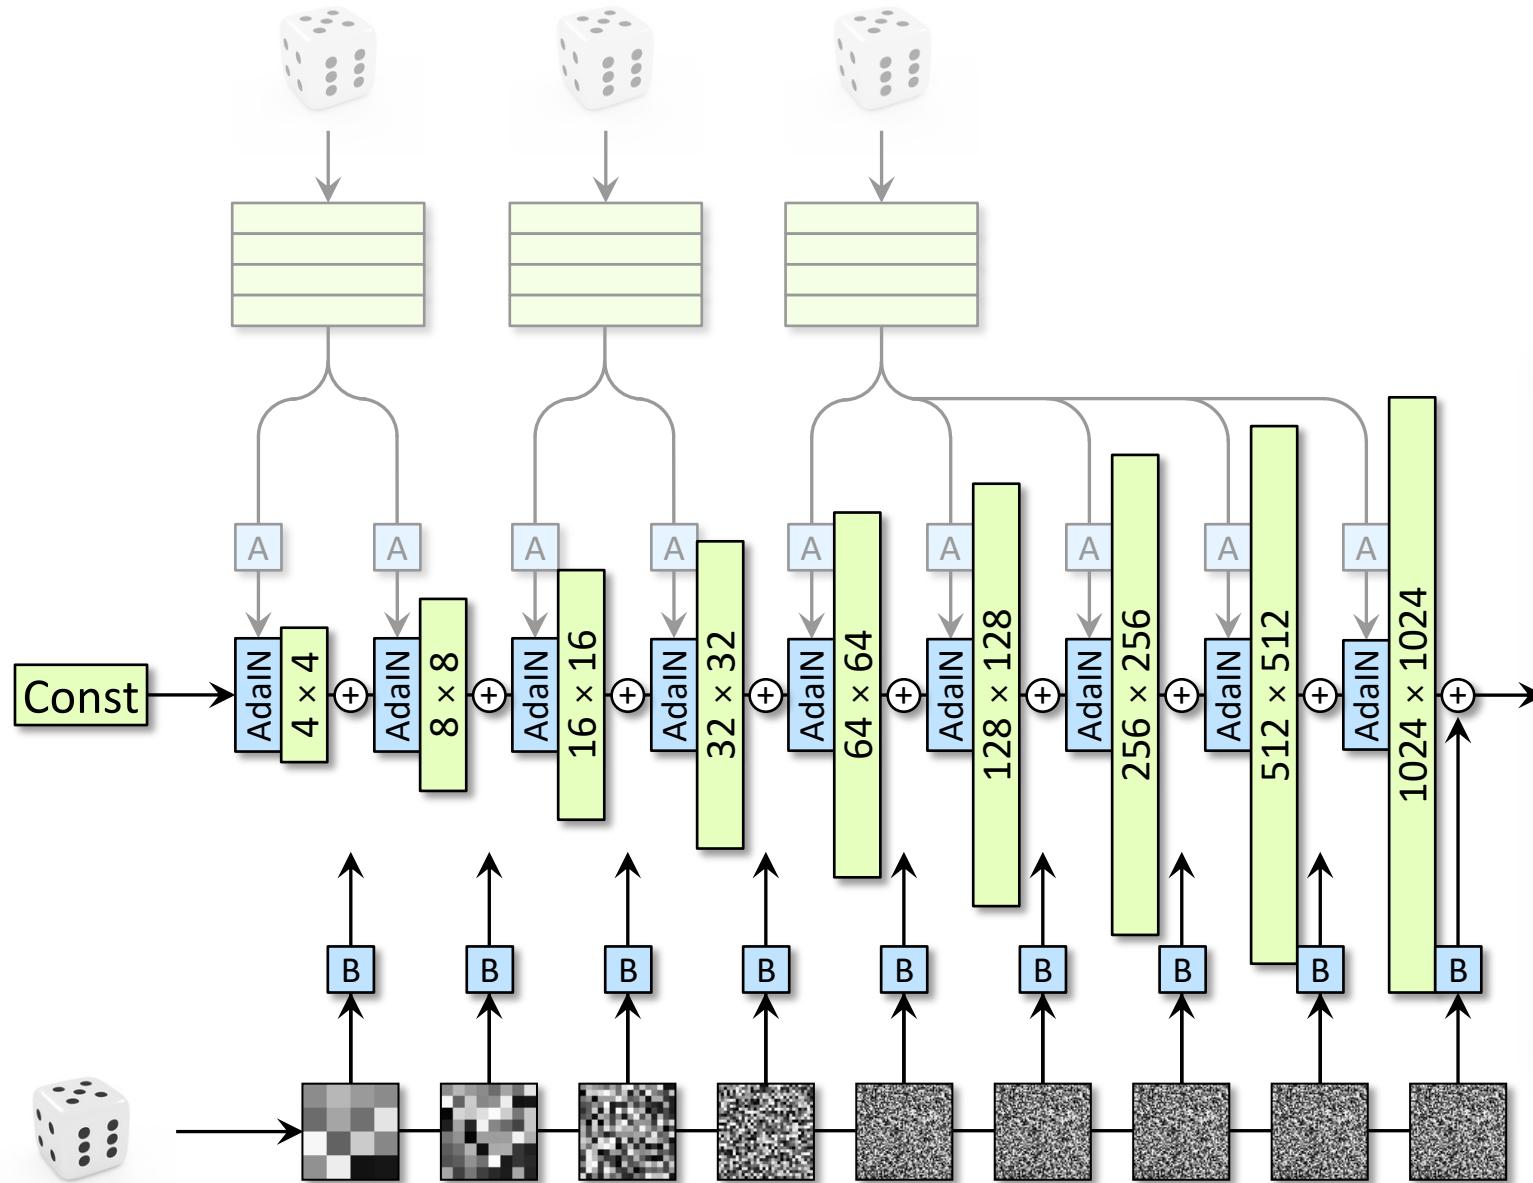
StyleGAN



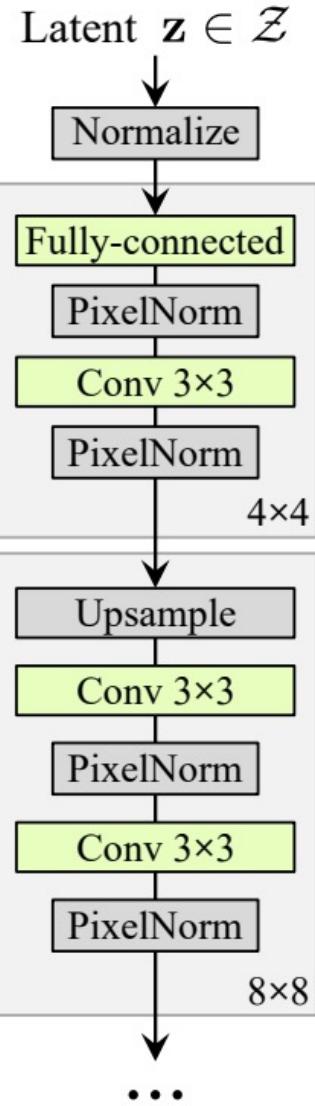
StyleGAN



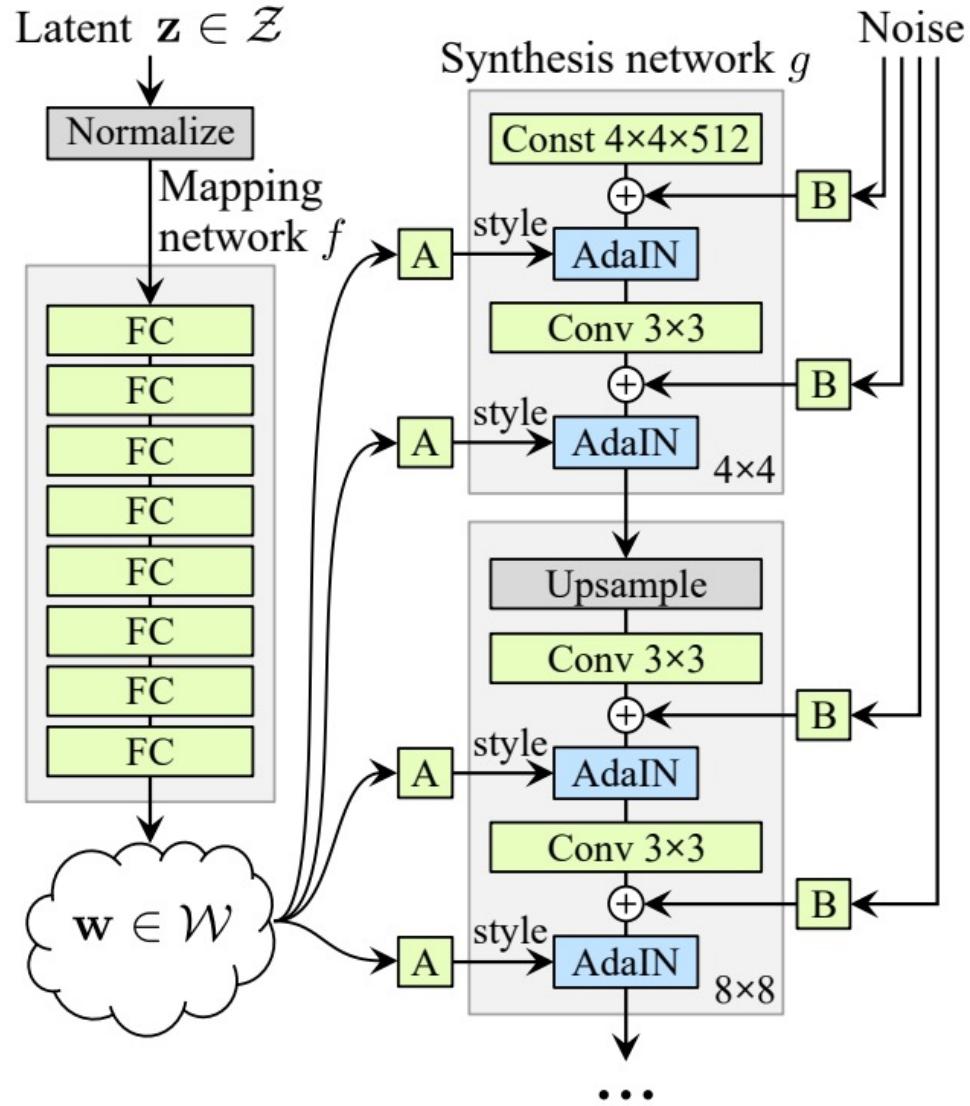
StyleGAN



StyleGAN



(a) Traditional



(b) Style-based generator

StyleGAN- Experiments

CelebA-HQ



30,000 images

Limited variation

Image quality not always ideal

Flickr-Faces-HQ



70,000 images

Varied age, ethnicity, background, ...

Consistent high quality

StyleGAN- Experiments

Coarse styles
 $(4^2 - 8^2)$



StyleGAN- Experiments

Source A: gender, age, hair length, glasses, pose



Source B:
everything
else

Result of combining A and B

StyleGAN- Experiments

Coarse styles
 $(4^2 - 8^2)$



Middle styles
 $(16^2 - 32^2)$



Fine styles
 $(64^2 - 512^2)$

StyleGAN- Experiments



Cars

Thank you!
Q & A